

resource-request algo for process P_i ,
 request = request vector for process P_i .

If $\text{Request}_i(j) = k$ then process P_i wants k instances of resource type R_j .

1) If $\text{request}_i \leq \text{Need}_i$ go to step 2. otherwise raise error condition, since process has exceeded its maximum claim.

2) If $\text{request}_i \leq \text{Available}$, go to step 3. otherwise P_i must wait, since resources are not available.

3) Relocate requested resources to P_i by modifying the state as follows.

$$\text{Available} = \text{Available} - \text{Request}_i;$$

$$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i;$$

$$\text{need}_i = \text{need}_i - \text{Request}_i;$$

If safe \Rightarrow the resources are allocated to P_i .

If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored.

* $P \rightarrow 10$ instances $B \rightarrow 5$ instances $C \rightarrow 7$ instances
 $(P_1, P_2, P_3, P_4, P_5)$

	Allocation			<u>Max</u>	Available	
	A	B	C	A	B	C
P ₀	0	1	0	7	5	3
P ₁	2	1	1	3	2	2
P ₂	3	0	2	9	0	2
P ₃	2	1	1	2	2	2
P ₄	1	1	2	4	3	3

Need

	A	B	C
P ₀	7	4	3
P ₁	1	1	1
P ₂	6	0	0
P ₃	0	1	1
P ₄	3	2	1

P₂

Allocation

P _i	A	B	C
P ₀	2	1	1
P ₁	1	1	1
P ₂	2	1	1
P ₃	0	1	1
P ₄	3	2	1

Available

P _i	A	B	C
P ₀	2	1	1
P ₁	1	1	1
P ₂	2	1	1
P ₃	1	0	0
P ₄	0	1	1

need

P _i	A	B	C
P ₀	1	1	1
P ₁	1	1	1
P ₂	1	1	1
P ₃	0	1	1
P ₄	1	1	1

New all.

Available

P _i	A	B	C
P ₀	2	1	1
P ₁	1	1	1
P ₂	2	1	1
P ₃	1	0	0
P ₄	0	1	1

new available

P _i	A	B	C
P ₀	1	1	1
P ₁	2	2	2
P ₂	1	1	1
P ₃	0	1	1
P ₄	1	1	1

P₃

Allocation

P _i	A	B	C
P ₀	2	1	1
P ₁	1	1	1
P ₂	2	1	1
P ₃	1	0	0
P ₄	0	1	1

available

P _i	A	B	C
----------------	---	---	---

need

P _i	A	B	C
----------------	---	---	---

New all.

Available (dummy exec)

P _i	A	B	C
----------------	---	---	---

New Available (after exec)

P _i	A	B	C
----------------	---	---	---

P2	Allocation			Available			Need		
	A	B	C	A	B	C	A	B	C
	3	0	2	6	3	3	6	0	0

new Allo. Available

A	B	C	A	B	C
9	0	2	0	3	3

new Avail

A	B	C
9	3	8

P3	Allocation			Need		
	A	B	C	A	B	
	0	1	0	9	3	5
				7	4	3

Not safe state

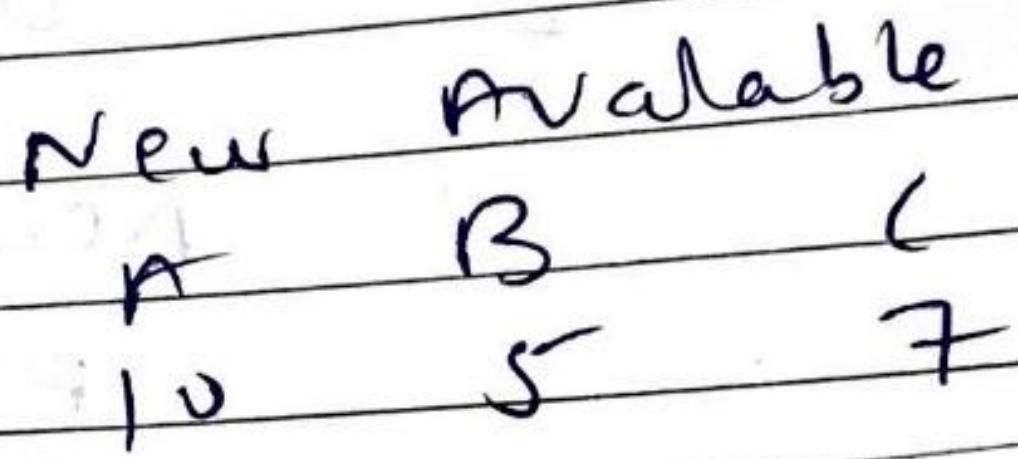
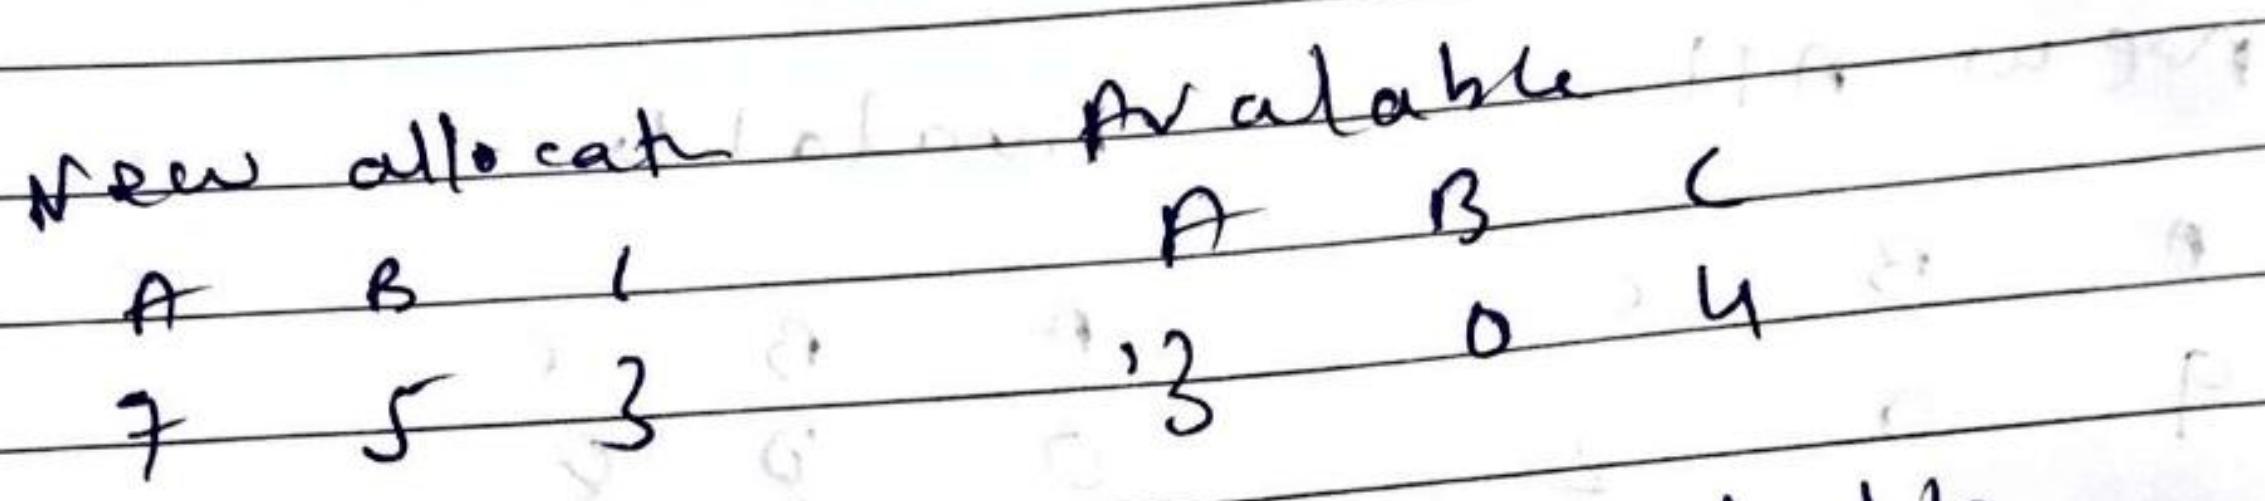
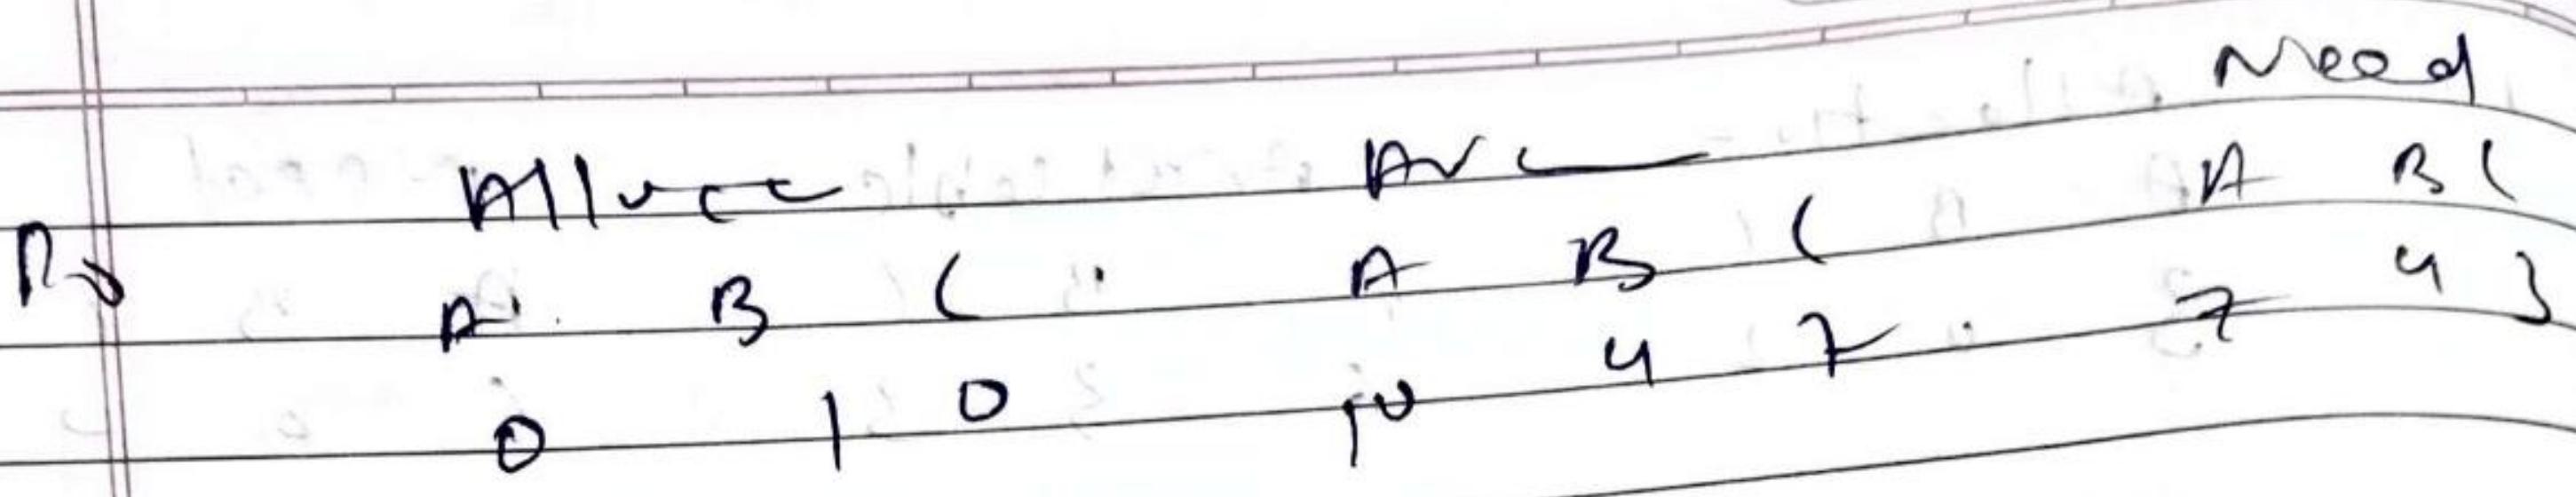
P4	allo.			available		
	A	B	C	A	B	
	1	1	2	9	3	5
				3	2	1

New Allo. Avail

A	B	C	A	B	C
4	3	3	6	1	4

New Avail

A	B	C
10	4	7



Methods for handling deadlocks

- Deadlock prevention

- mutual exclusion

- hold and wait

- no preemption

- circular wait

Deadlock avoidance

- Resource allocation graph (with claim edges)

Banerjee's algorithm

- safe state

- Data structure used

- safety algorithm

- Resource request algorithm

o Deadlock detection :- Allow system to enter deadlock state

Detection algorithm

Recovery scheme

o Data structures used :-

Available :- A vector of length m indicates no of available resources of each type.

Allocated : An $n \times m$ matrix defines no of resources of each type currently

Request:

o Detection algorithm : 1) Let work and finish be vectors of length m and n, respectively. Initialize

a) $\text{work} = \text{Available}$

b) For $i = 1, 2, \dots, n$, if $\text{Allocation}_{i,i} > 0$ then $\text{Finish}[i] = \text{false}$, else $\text{Finish}[i] = \text{true}$

2) find an index i such that both $\text{Finish}[i] = \text{false}$

$\text{Request}_i \leq \text{work}$

If no such i exists, go to step 4

3) $\text{work} = \text{work} + \text{Allocation}_i$

$\text{Finish}[i] = \text{true}$

go to step 2

4) If $\text{Finish}[i] = \text{false}$, for some i , $1 \leq i \leq n$, then the system is in

deadlock state

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	0	0	0	0	1	0
P1	2	0	0	2	0	2	1	0	0
P2	3	0	3	0	0	0	0	0	0
P3	2	1	1	1	0	0	0	0	0
P4	0	0	2	0	0	2	0	0	0

1) In this, work = [0, 0, 0] and finish = [false, false, false, false, false]

(P0, P2, P1, P3, P4)

P0	New Available	A	B	C
		0	1	0

P2	New Available	A	B	C
		3	1	3

P1	New Available	A	B	C
	A : B +	0	0	0
New All.	Nearest	3	1	3
A B C	A B C			
2 0 2	2 0 2			
4 0 2				

New Avail	A	B	C
1 1 1			

New avl
1 1 1

P ₃	New Allocated			Request			Available			
	A	B	C	A	B	C	A	B	C	
	3	1	1		1	0	0	5	1	3
				Available						
<hr/>										
				A	B	C				
				4	1	3				

New Allocation

A	B	C
7	2	4

P ₄	New Allocated	Request	Available
	A B C	A B C	A B C
	0 0 4	0 0 2	7 2 4

Available

A	B	C
7	2	2

New Allocation

A	B	C
7	2	6

• recovery from Deadlock :

- (1) Process termination
- (2) Resource preemption

→ About all deadlock processes
kill one process at a time until
deadlock cycle is eliminated.

Page No.	
Date	

In which order to abort?

- Priority of process
- How long process has computed and how much longer to completion.
- Resources the process has used
- Resources the process needs to complete its task
- How many processes will need to be terminated.
- Is process interactive or batch?

Q1) d) ✓

Q2) c) X

Q11) d) X

Q3) a) X

Q12) c) ✓

Q4) a) ✓

Q13) a) ✓

Q5) c) ✓

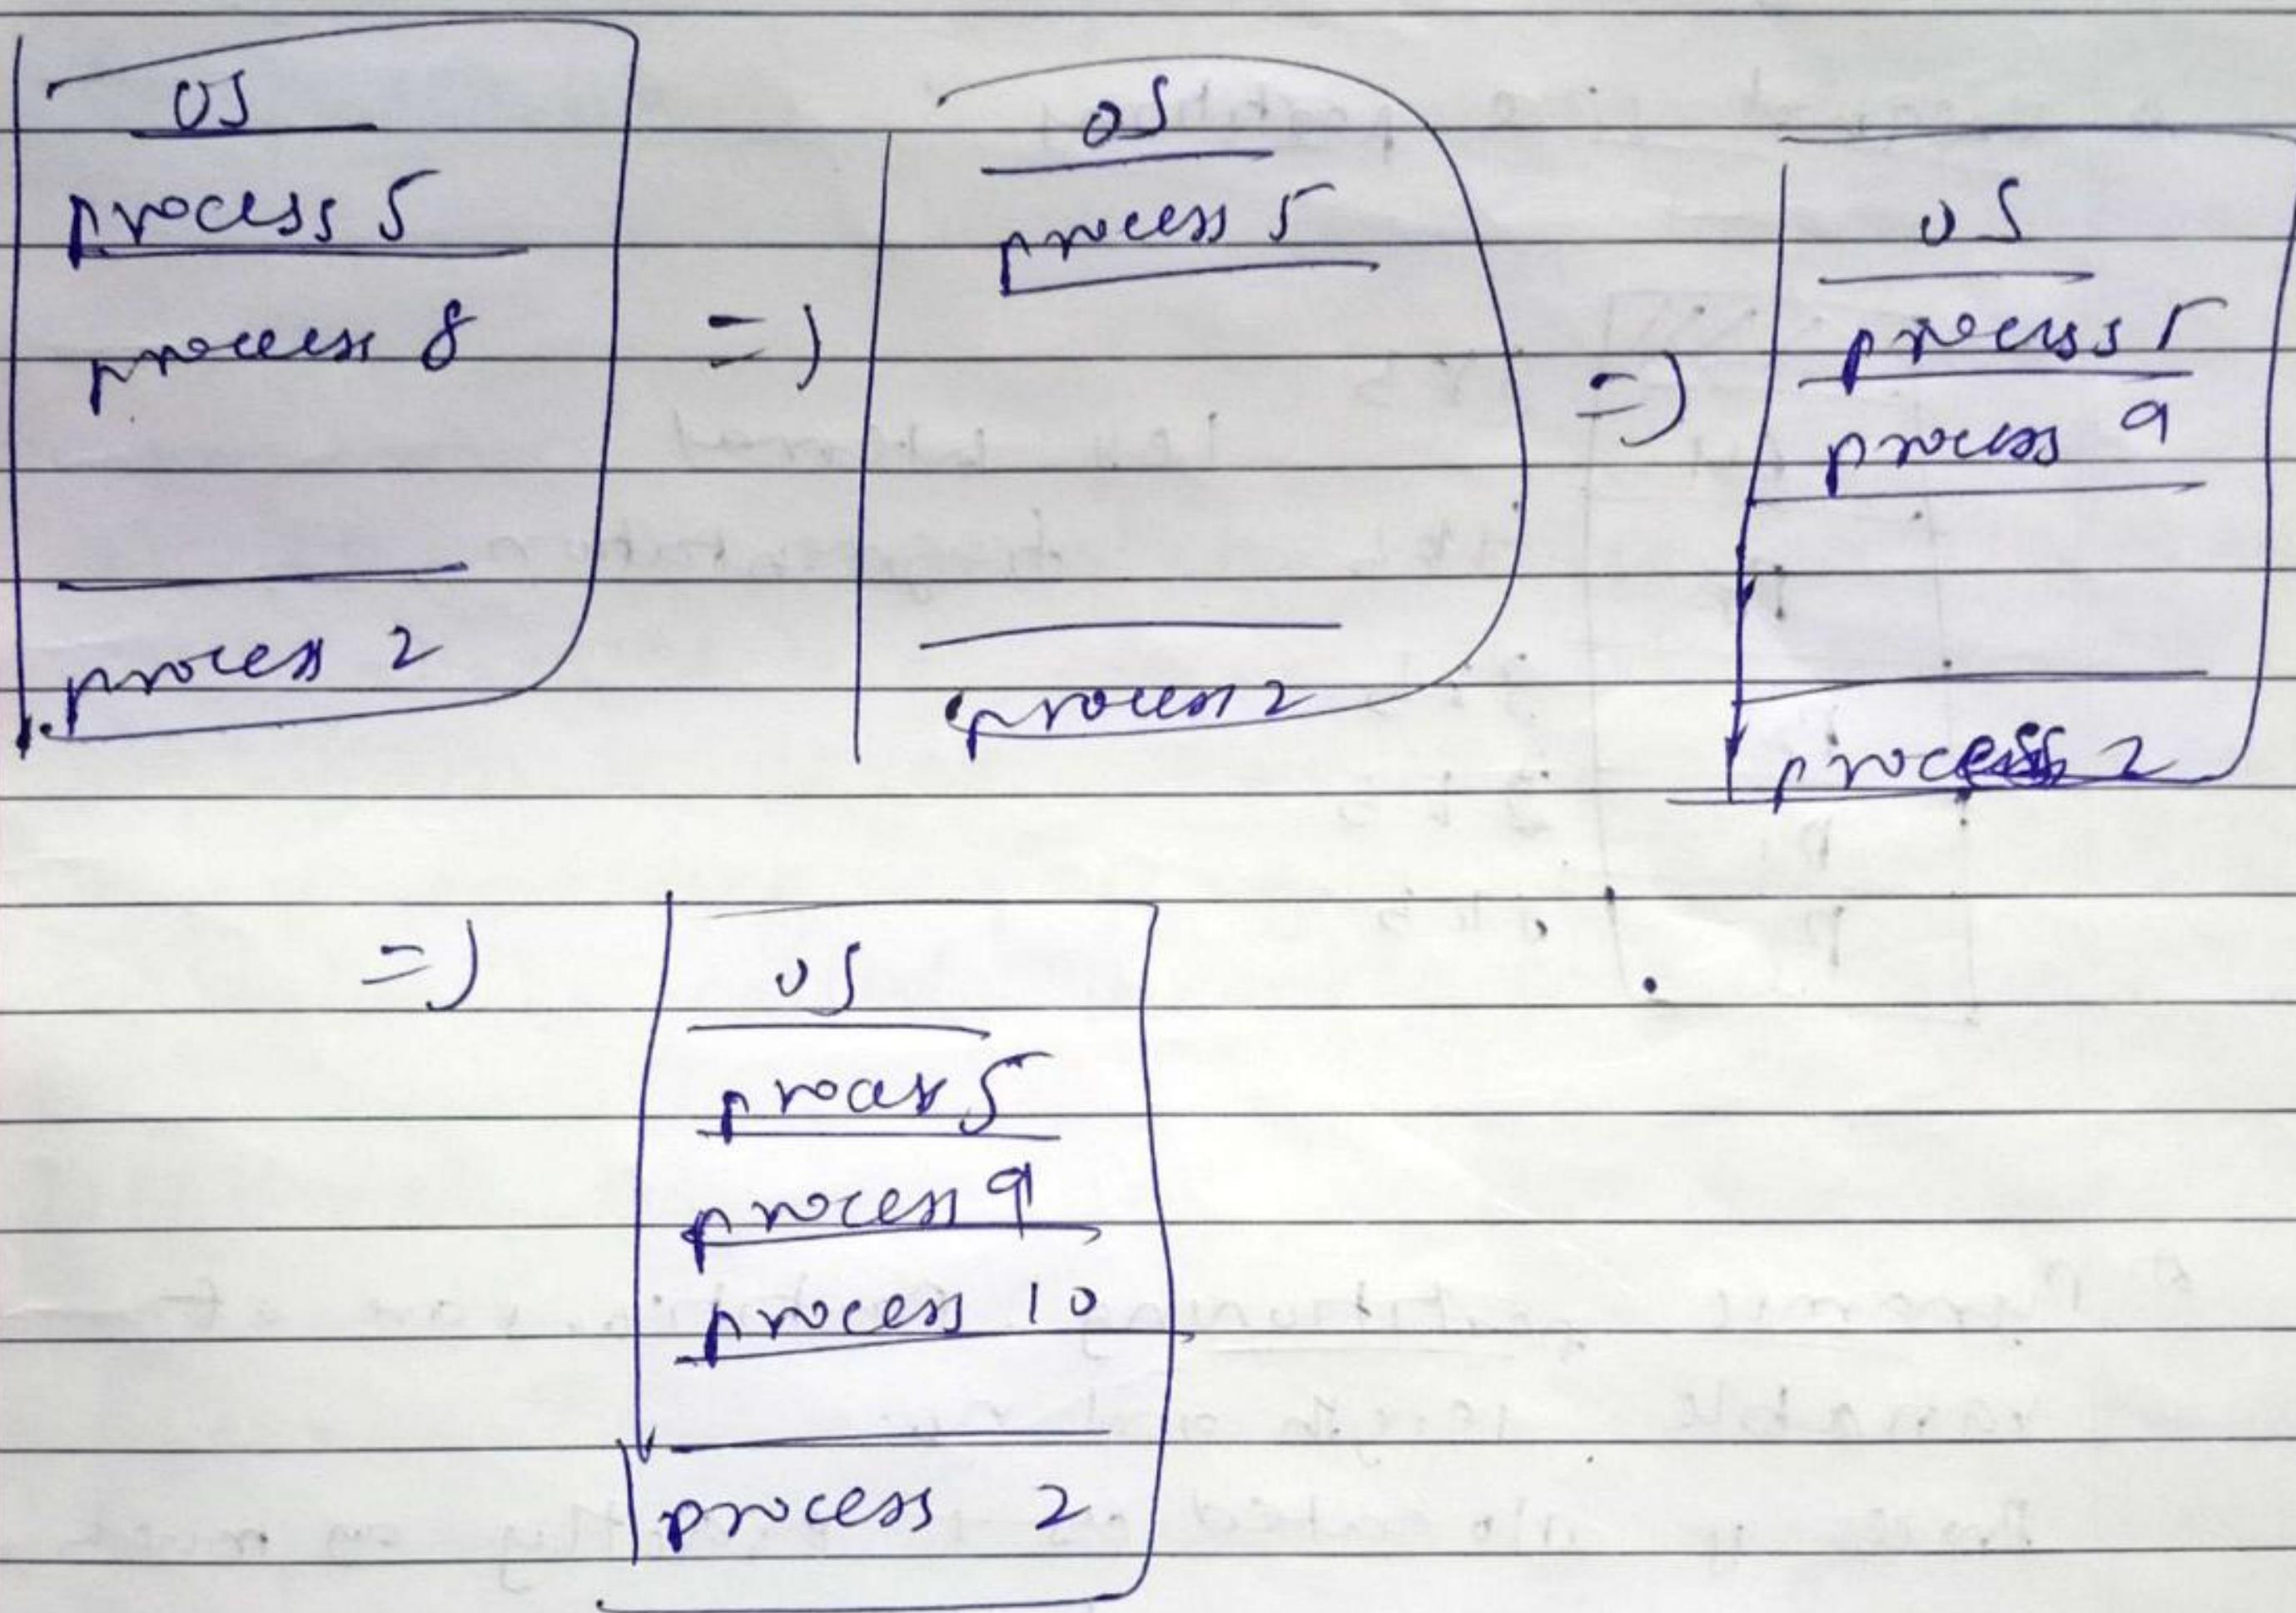
Q6) a) ✓

Q7) b) a) ✓

Contiguous Allocation

- one-block of available memory

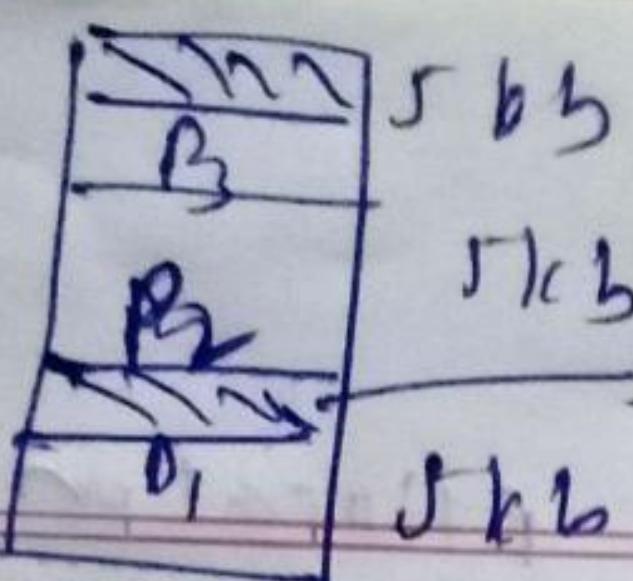
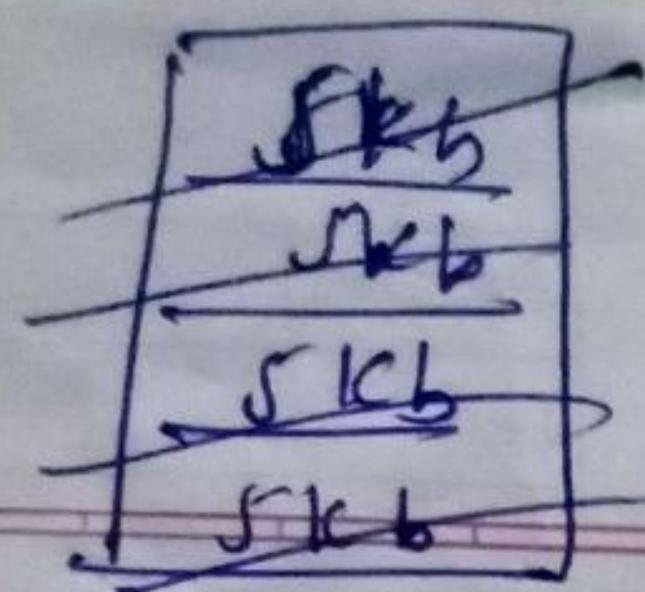
when a process arrives, it is allocated memory from a large hole enough to accommodate it
operating system maintains info about:
a) allocated partitions b) free partitions (holes)



- partitioning : 1) Fixed partitioning
2) Dynamic partitioning

Fixed partitioning:- 1) Equal-size partitions
2) Unequal-size partitions

Any process whose size is less than or equal to the partition size can be loaded into an available partition

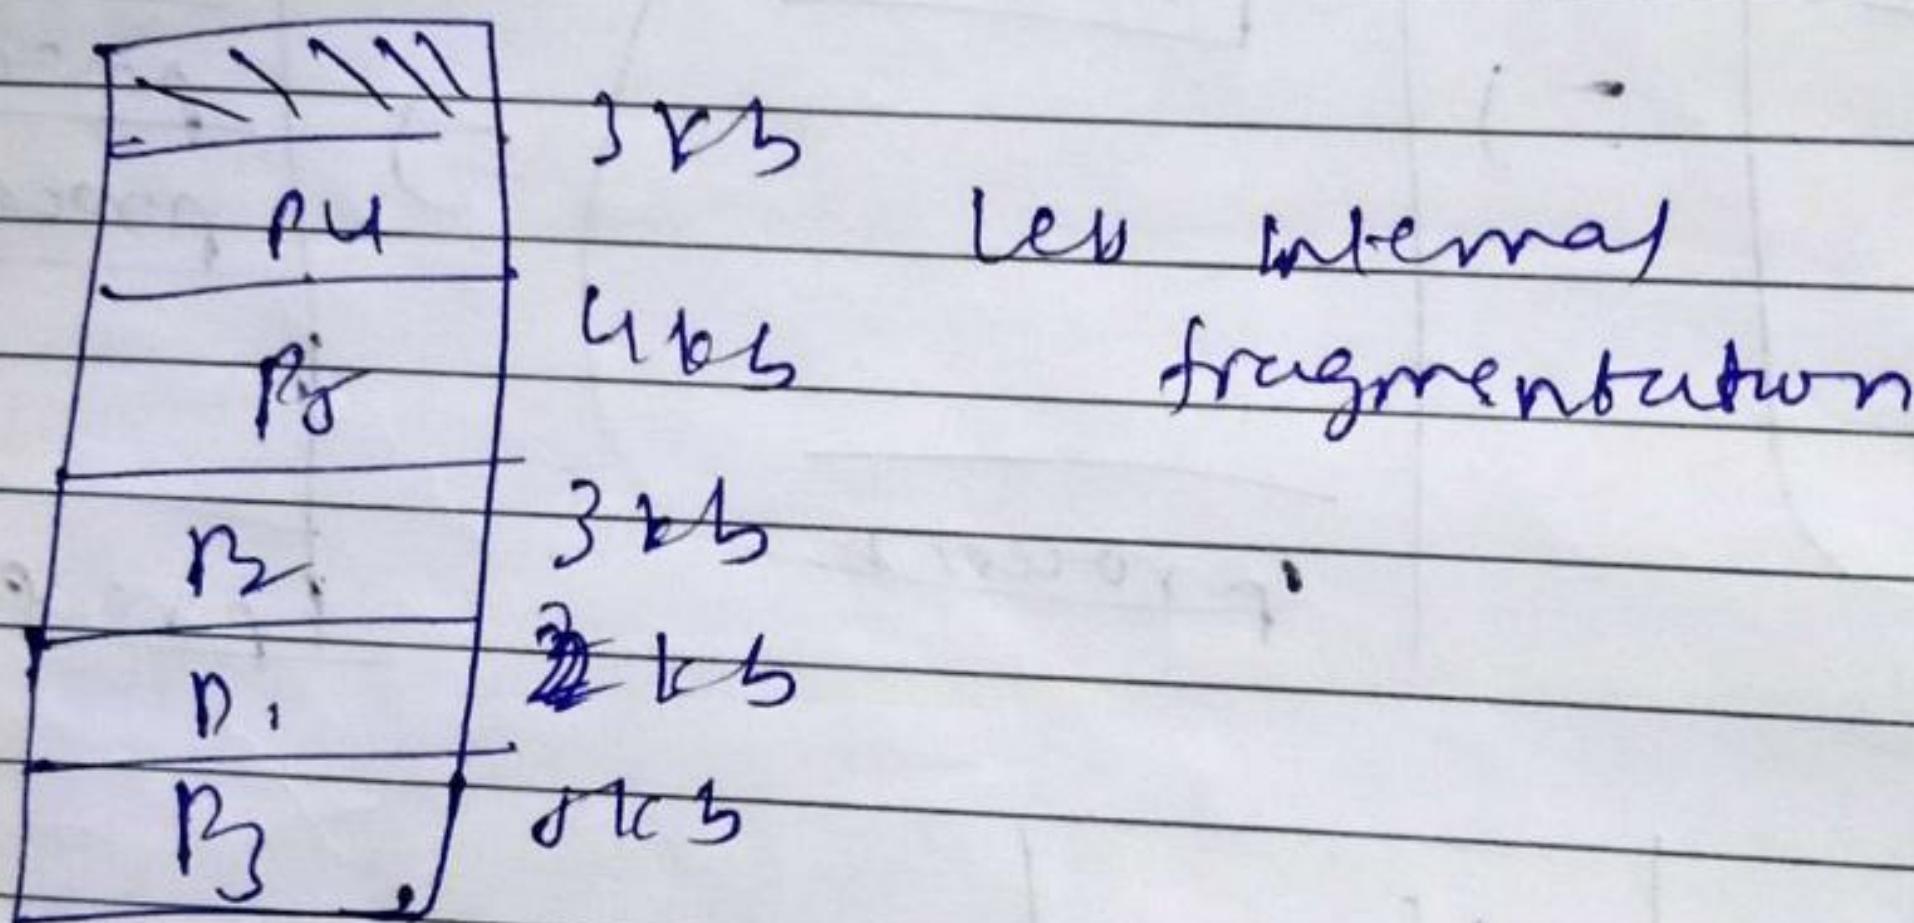


Internal fragmentation

If all partitions are full, the OS can swap a process out of a partition.

Main memory use is inefficient. Any program no matter how small, occupies an entire partition. This is called internal fragmentation.

* Unequal size partitions:



less internal fragmentation

* Dynamic partitioning: partitions are of variable length and no.

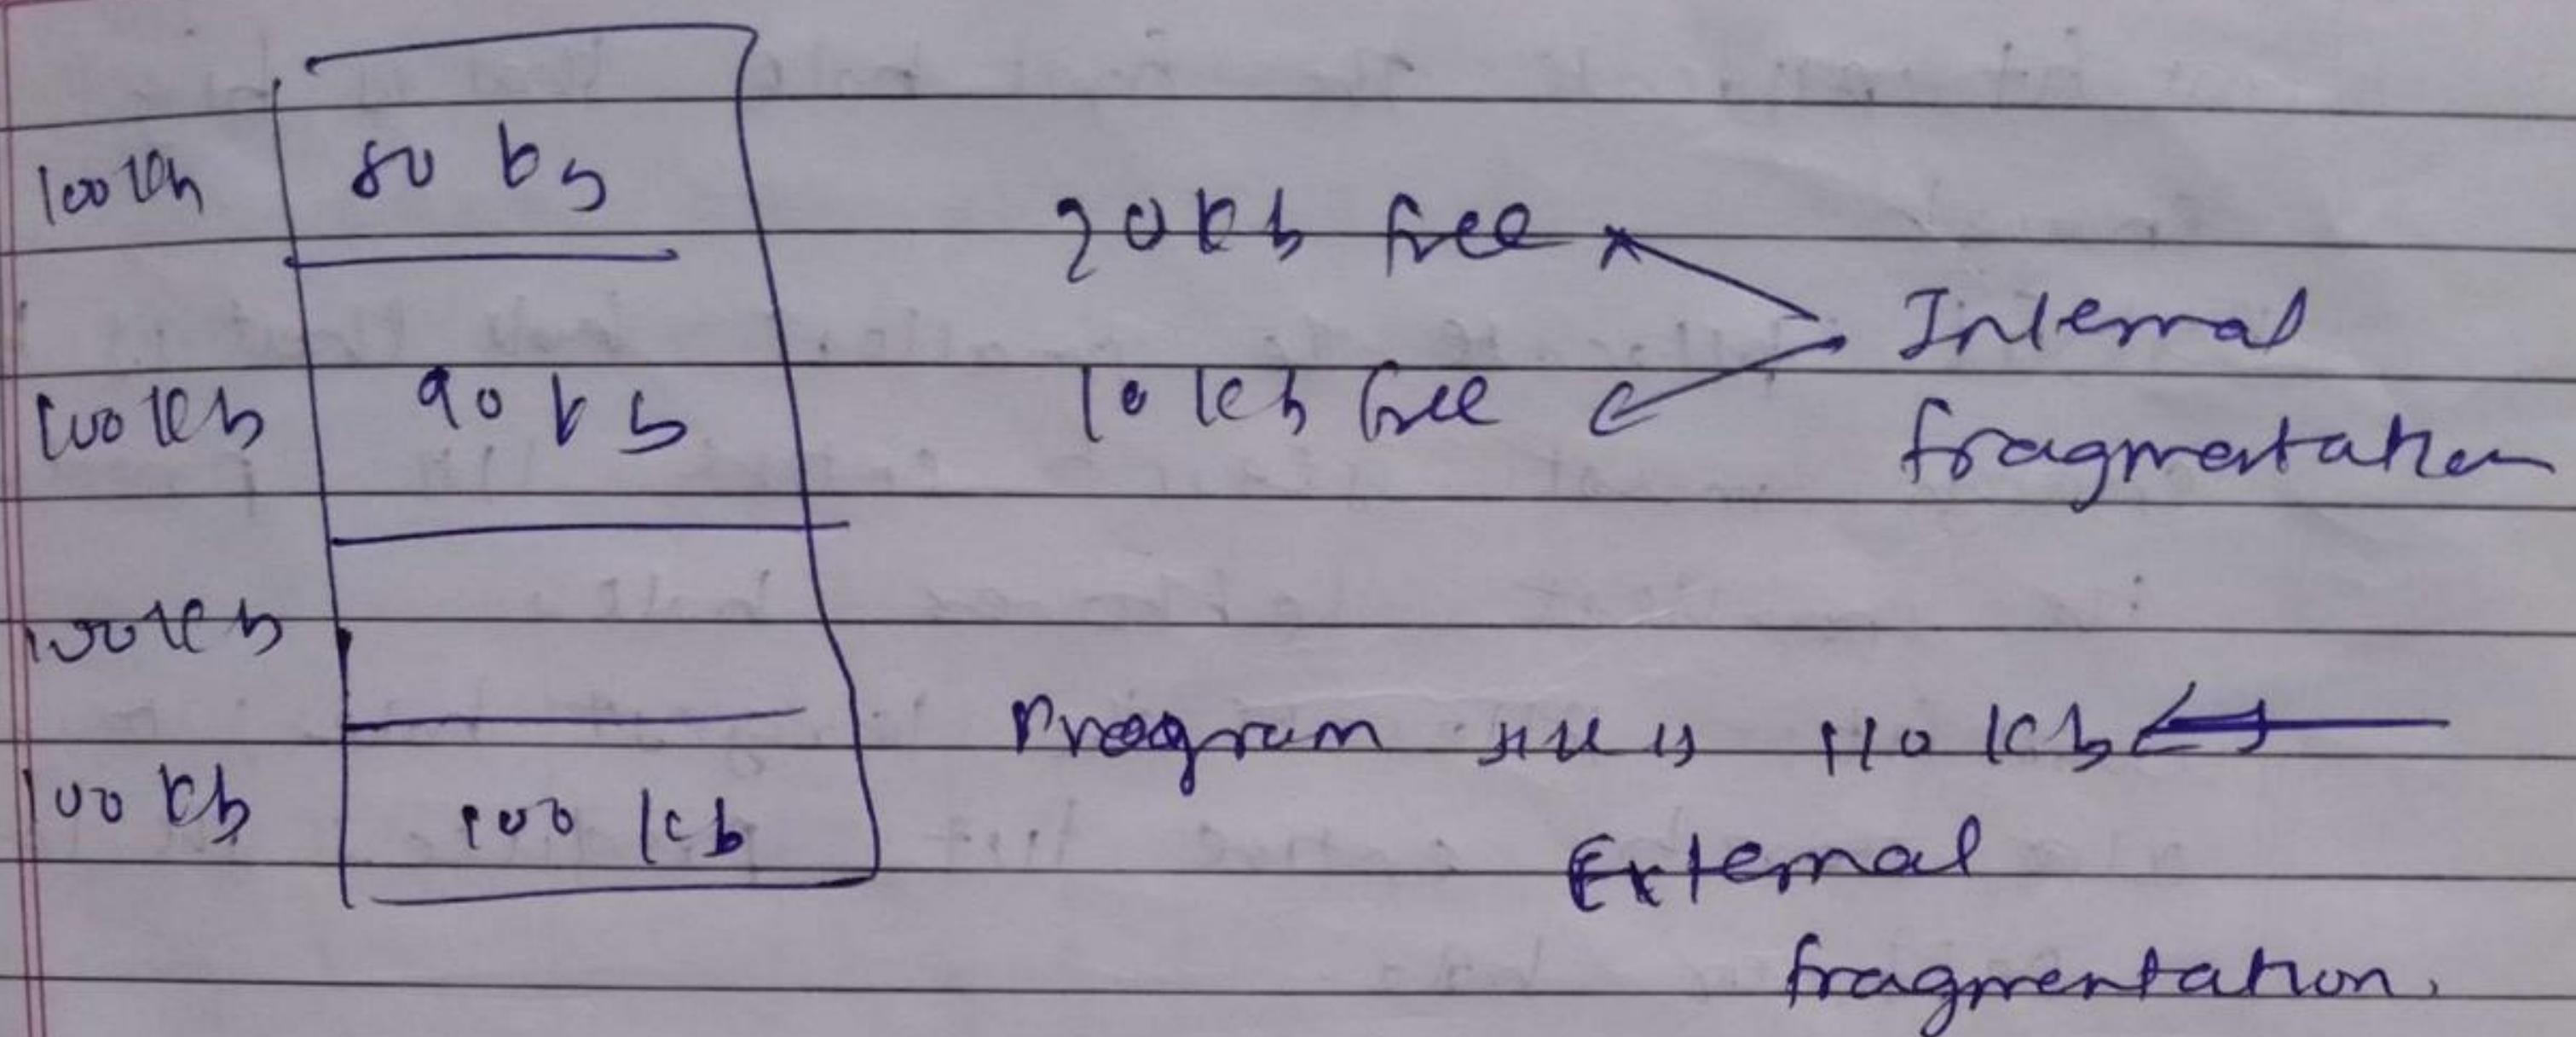
Process is allocated as + exactly as much memory is required.

Eventually get holes in the memory. This is called external fragmentation.

must we compaction to shift

are contiguous and all free processes so they memory is in one block

processes so they memory is in one block



Fragmentation

1) Internal fragmentation: when a program is allocated to a memory block, if that program is lesser than this memory block and remaining space goes wasted, this situation is called internal fragmentation.

2) External fragmentation :- when a memory cannot be allocated to a requesting process because the memory blocks are very small. The process size is more than any available hole.

Reduce external fragmentation by compaction

- shuffle memory contents to place all free memory together in one large block
- compaction is possible only if relocation is dynamic, and is done at execution time.

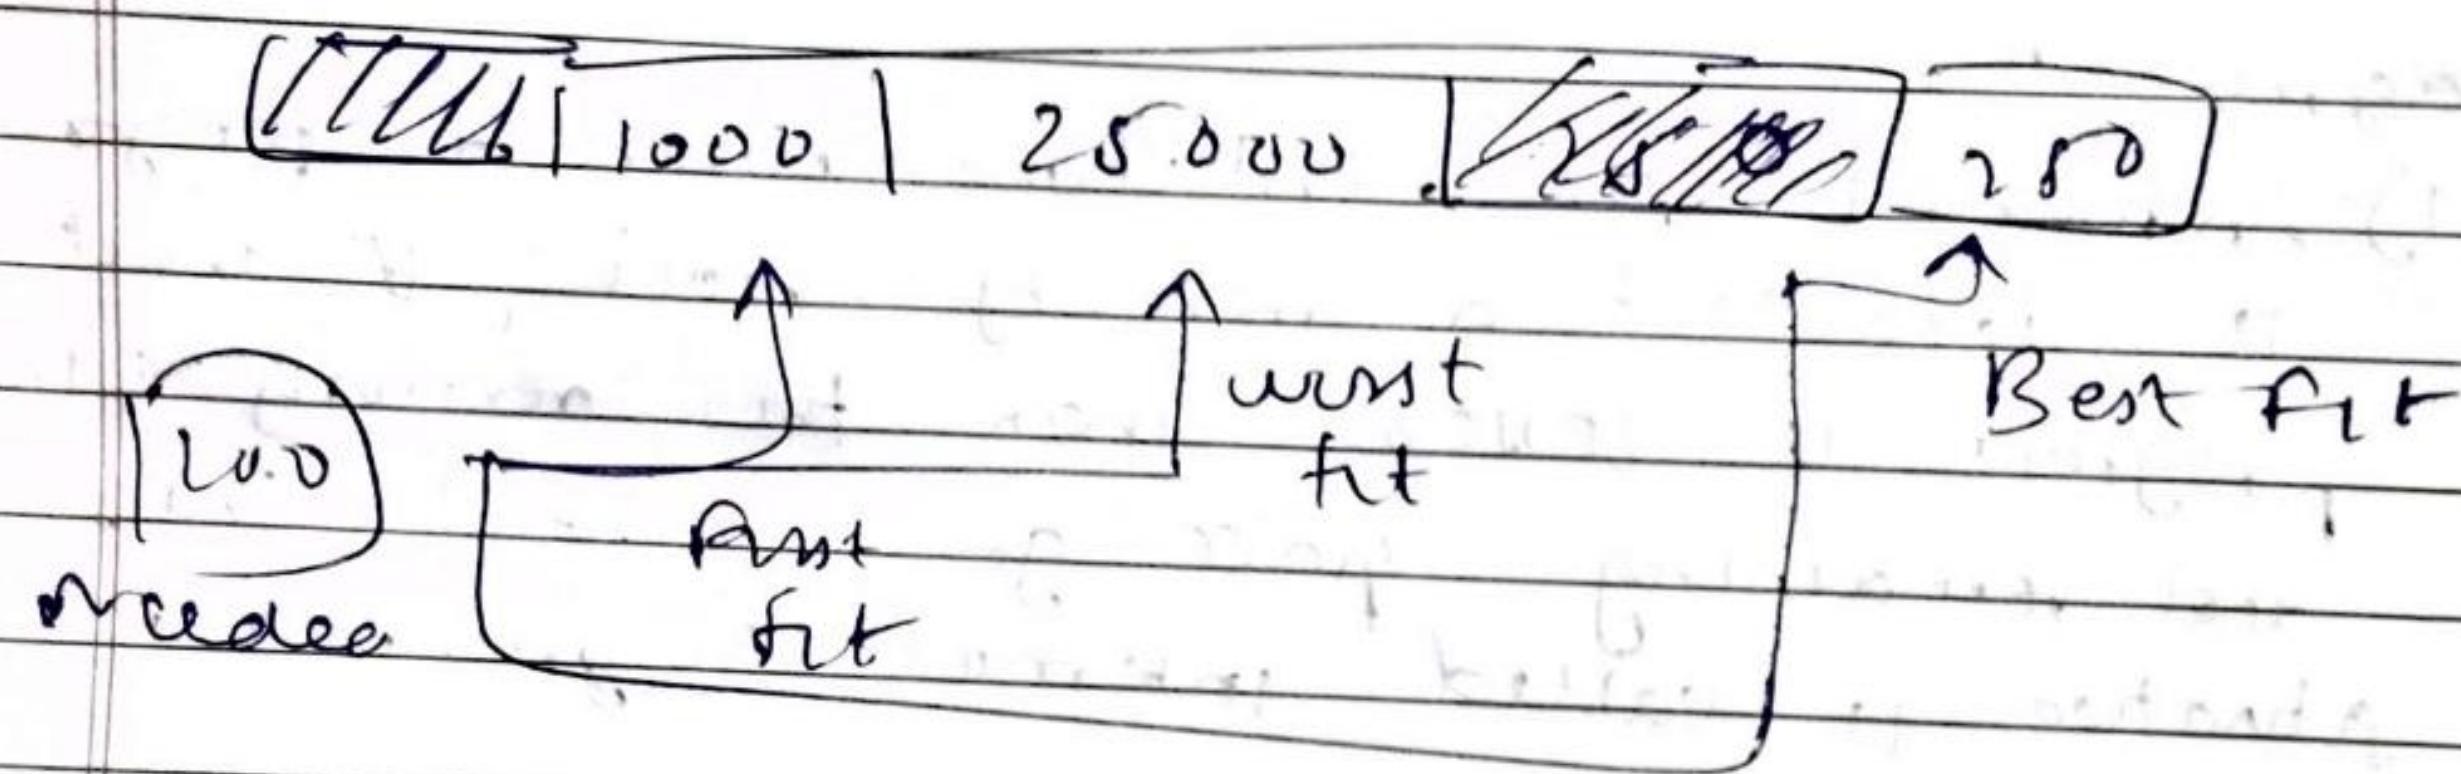
Dynamic Storage-Allocation problem :-

How to satisfy a request of size n from a list of free holes

First fit: Allocate the first hole that is big enough.

Best fit: Allocate the smallest hole that is big enough. must search entire list. produces the smallest leftover hole.

Worst fit: Allocate the largest hole; must also search entire list. produces the largest leftover hole.

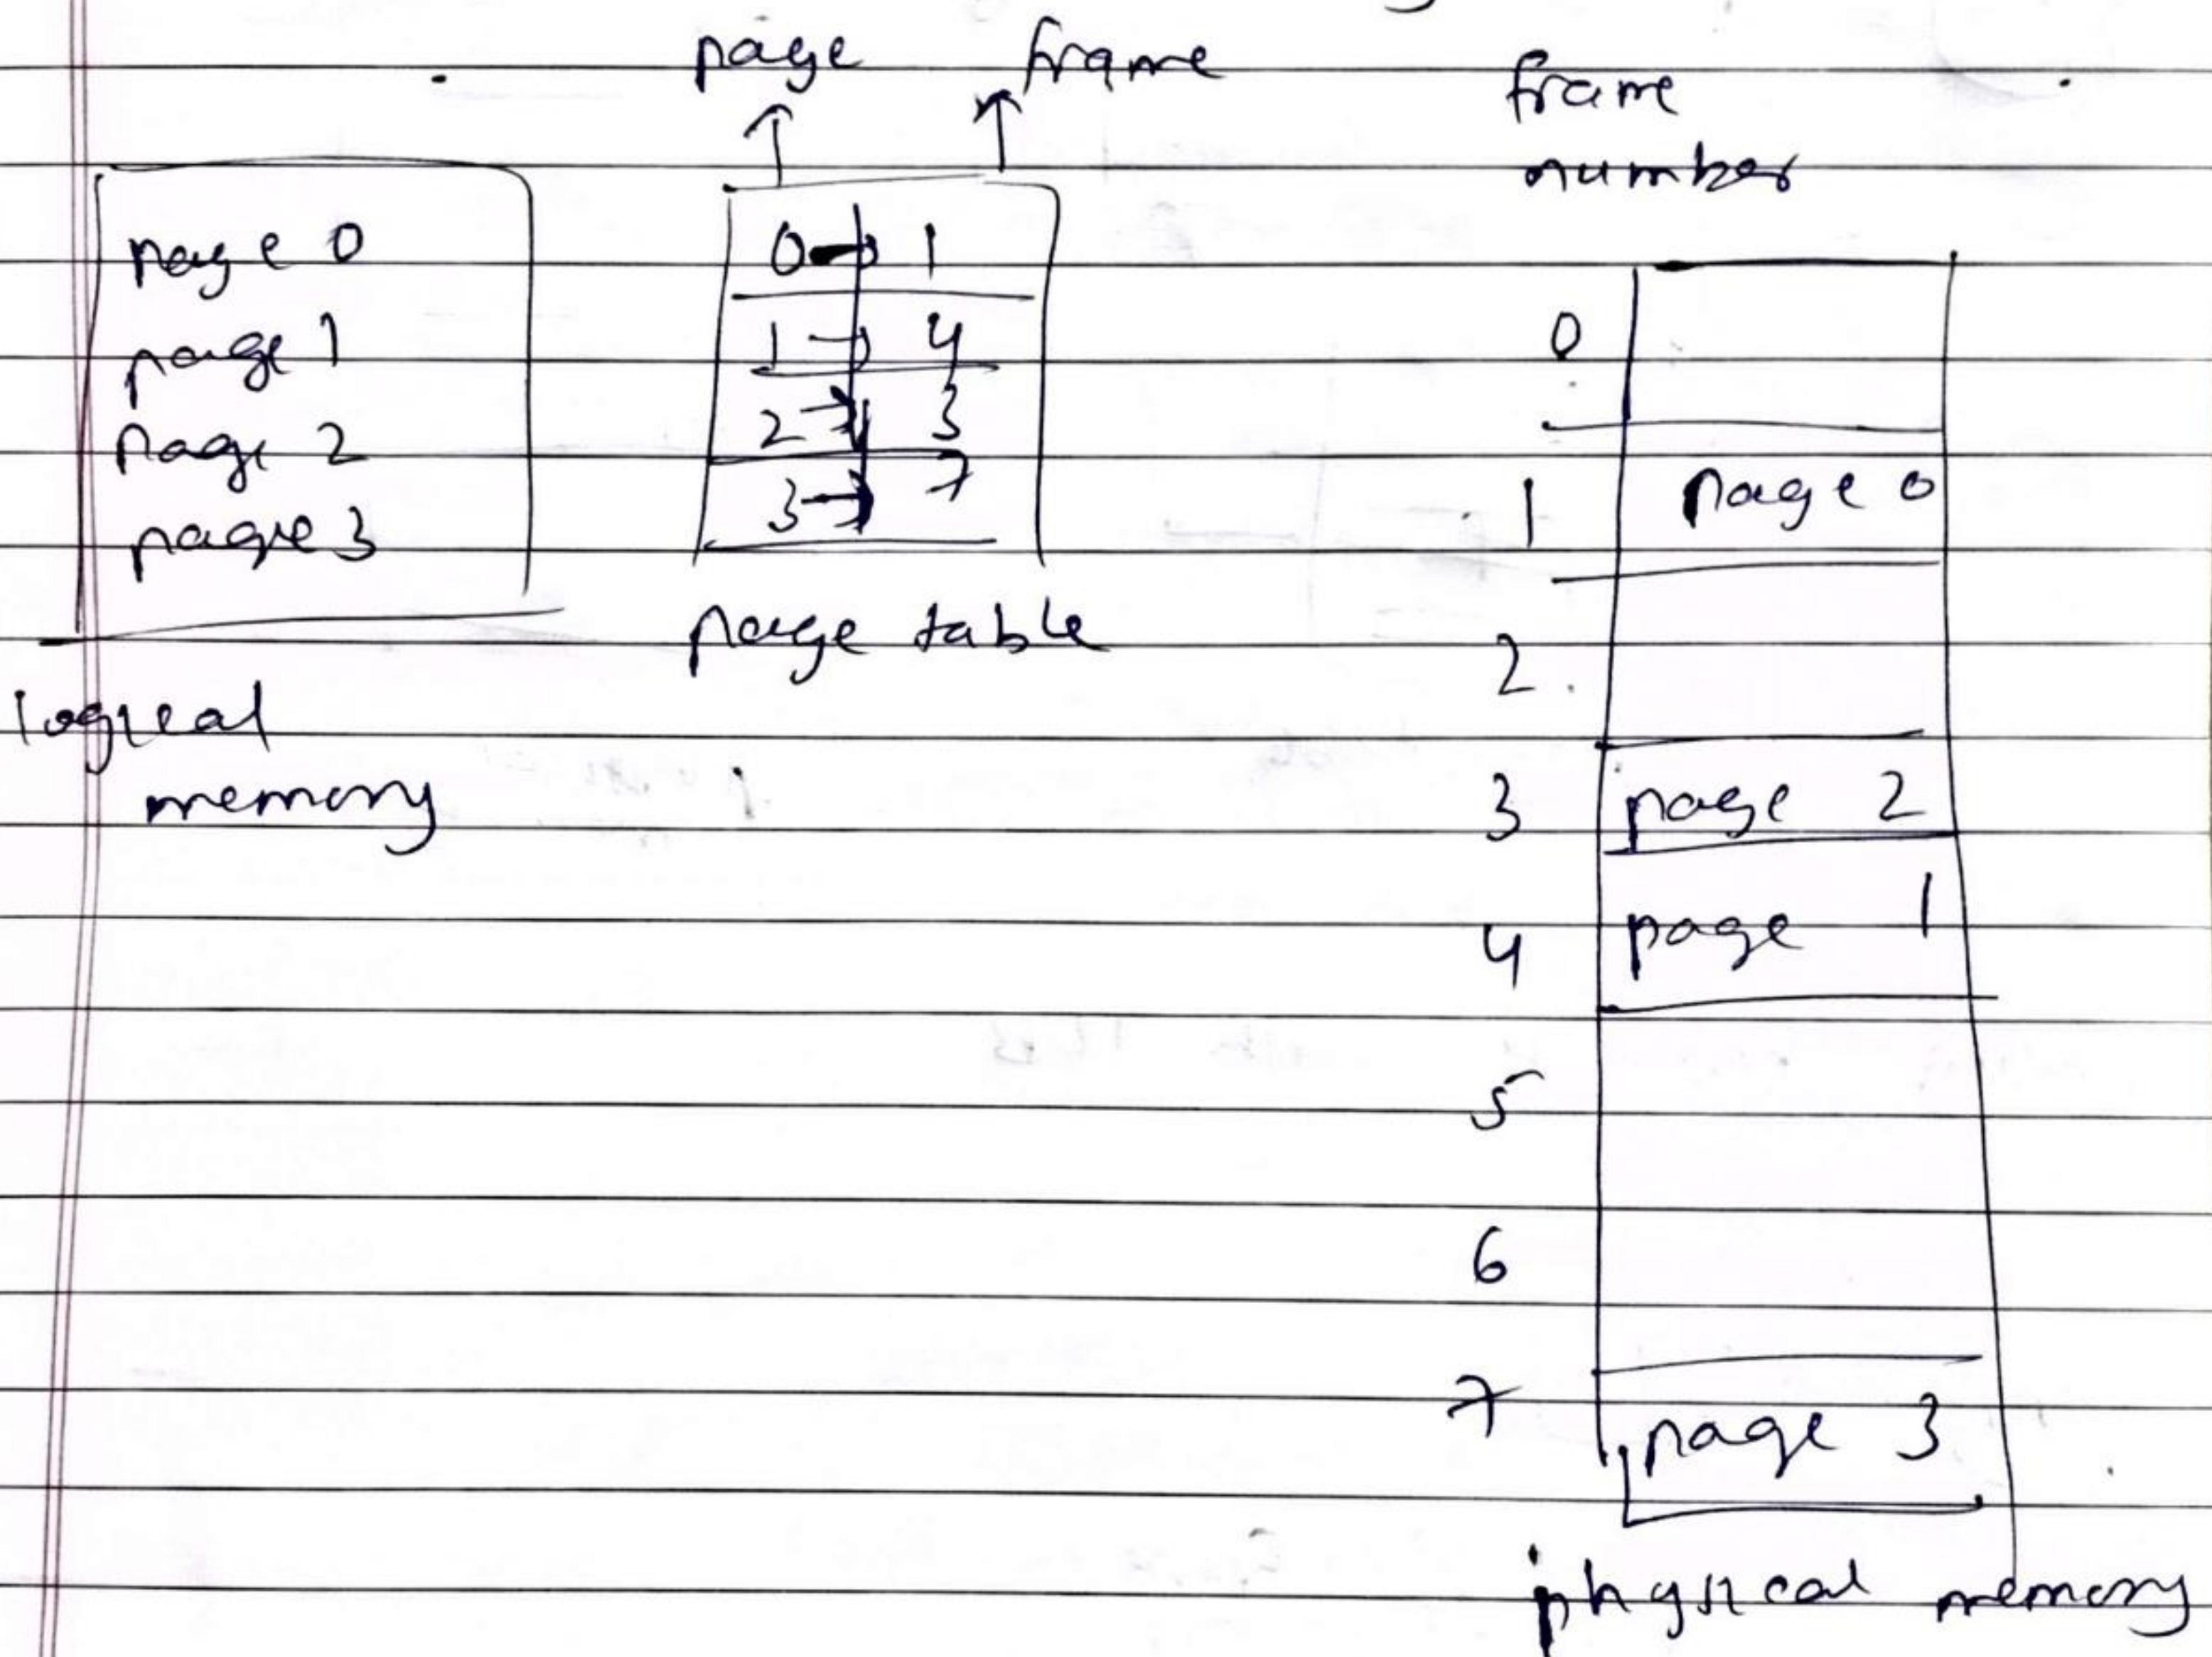


A Paging: Paging is a memory management scheme that permits the physical address space of a process to be non contiguous. Divide physical memory into fixed-size blocks called frames. Divide logical memory into blocks of same size called pages.

When a process is to be executed its pages are loaded loaded into any available memory frames. Set up a page table to translate logical addresses to physical addresses. Reduces internal fragmentation.

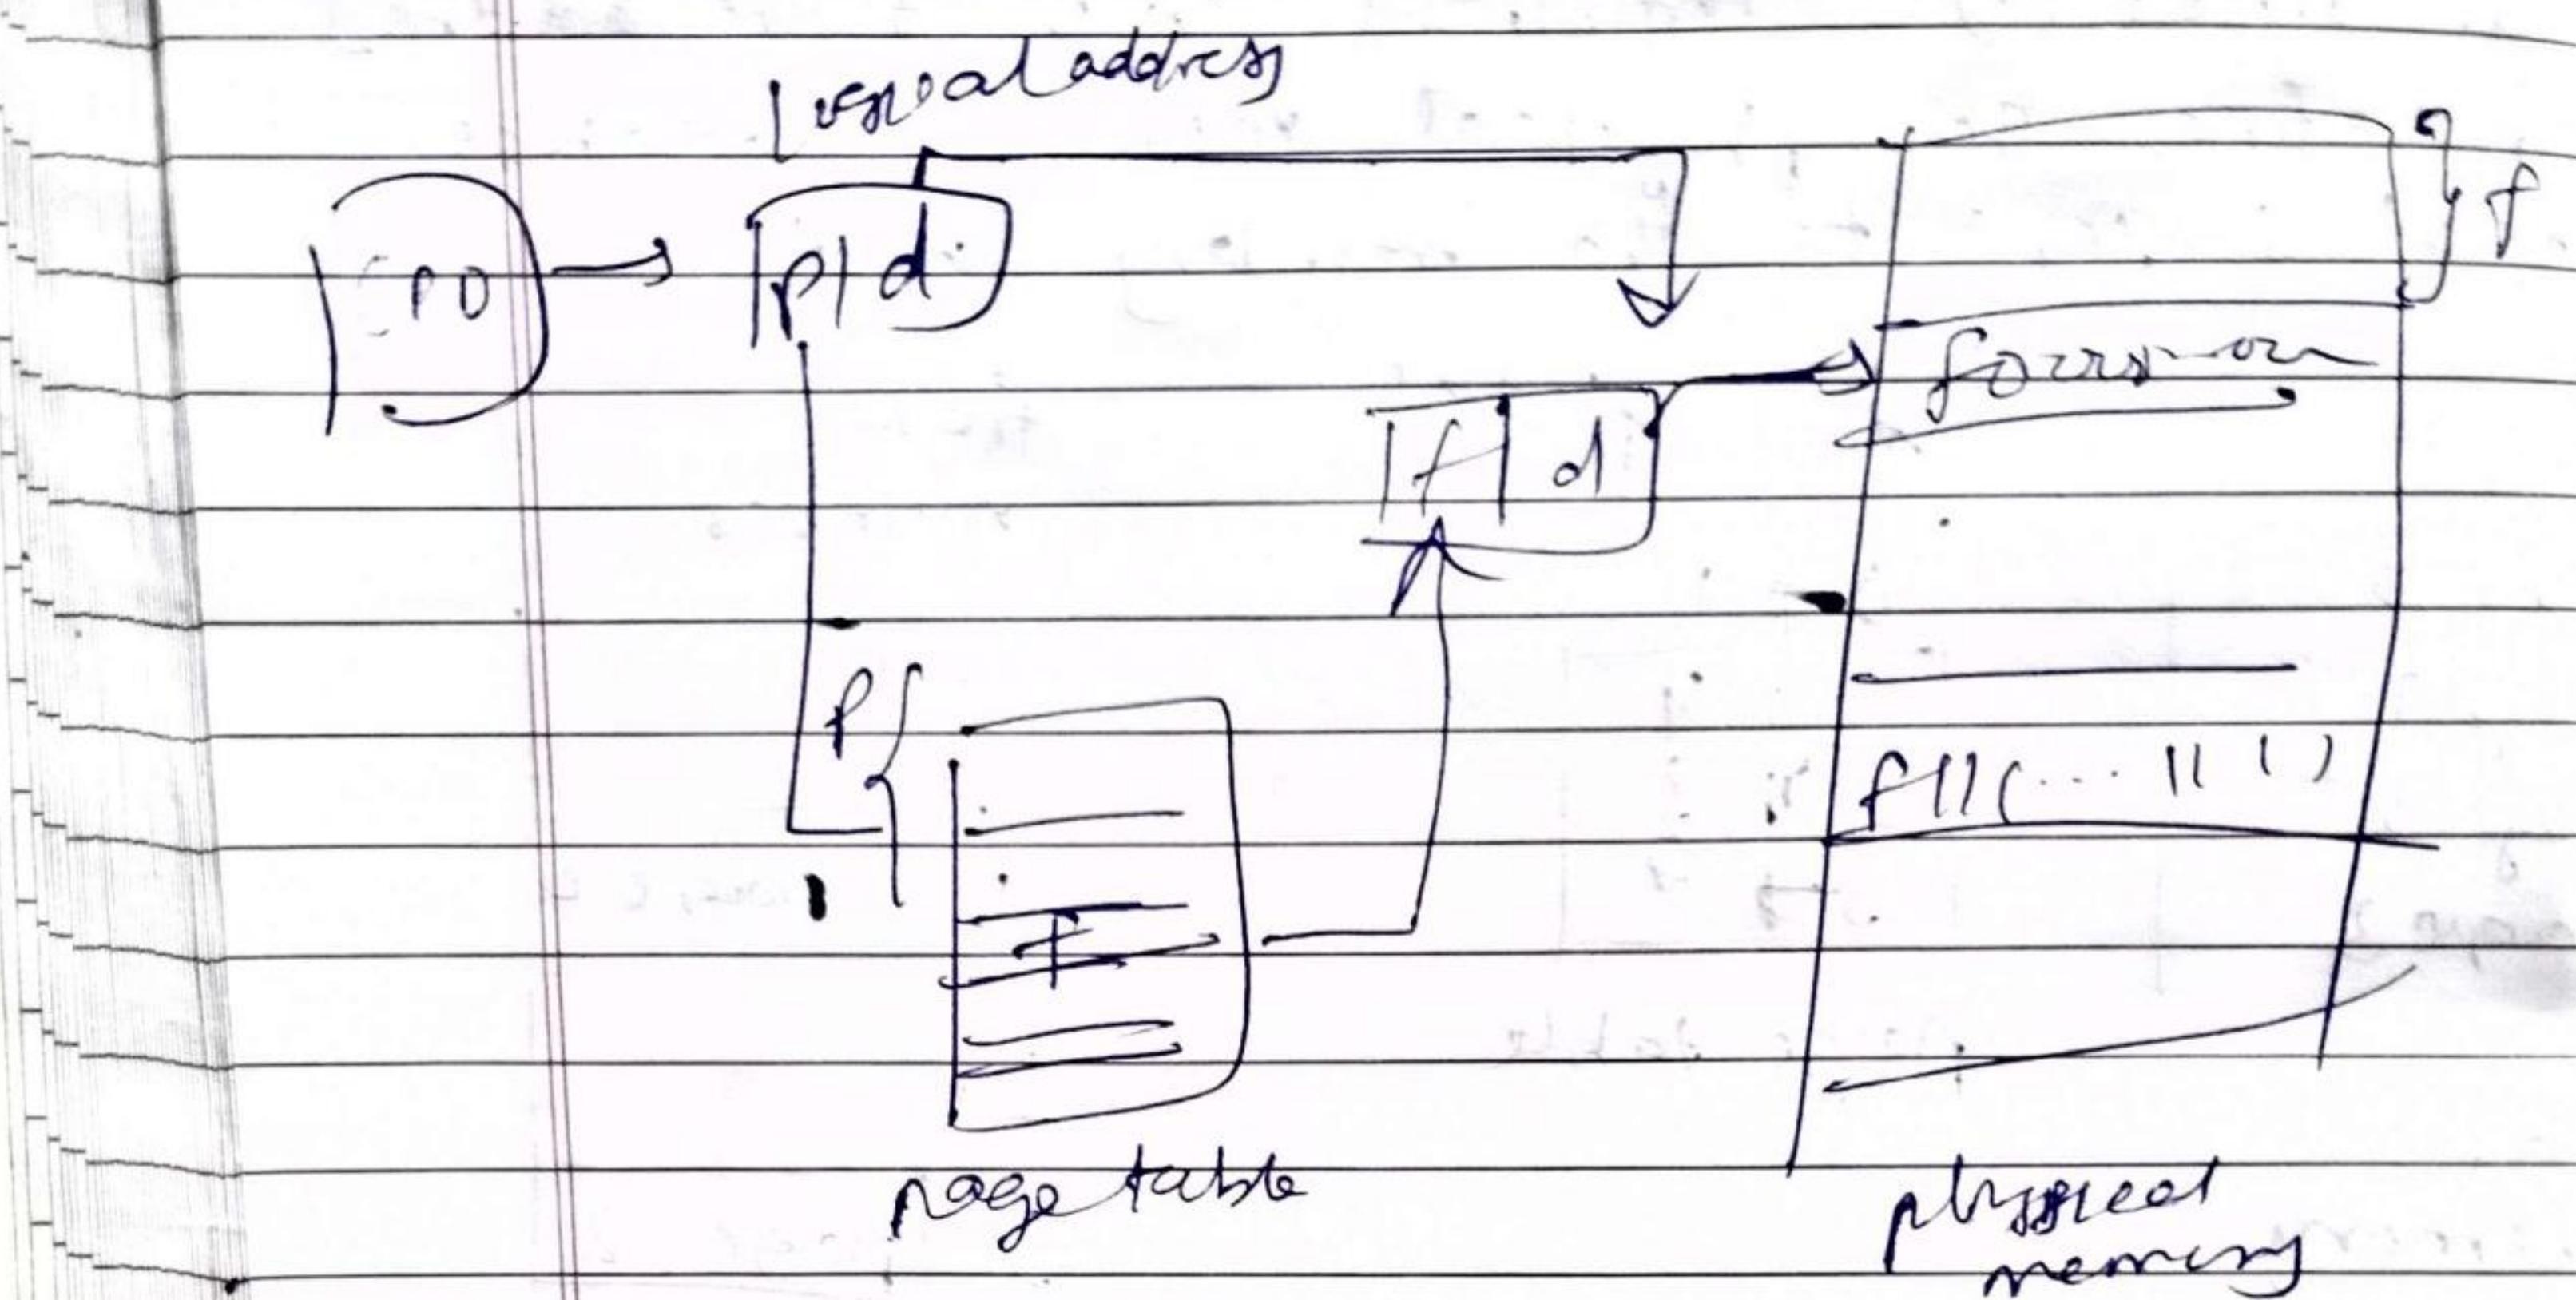
Address generated by CPU is divided into: page numbers (p) - used as an index into a page table which contains base address of each page in physical memory.

page offset (a) - combined with base address to define the physical memory address that is sent to the memory unit.

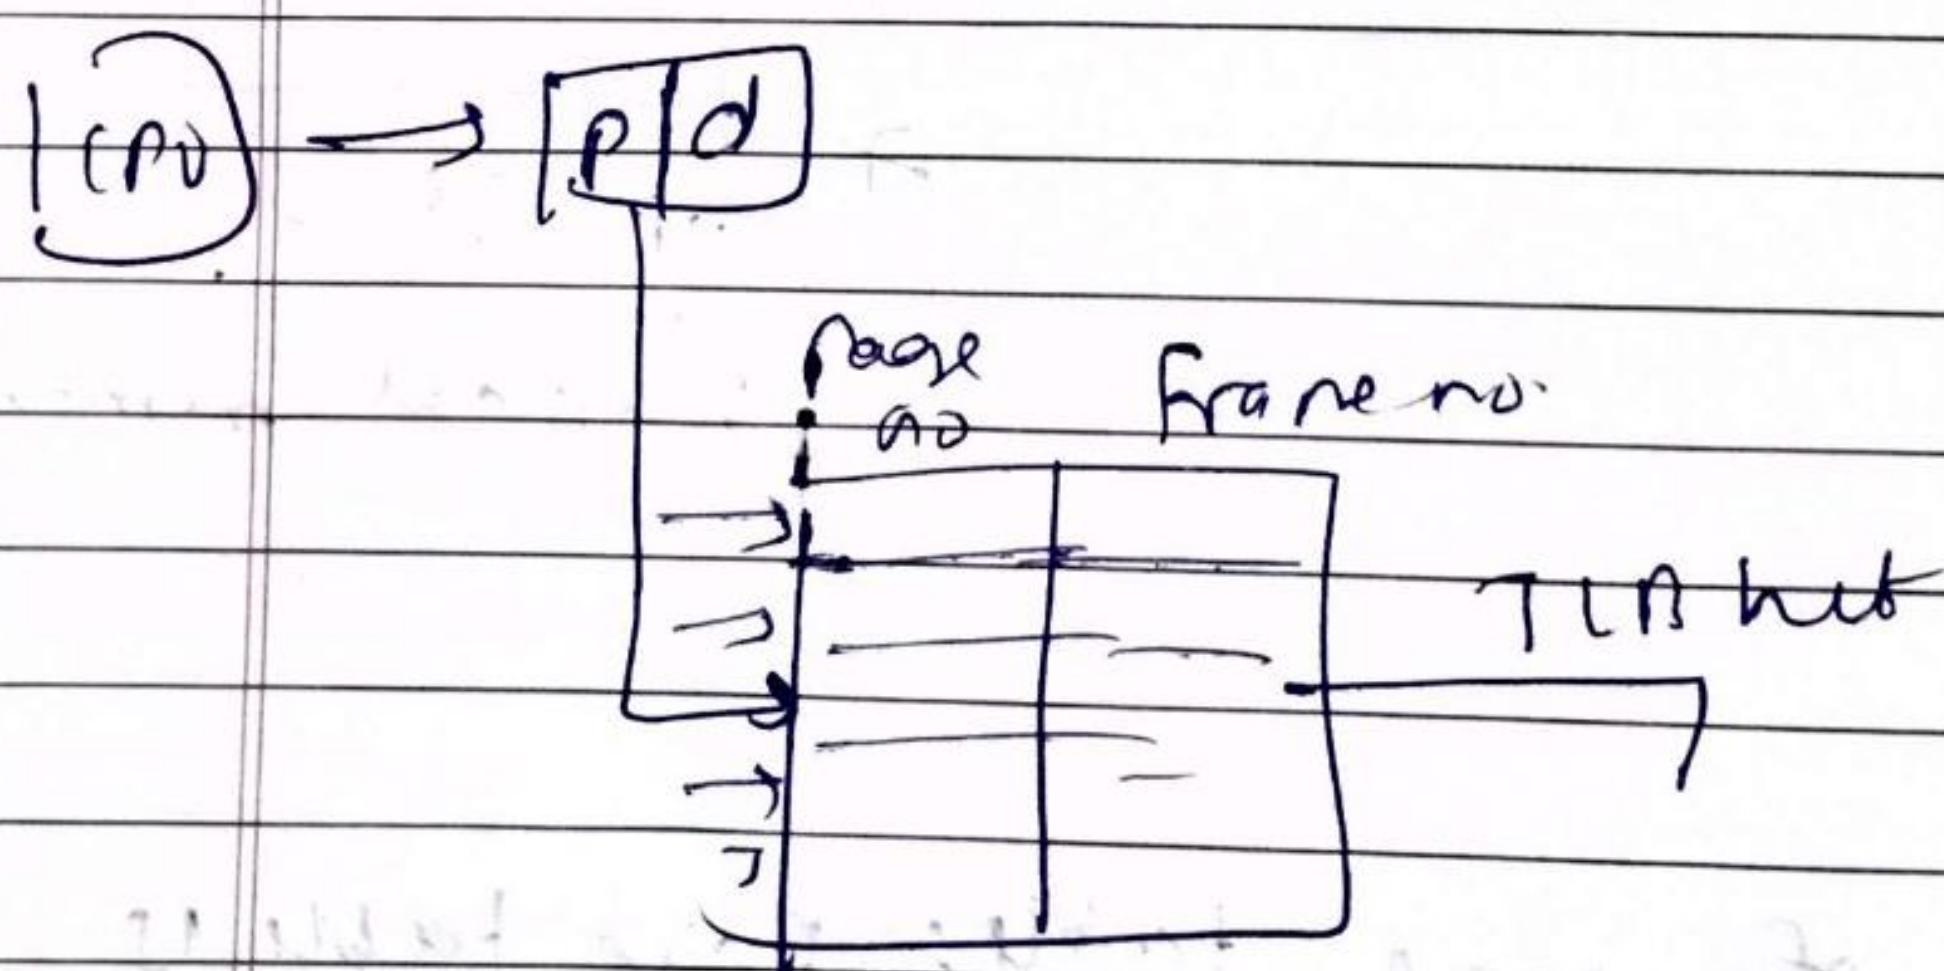


* free frames :-

- Implementation of page table:- Page table is kept in main memory.
Page table base register (PTBR) points to the page table.
Page table length register (PTLR) indicates size of the page table.
In this scheme every data/instruction access requires 2 memory accesses. one for the page table and one for the data/instruction.



paging hardware with TLB



O page table is kept in main memory
 page table base register (PTBR) points to the
 page table
 page table

This ↑ the time of access (as size of page table
 grows)

This problem can be solved by the use of
 a special lookup hardware called as
 TLB

hit ratio - percentage of times that a
 page number is found in TLB.

^(up) effective memory access time is reduced by
 TLB.

memory protection : memory protection implemented
 by associating protection bit with each frames
 valid invalid bit ~~attached~~ attached to each entry.

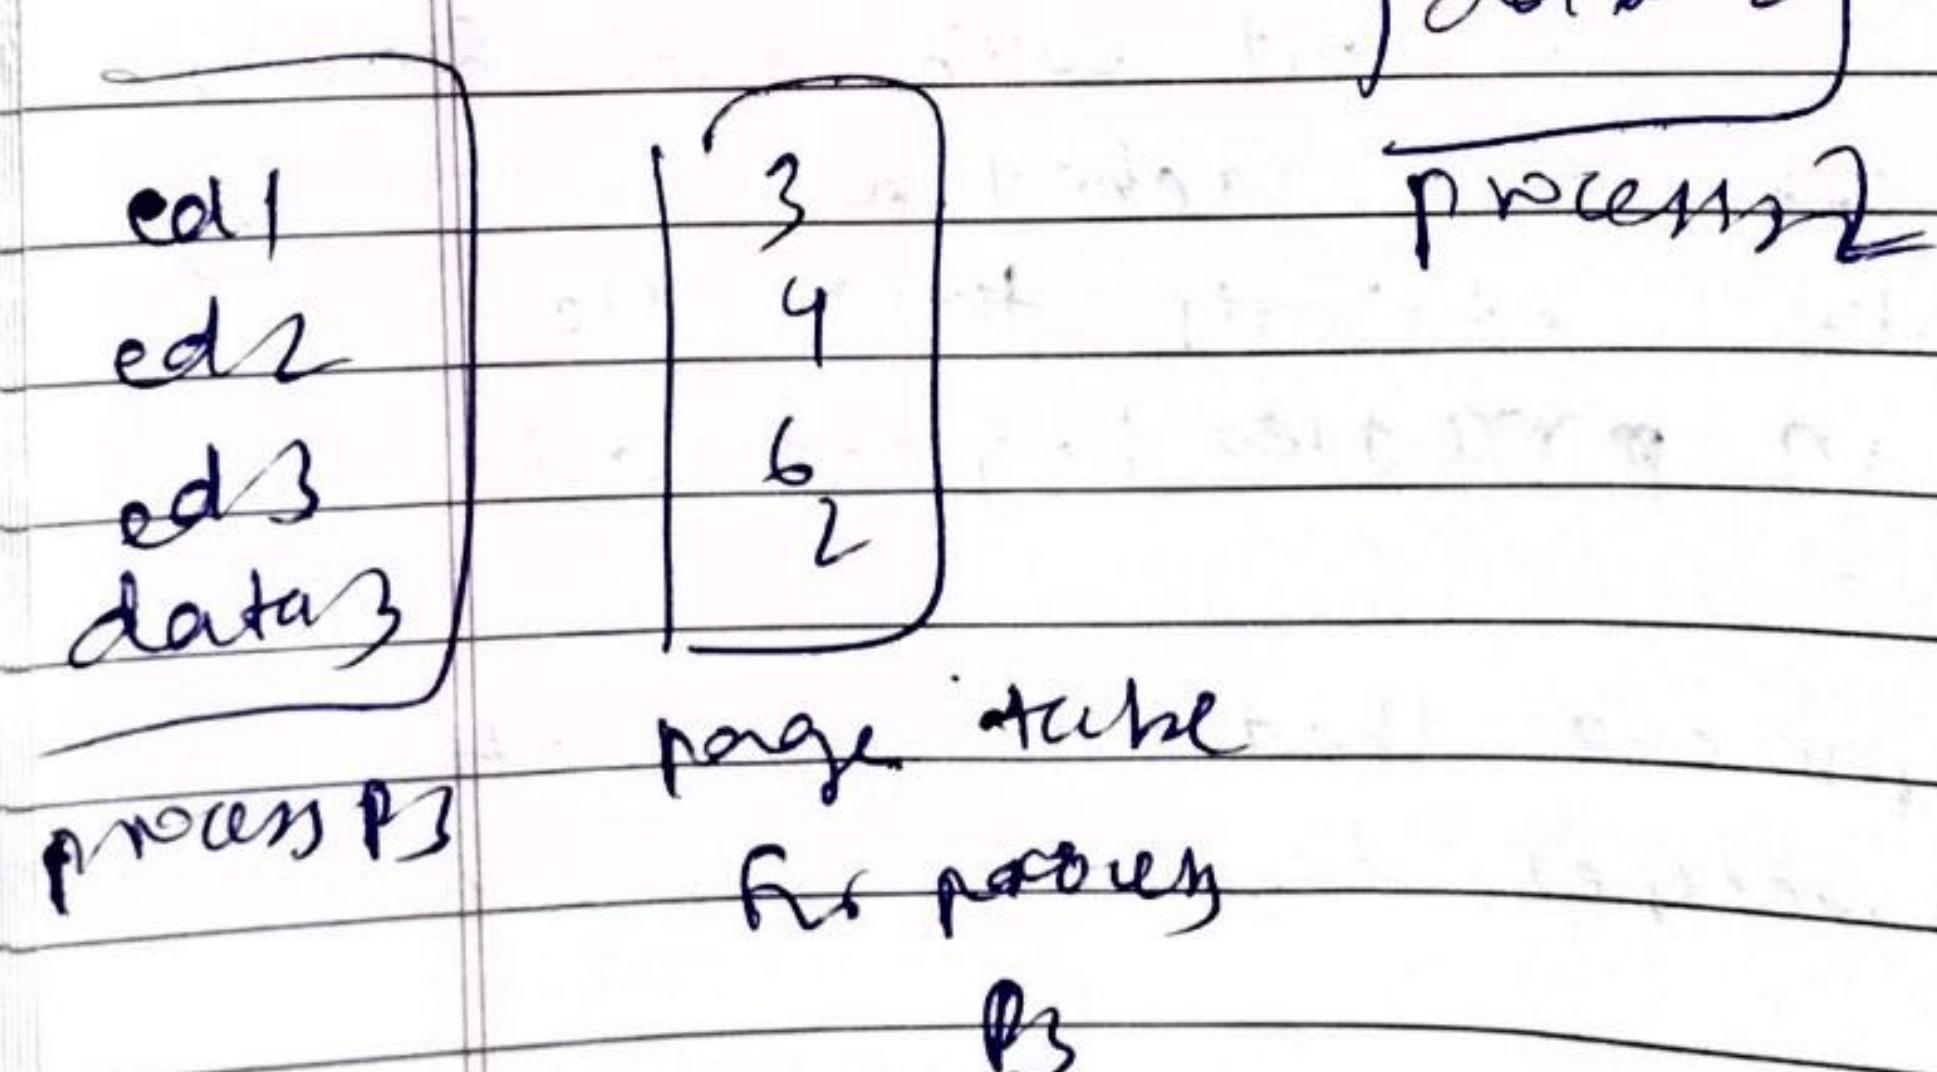
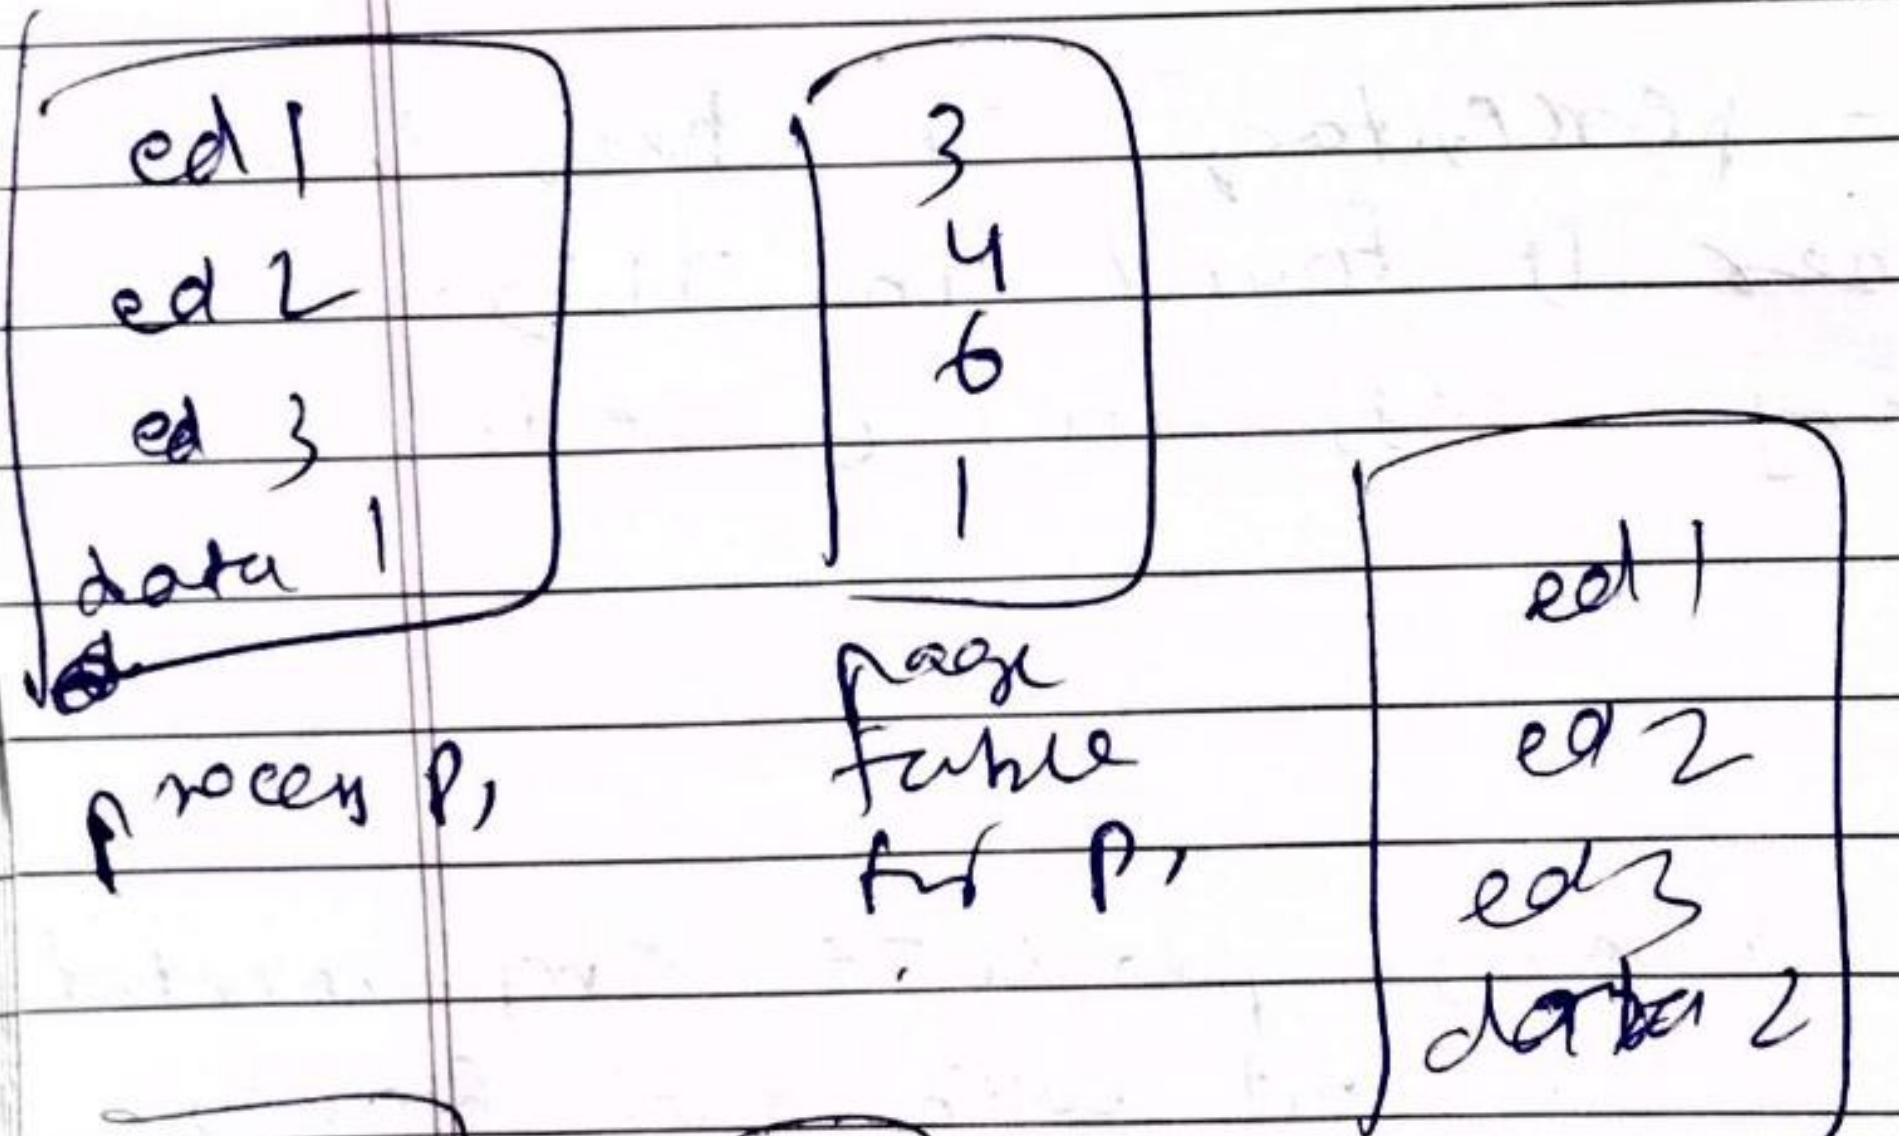
In page table: "valid" indicates that the
 associated page is in processes logical address
 space

"invalid" also is present that no associated
 page is not in processes logical address
 space.

* shared pages: - shared code

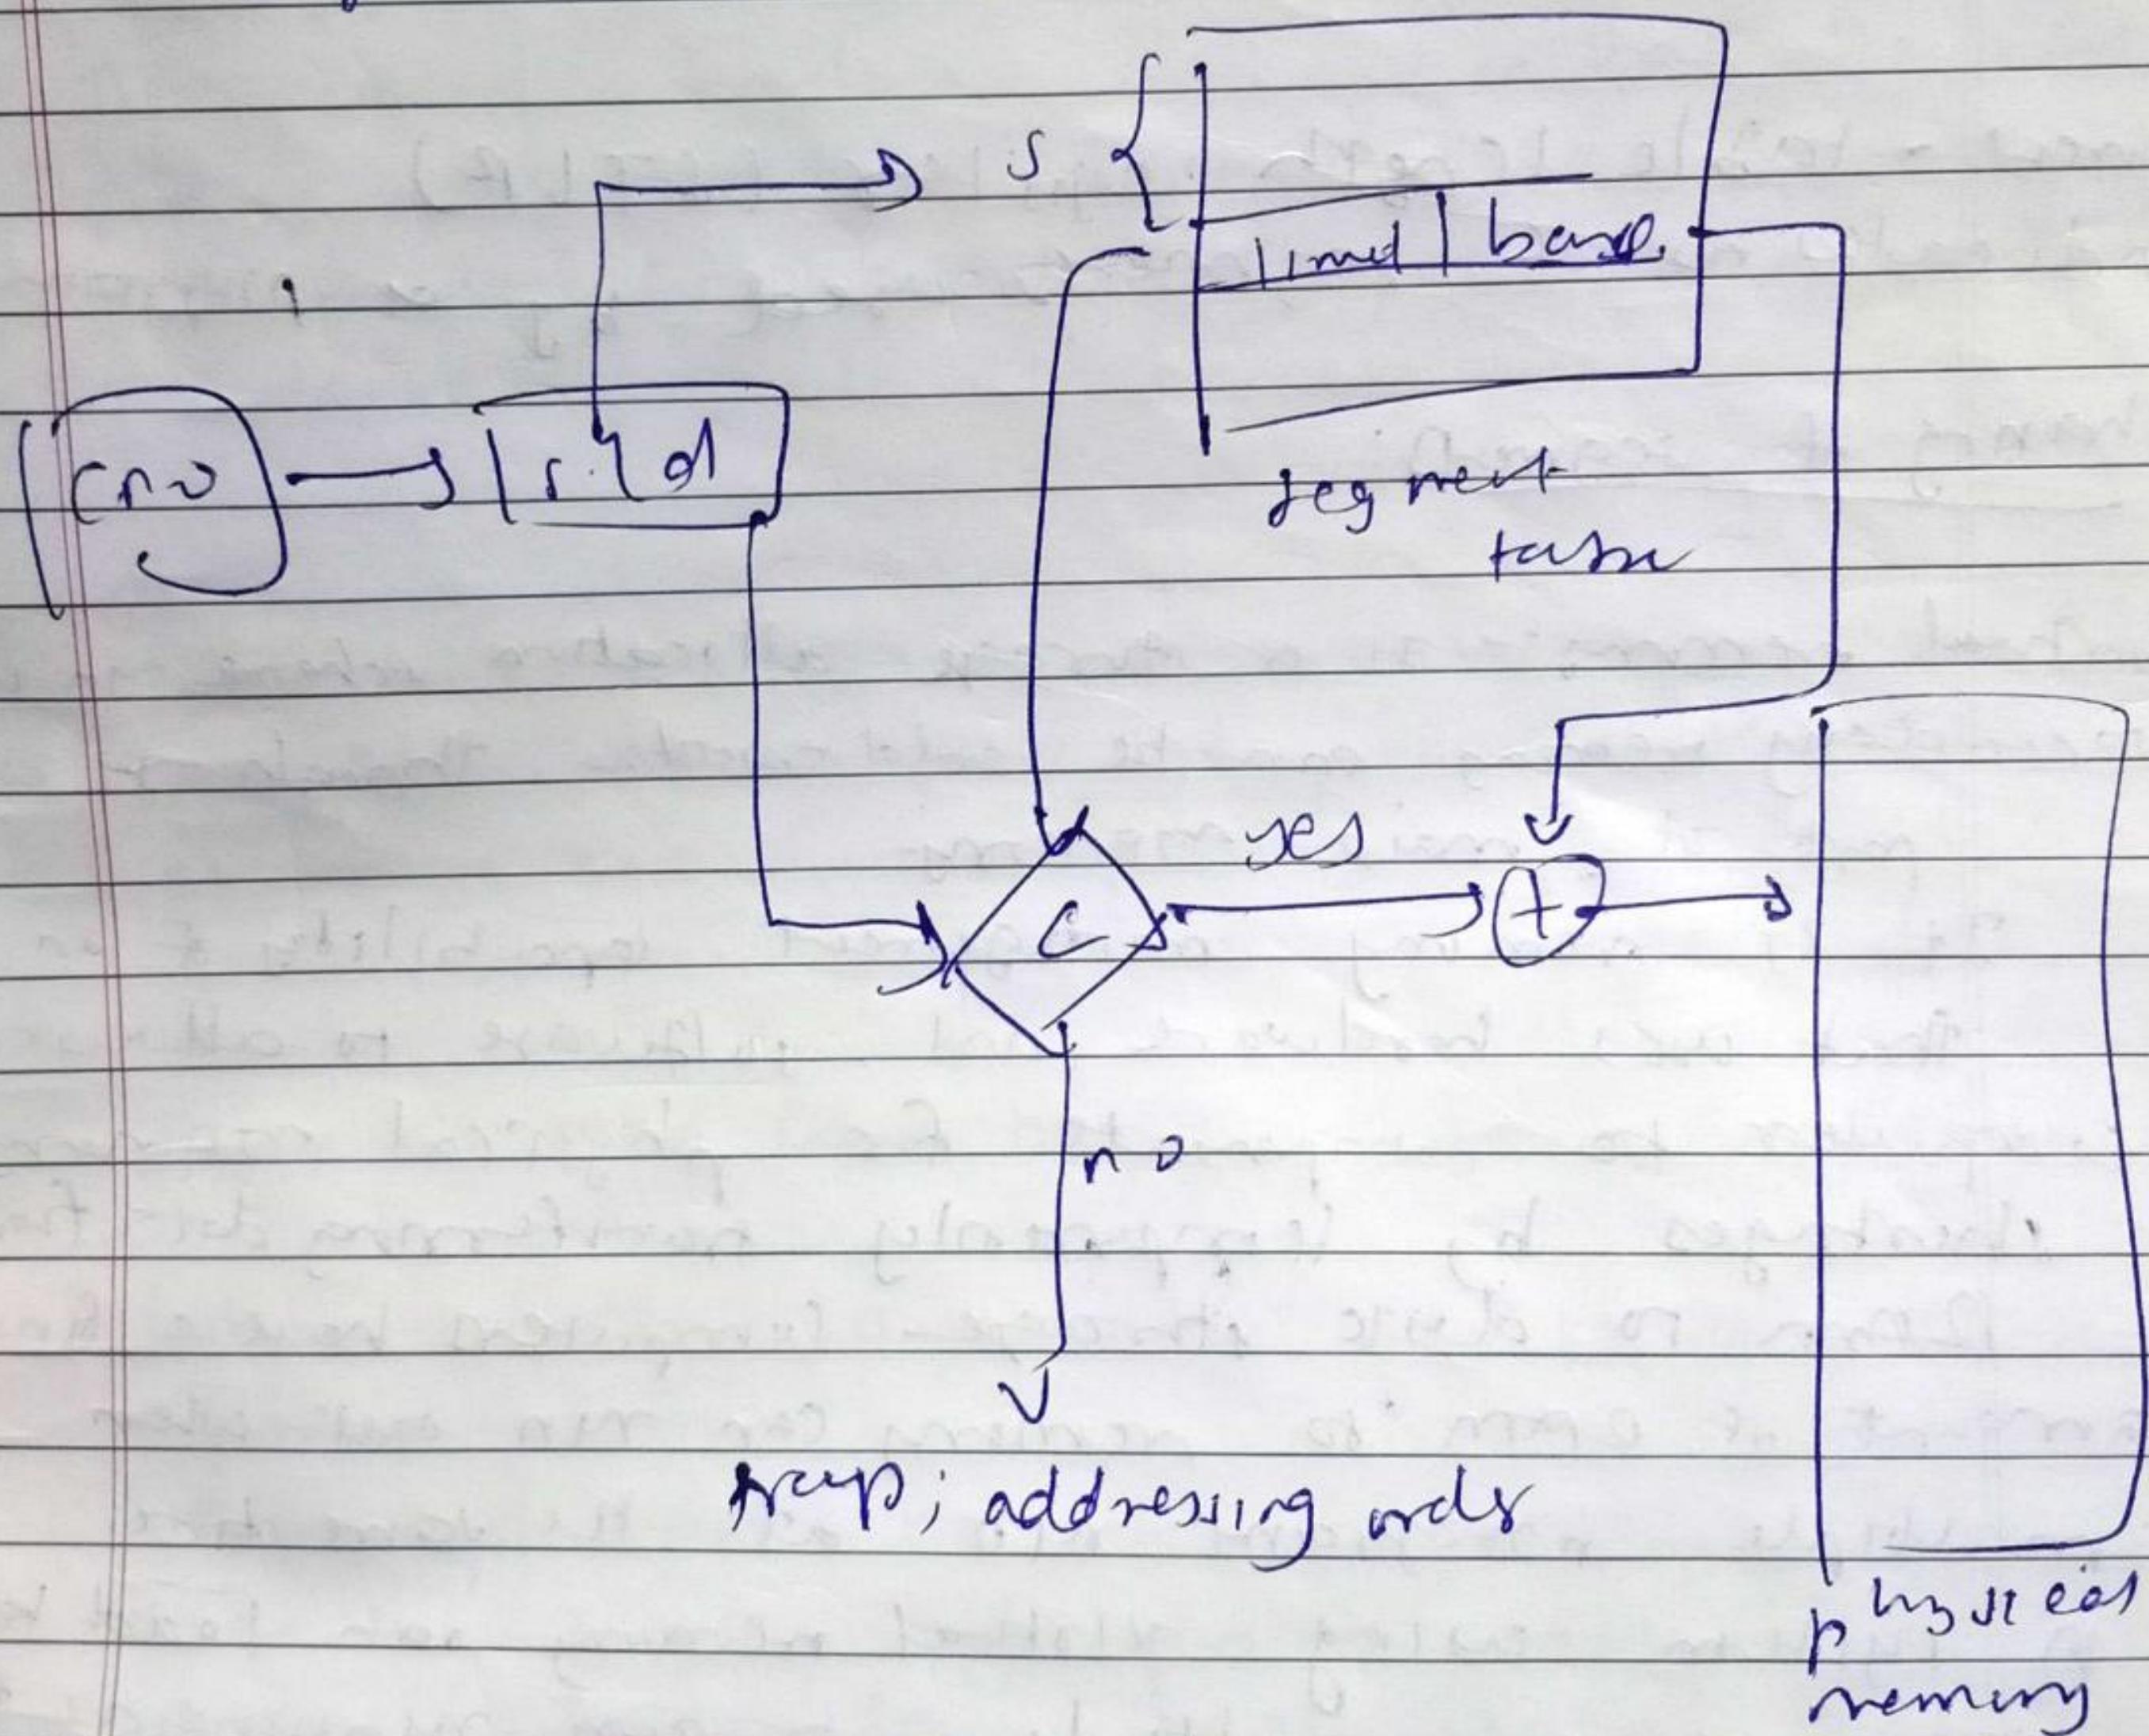
One copy of shared code among processes
 shared code must appear in same location in
 the logical address space of all processes
 private code and data.

- Each process keeps a separate copy of the code and data
- file pages for private code and data can appear anywhere in logical address space



• Segmentation: A memory management technique in which memory is divided into variable sized chunks which can be allocated to processes. Each chunk is called a segment.

Segmentation hardware



Segmentation architecture

logical address consists of 2 tuple
(segment-number, offset)

Segment table - maps 2 dimensional physical addresses over each table entry has

base - contains the starting physical

address where the segments reside in memory.
limit - specifies the length of the segment

segment-table base register (STBR) -

points to segment table's location in memory.

segment-table length register (STLR)

indicates no. of segments used by a program

Sharing of segments:

'Virtual memory' is a storage allocation scheme in which secondary memory can be addressed as though it were part of main memory.

It is memory management capability of an OS that uses hardware and software to allow a computer to compensate for physical memory shortages by temporarily transferring data from RAM to disk storage. Computers have a finite amount of RAM so memory can run out when multiple programs run at the same time.

A system using virtual memory can load larger programs as multiple programs running at the same time, allowing each one to operate as if it has infinite memory and without having to purchase more RAM.

The OS divides memory into page files or swap files.

Each page is stored on a disk and when the page is needed, the OS copies it from the disk to main memory and translates virtual

addresses into real addresses.

Page fault :- what happens if a process tries to access page that was not brought into memory? Access to a page marked invalid causes a page fault.

The paging hardware will notice that the invalid bit is set, causing a trap to the OS.

This trap is the result of the OS failing to bring the desired page to the memory.

Steps

Trap to OS

Save the user registers and process state

Determine that interrupt was a page fault

Determine the location of the page on disk

Issue a read from the disk to a free frame

while waiting, allocate the CPU to some other process.

Interrupt from the disk (I/O completed)

Save the registers and process state

Correct page table table to show that the page is now in memory

Wait for the CPU to be allocated to the process again.

Restore the user register, process state and new page table and resume the interrupted instruction.

what happens if there is no free frame?

- page replacement - find some page in memory, but not really in use, swap it out

- algorithm

- performance - want an algorithm which will result in min. no. of page faults

* Basic page replacement - find the location of the desired page on disk and a free frame

- If there is a free frame, use it

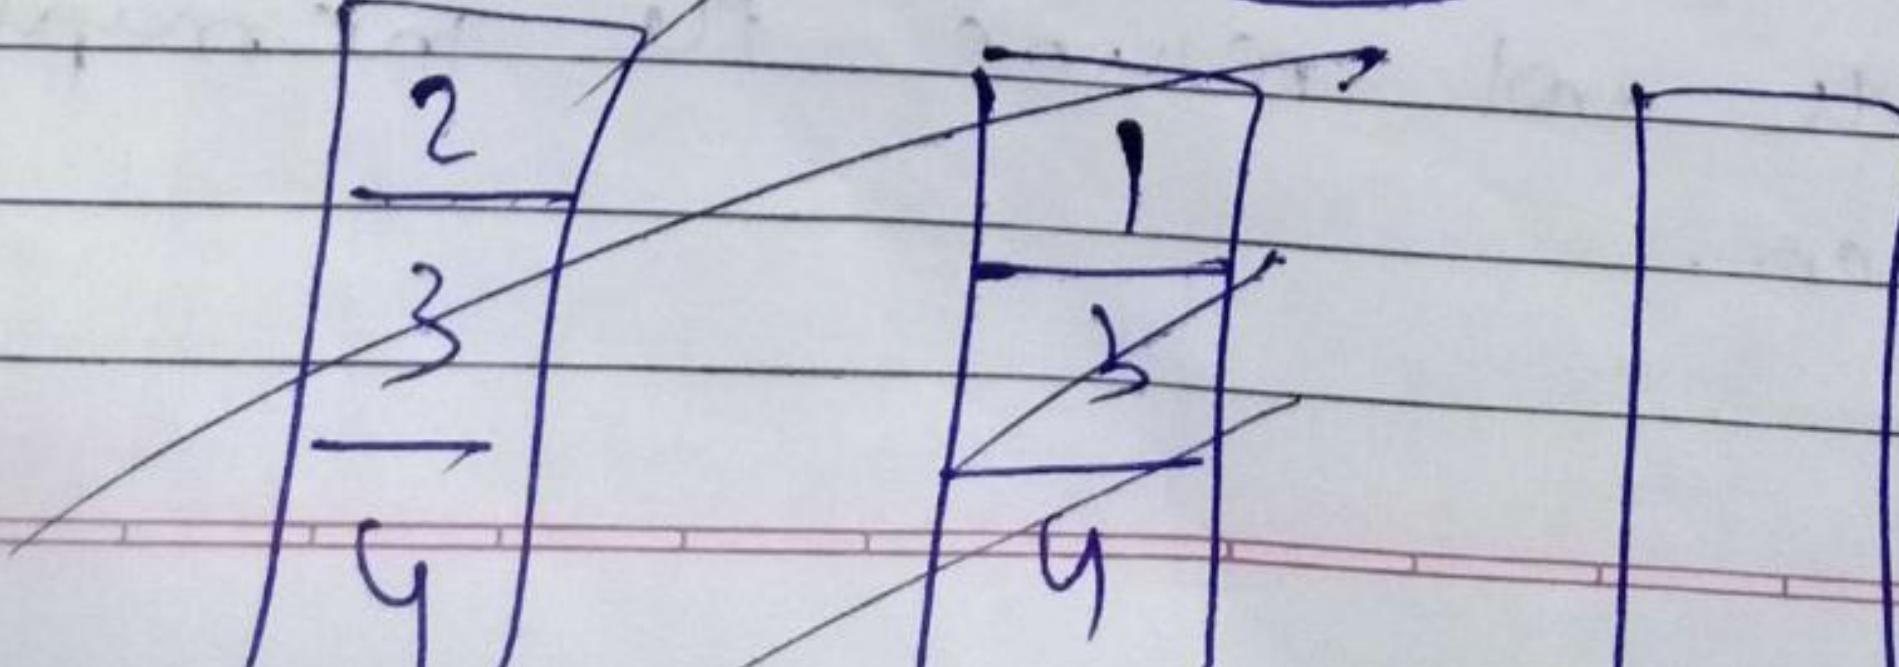
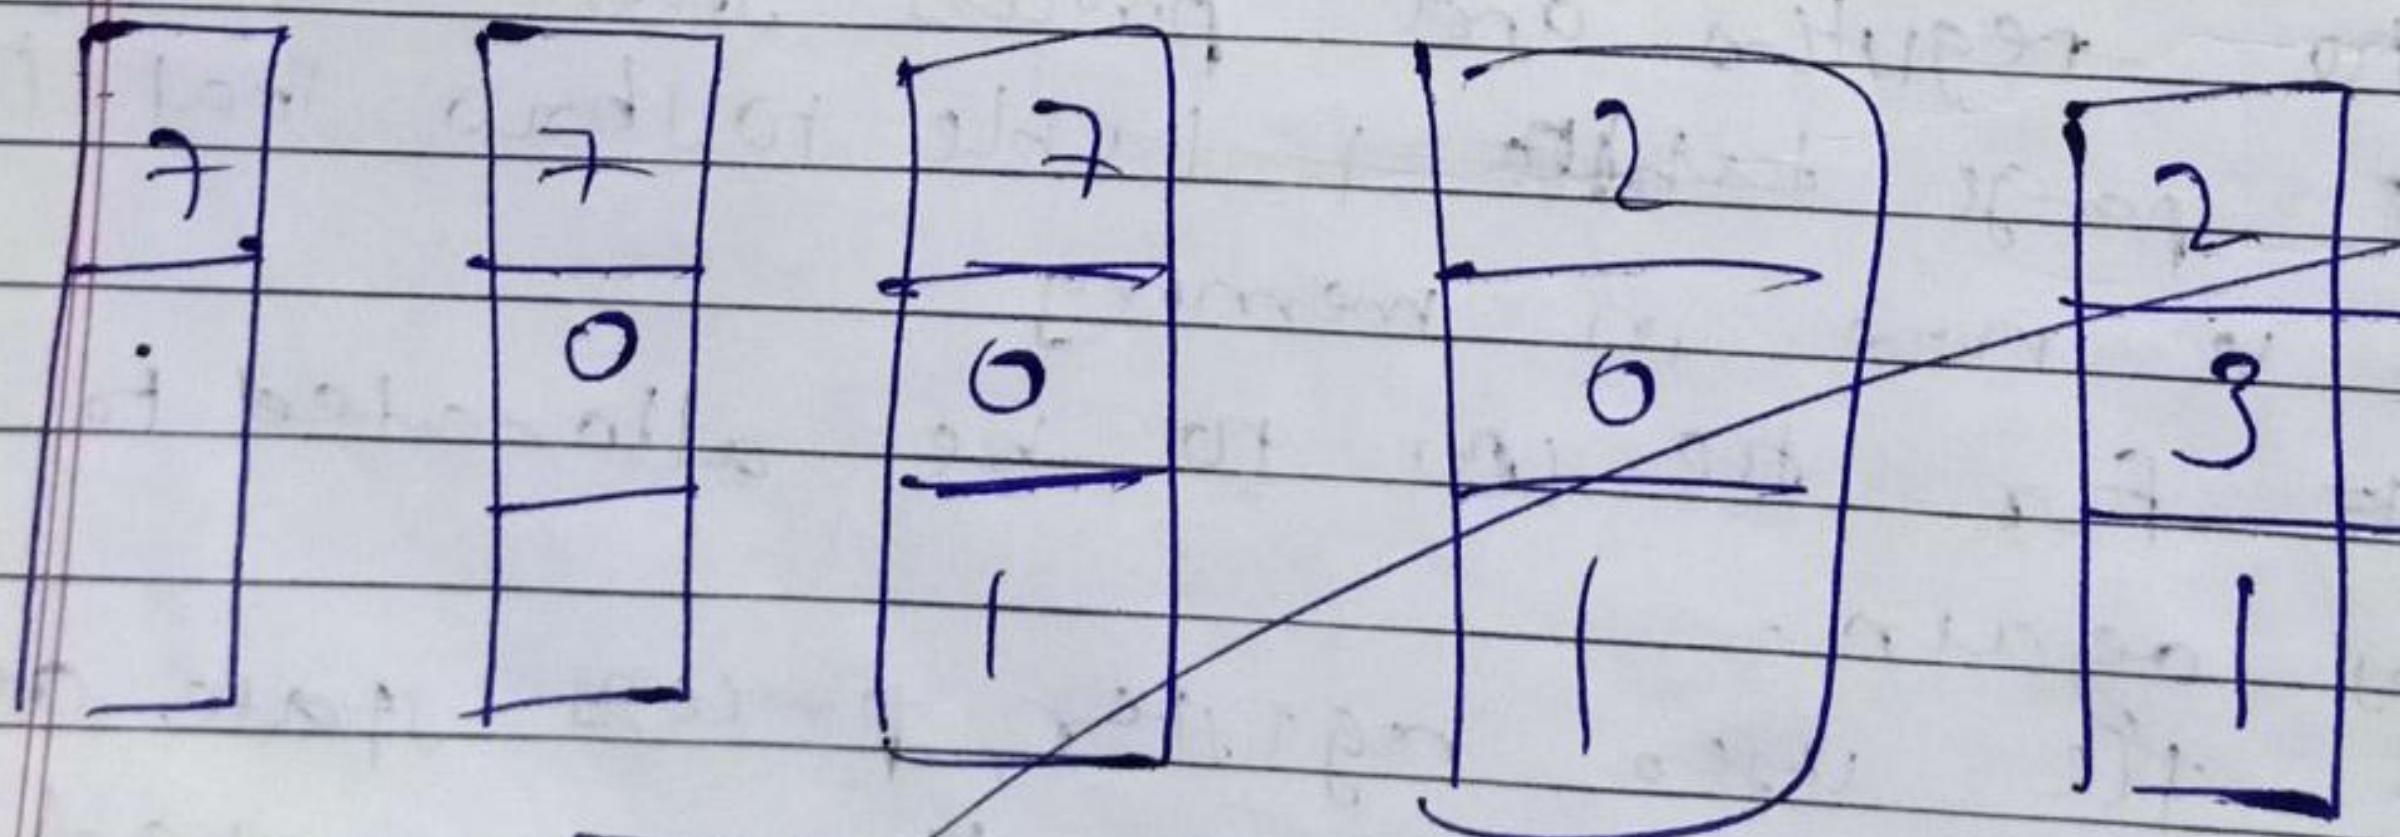
- If there is no free frame, use a page replacement algorithm to select a victim frame.

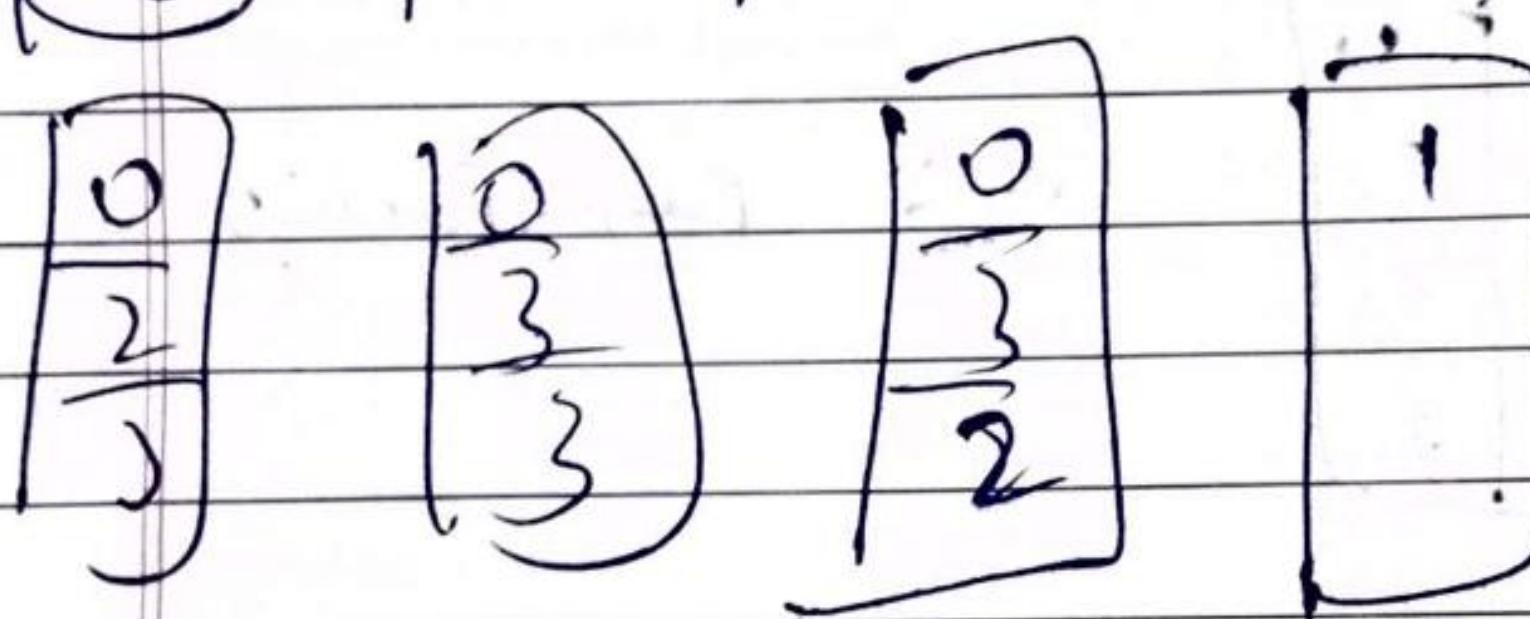
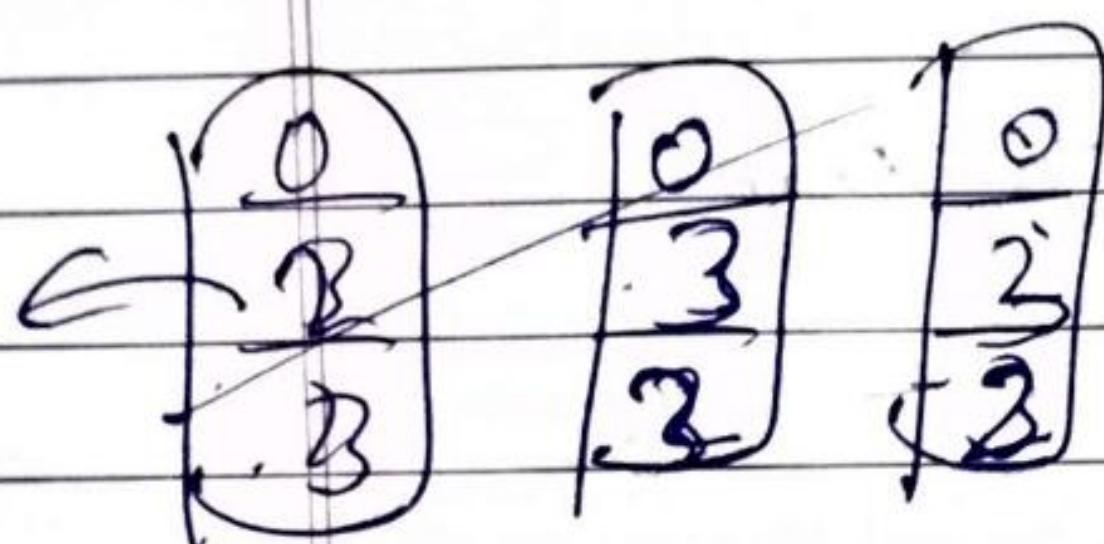
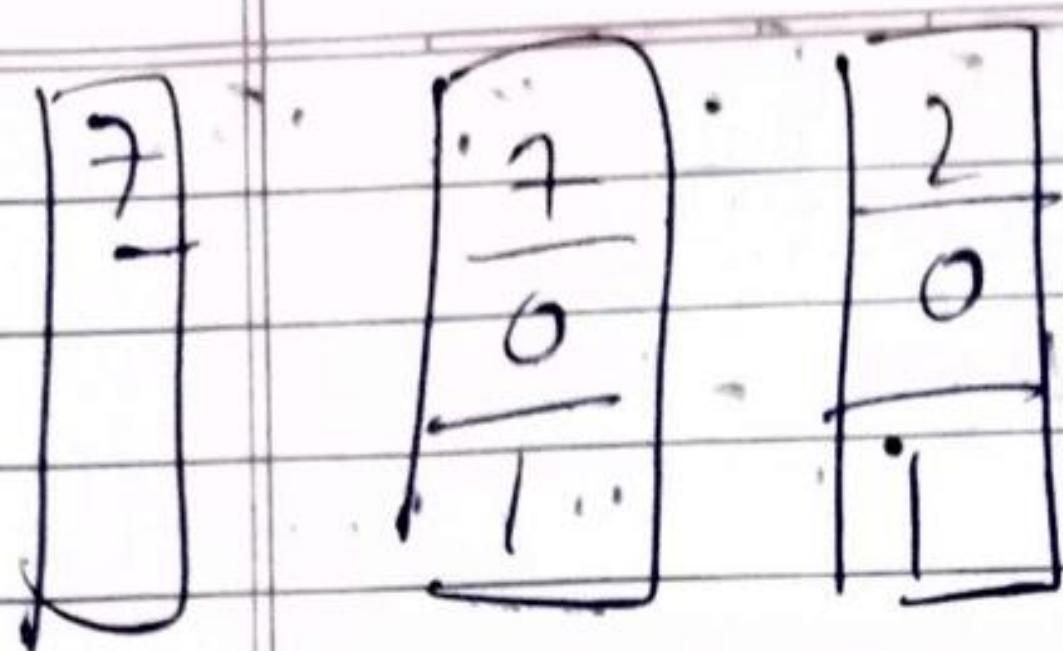
3) read the desired page into the (newly) free frame. Update the page table.

a) restart the process

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

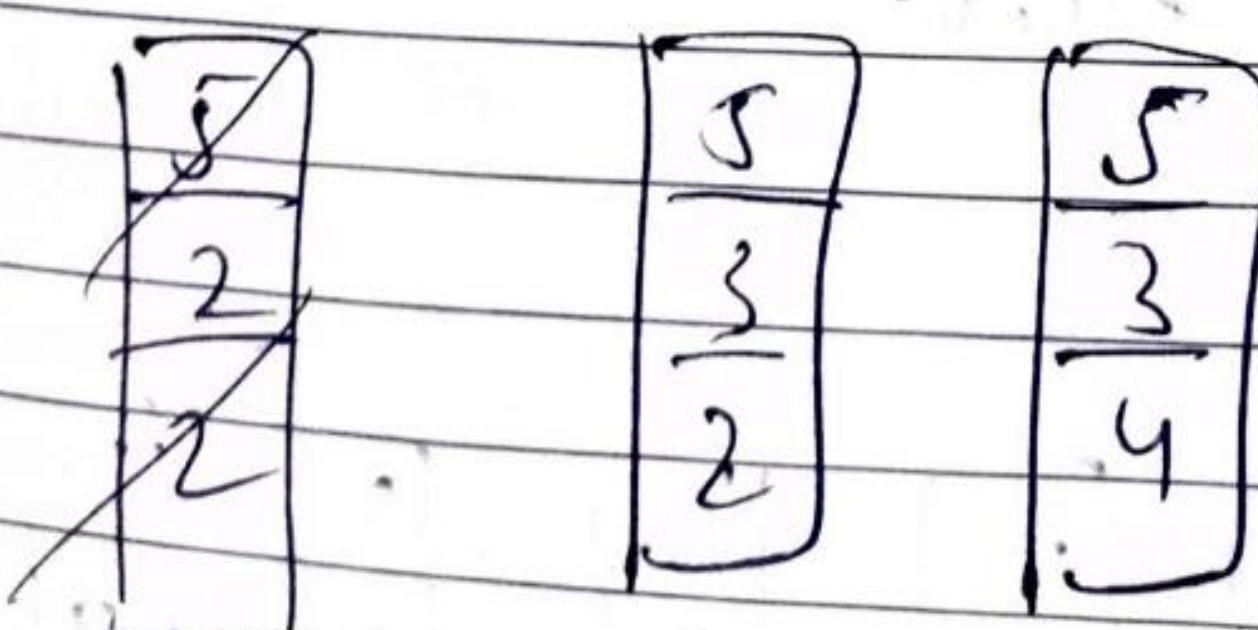
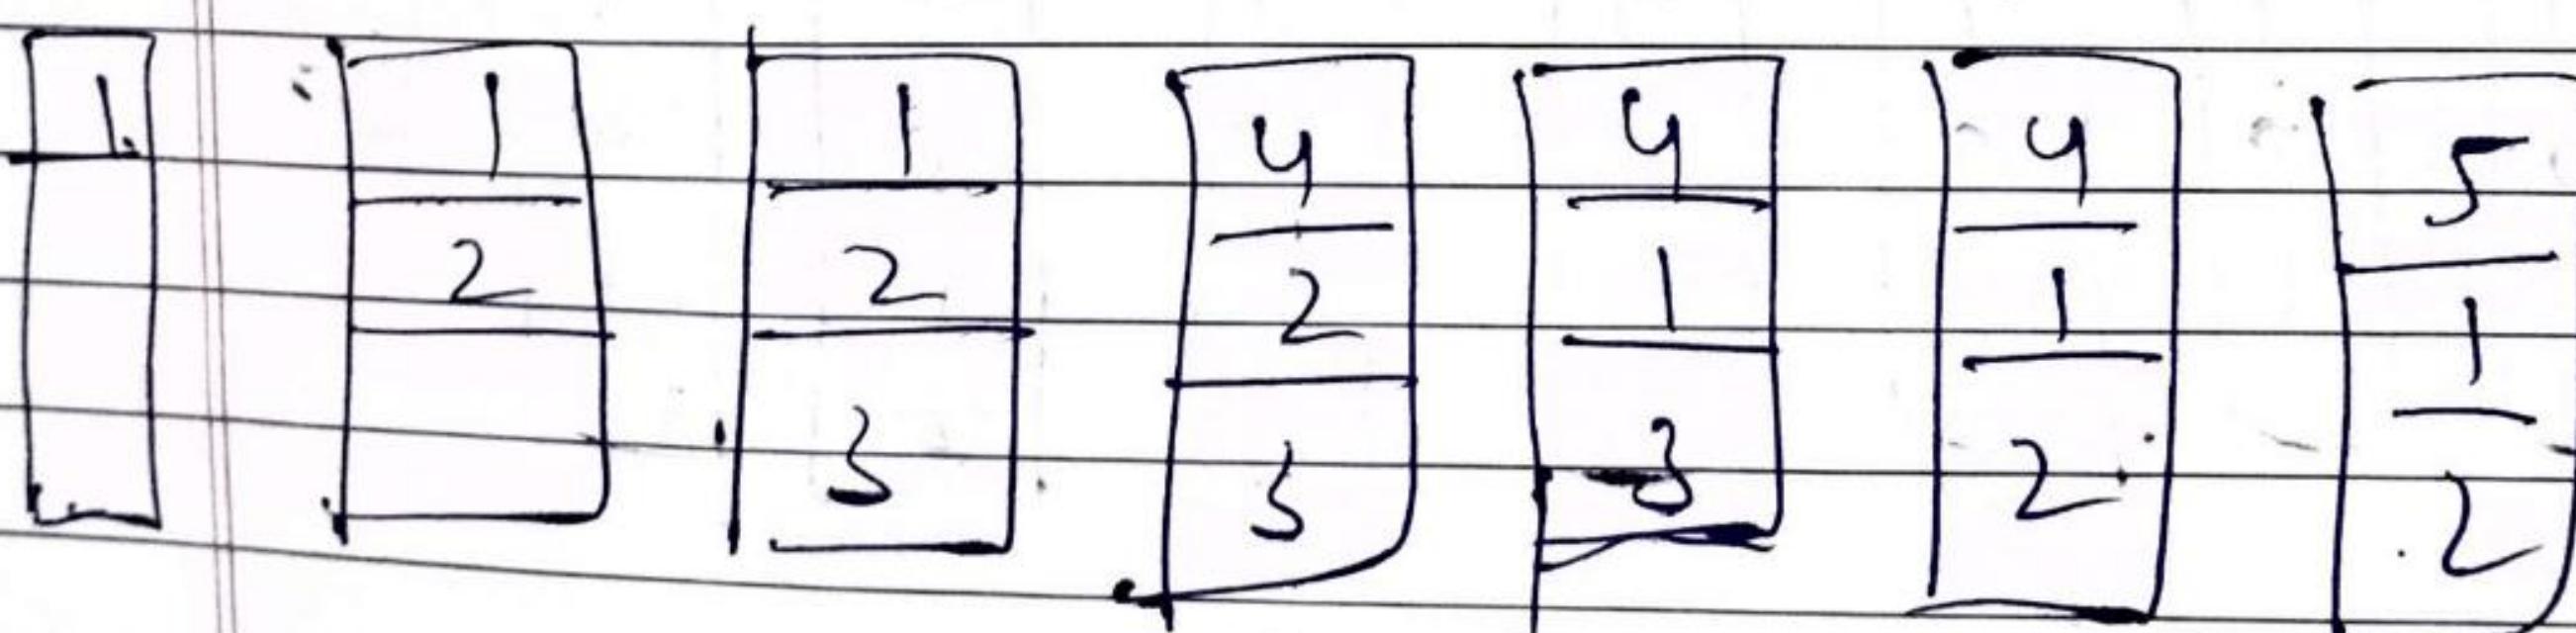




* FIFO

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

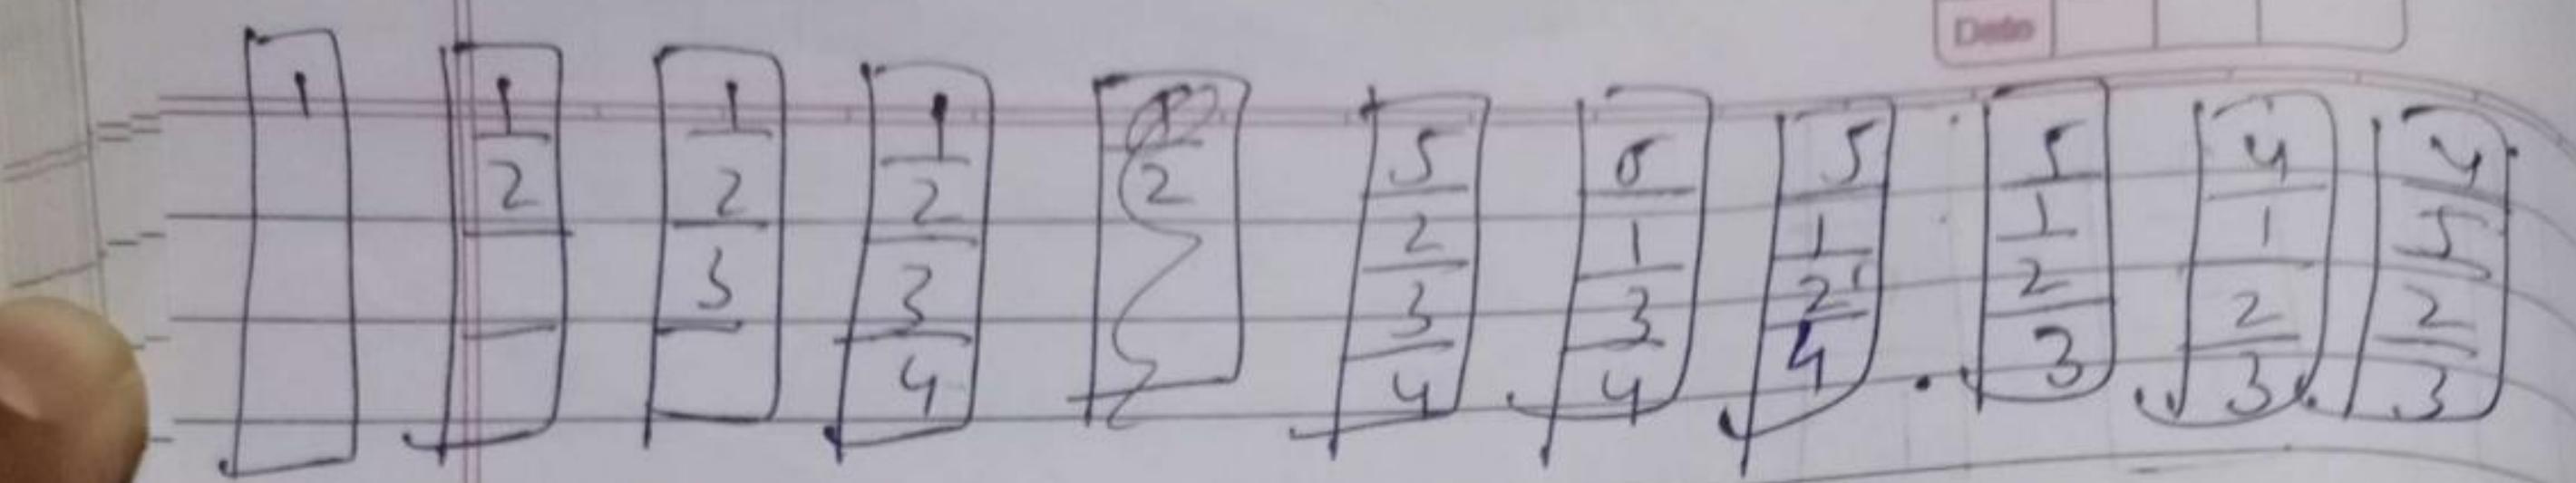
3 frames



≈ 9 page faults

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 1

Page No.		
Date		

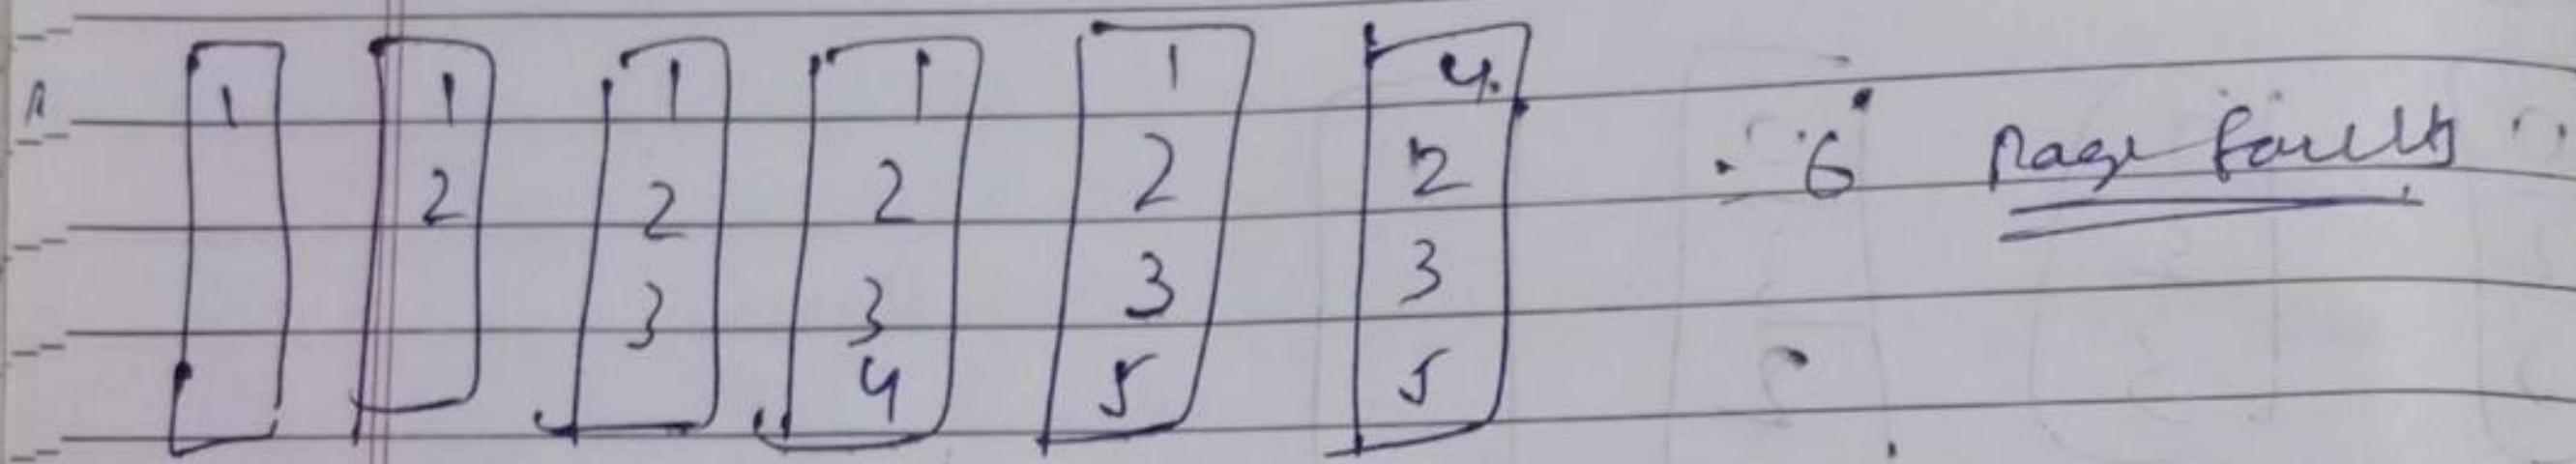


= 10

This is Belady's anomaly. Page fault increases as frames ↑

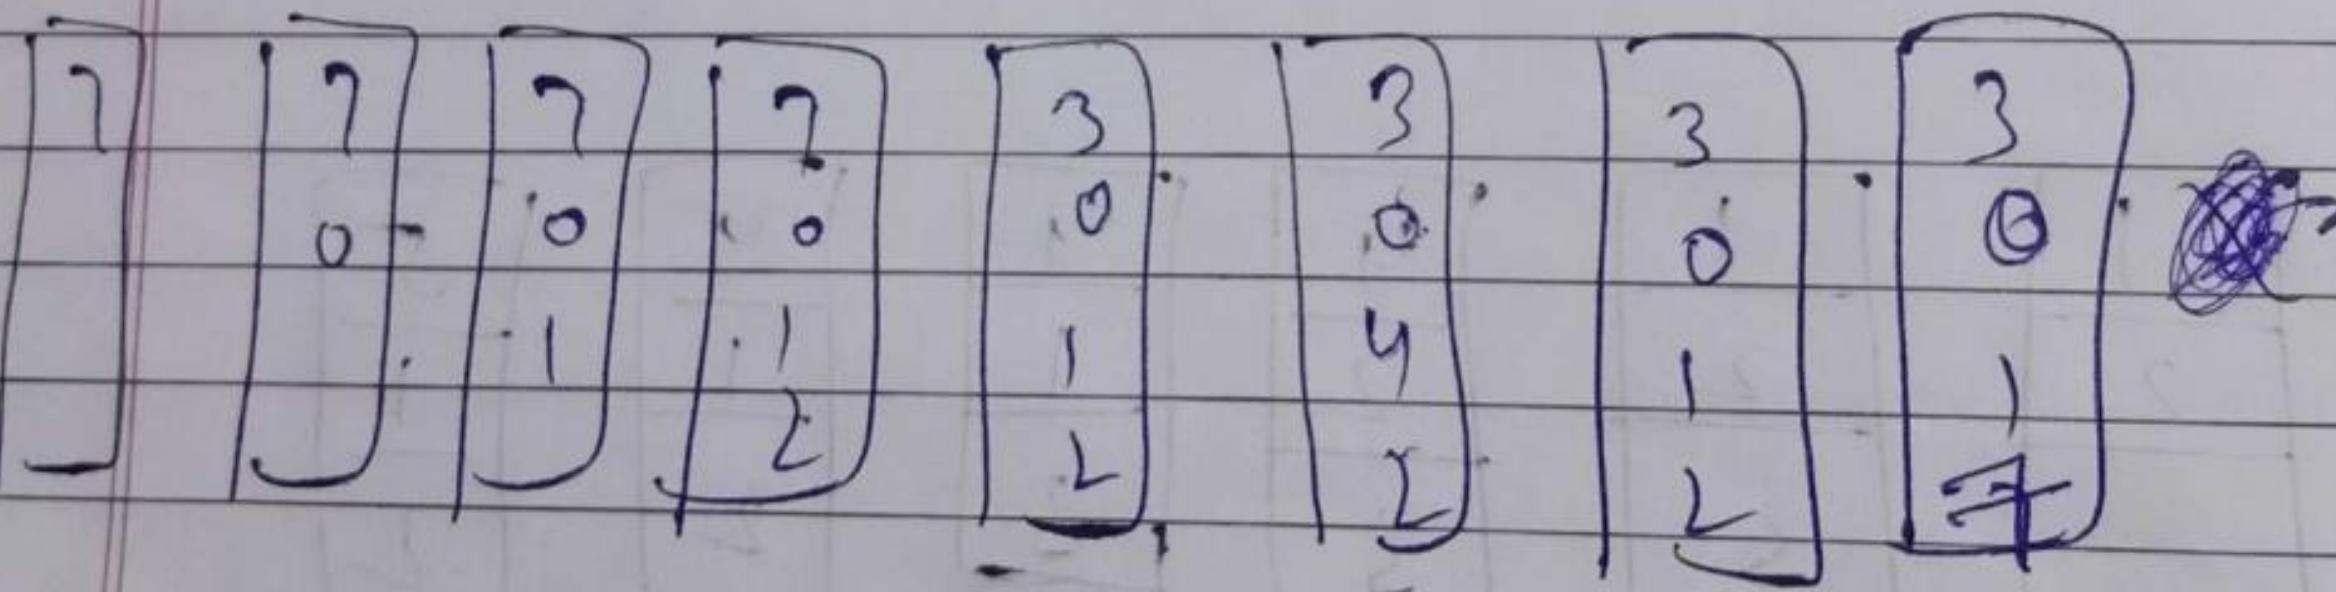
2) Optimal page replacement algorithm:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5, 2, 3



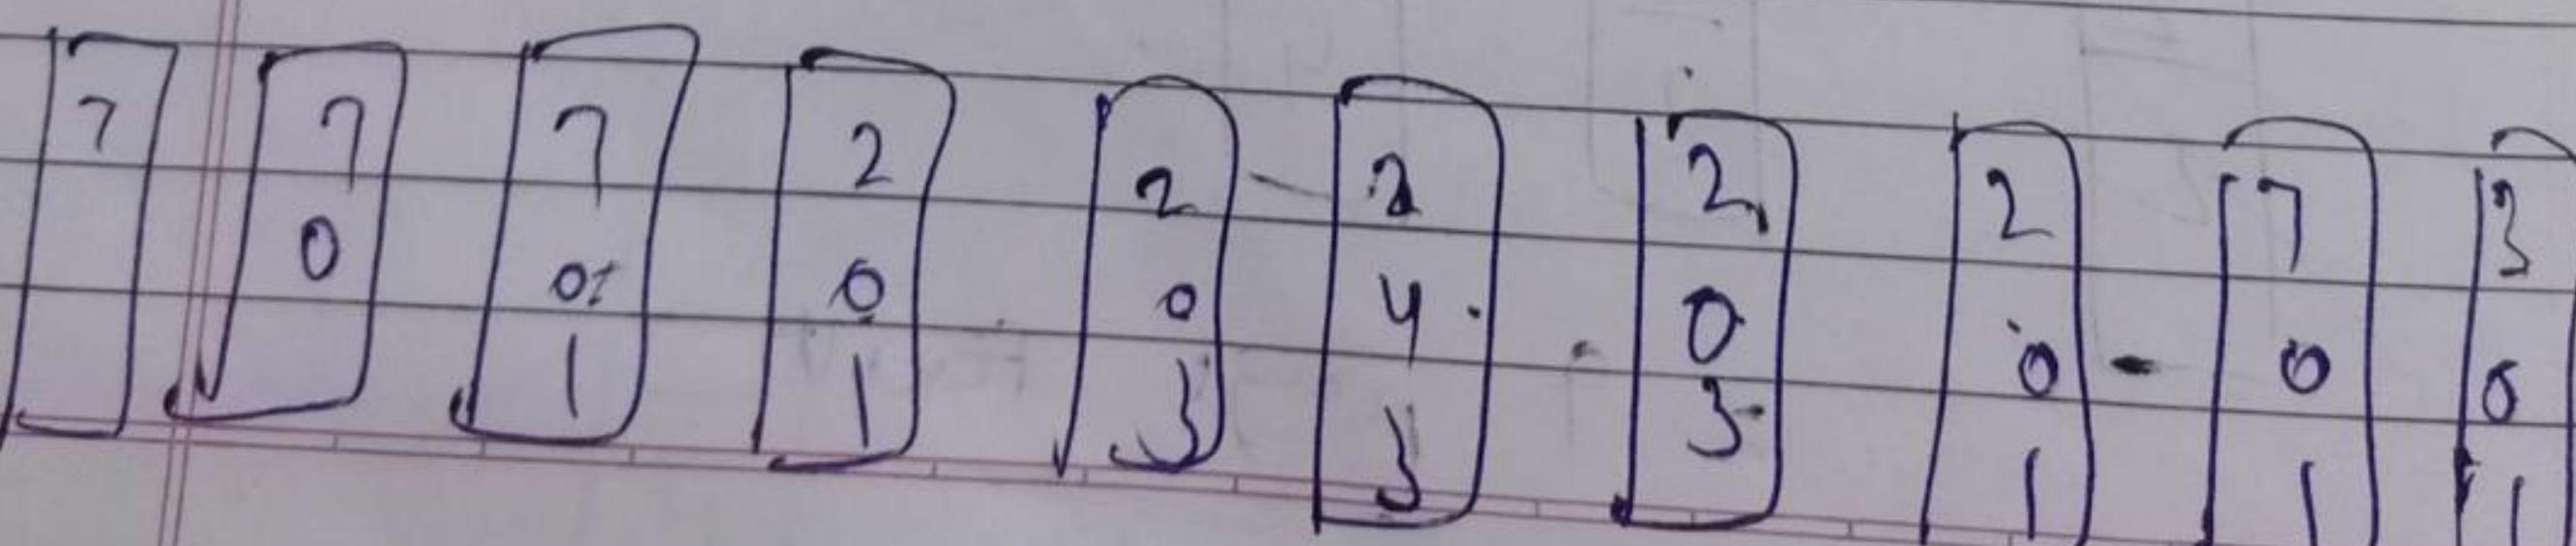
= 6 page faults

7, 0, 1, 2, 3, 0, 4, 1, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0
1, 3, 0, 1



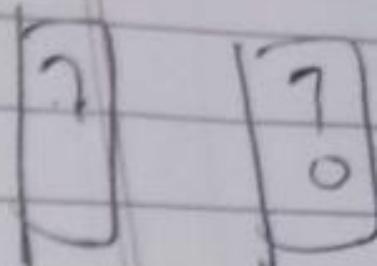
= 8 page faults

3 frames



= 10 page faults

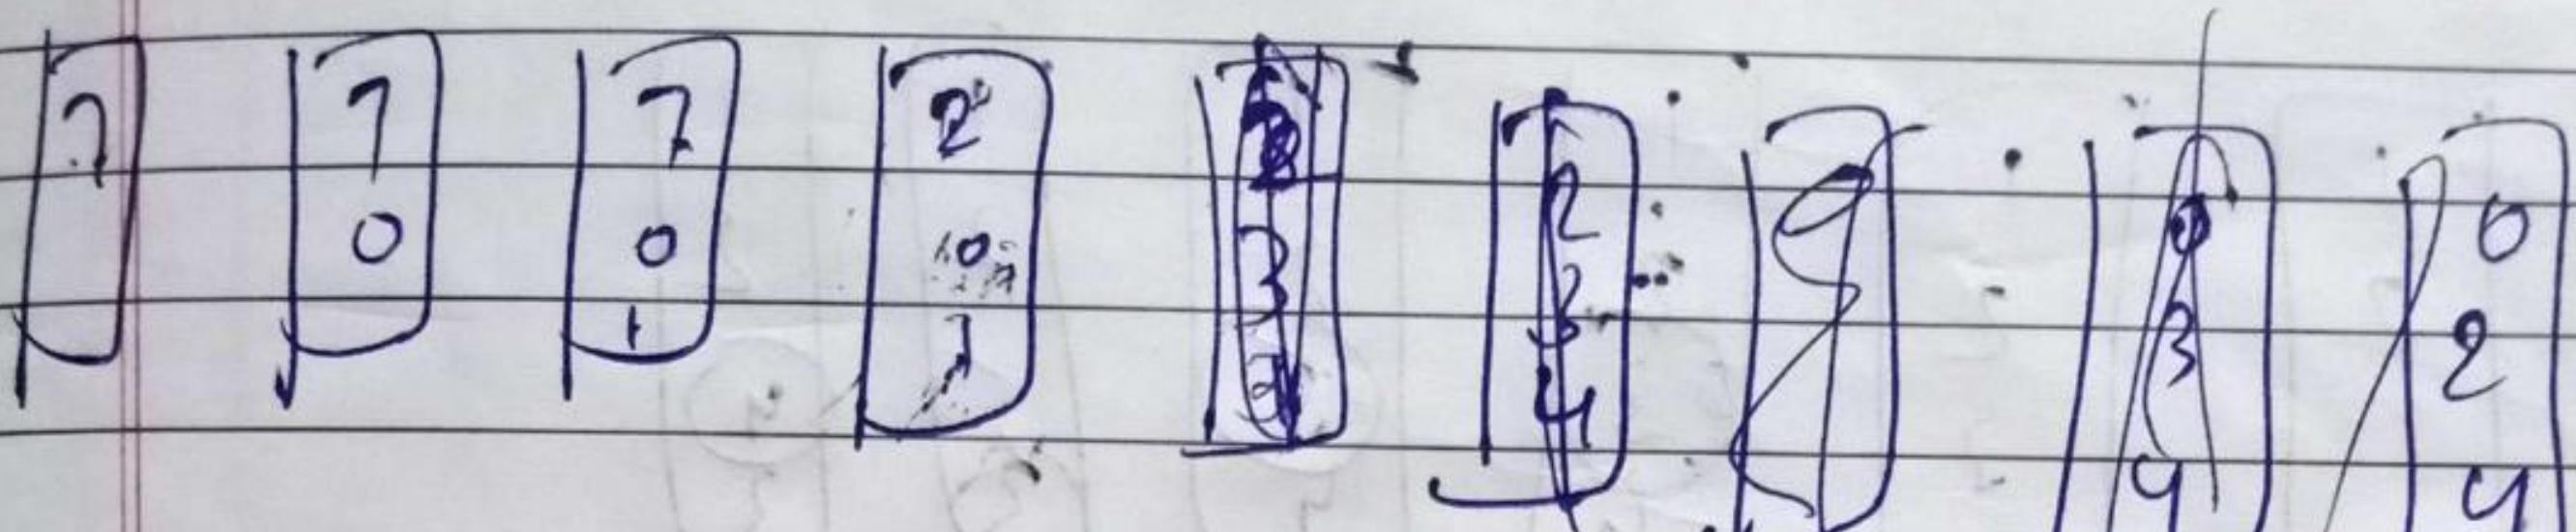
3) LRU (least recently used)
7, 0, 1, 1, 3 frames



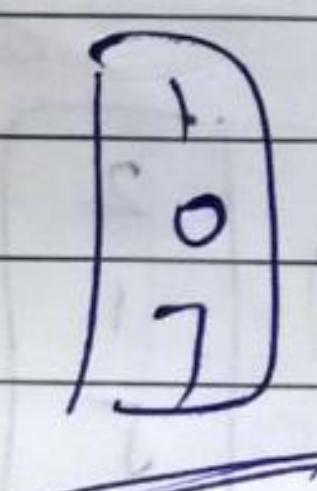
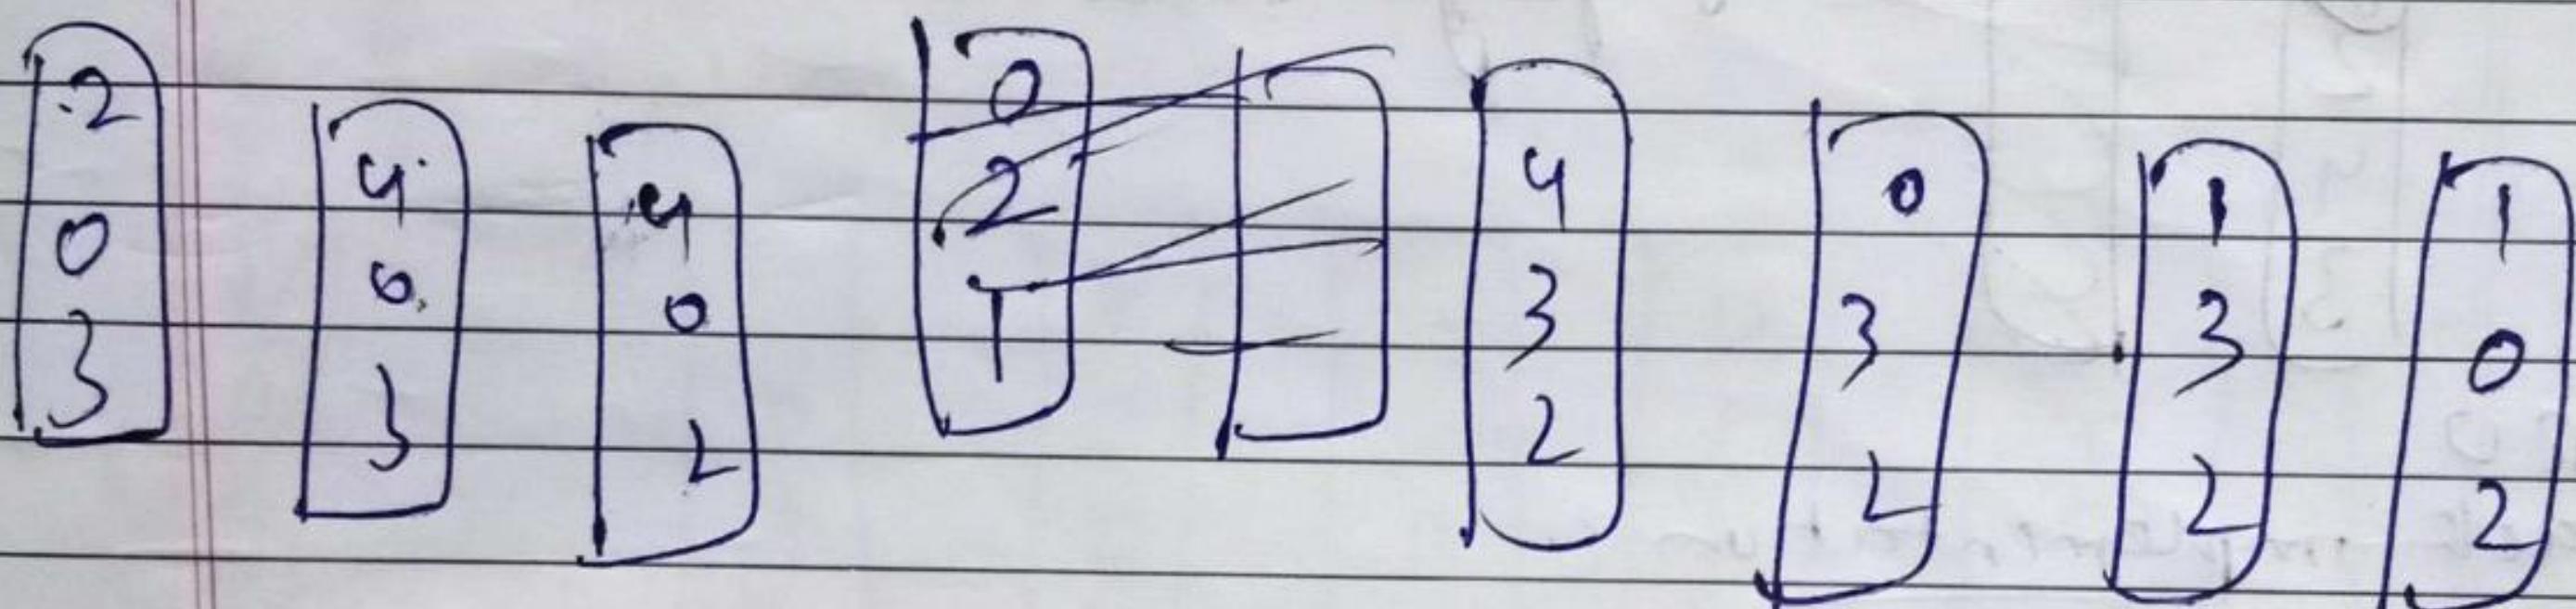
3) LRU (least recently used)

7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 7, 2, 1, 2,
0, 1, 7, 0, 1

3 frames

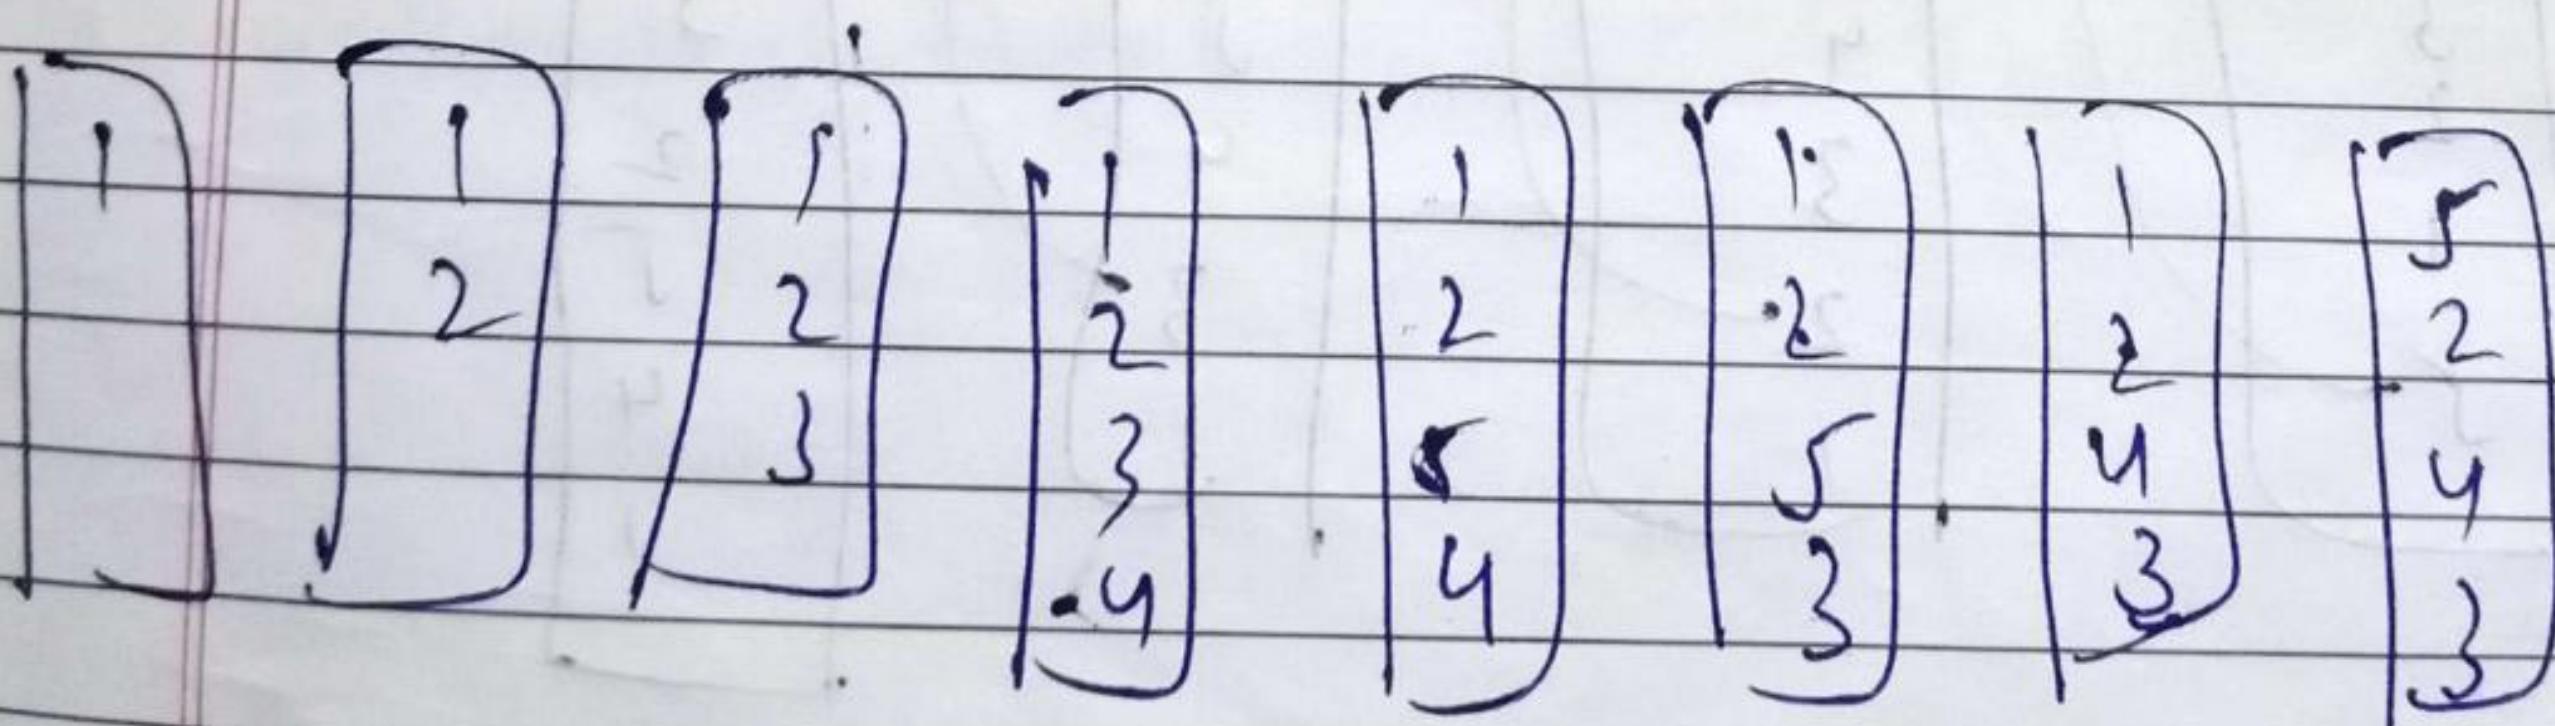


LRU page faults.



= 12 page faults

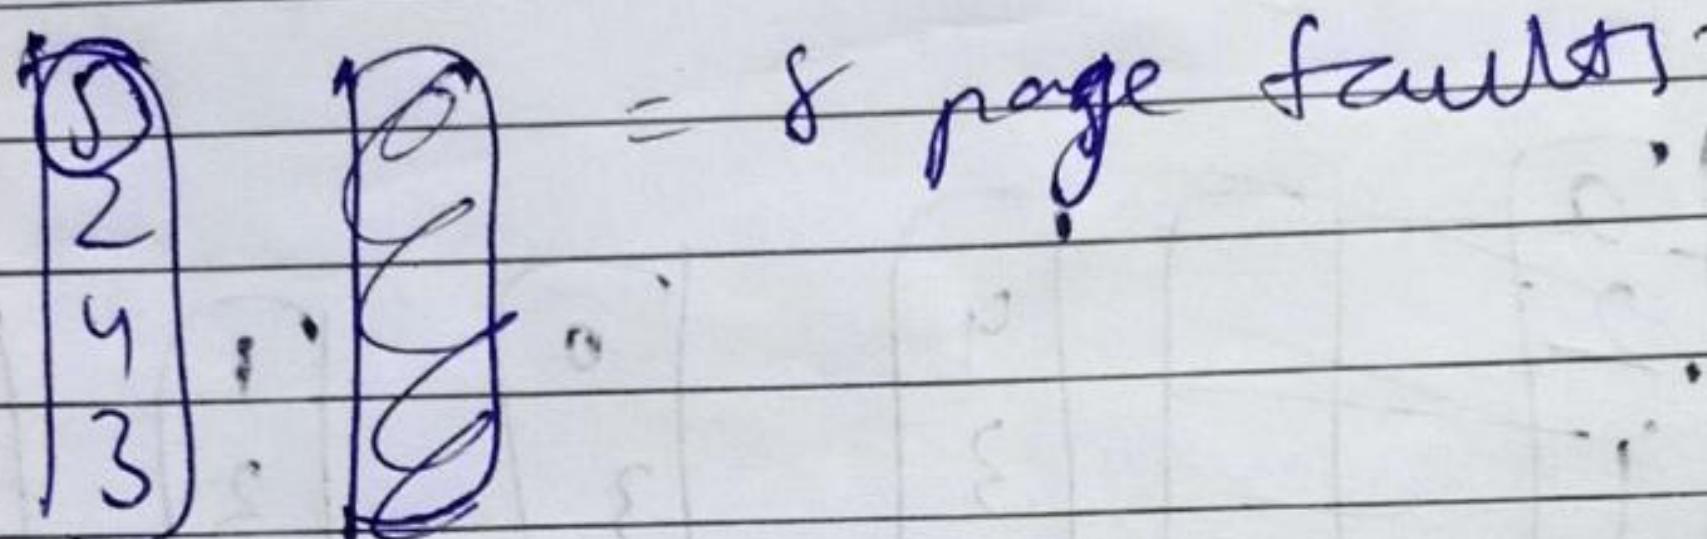
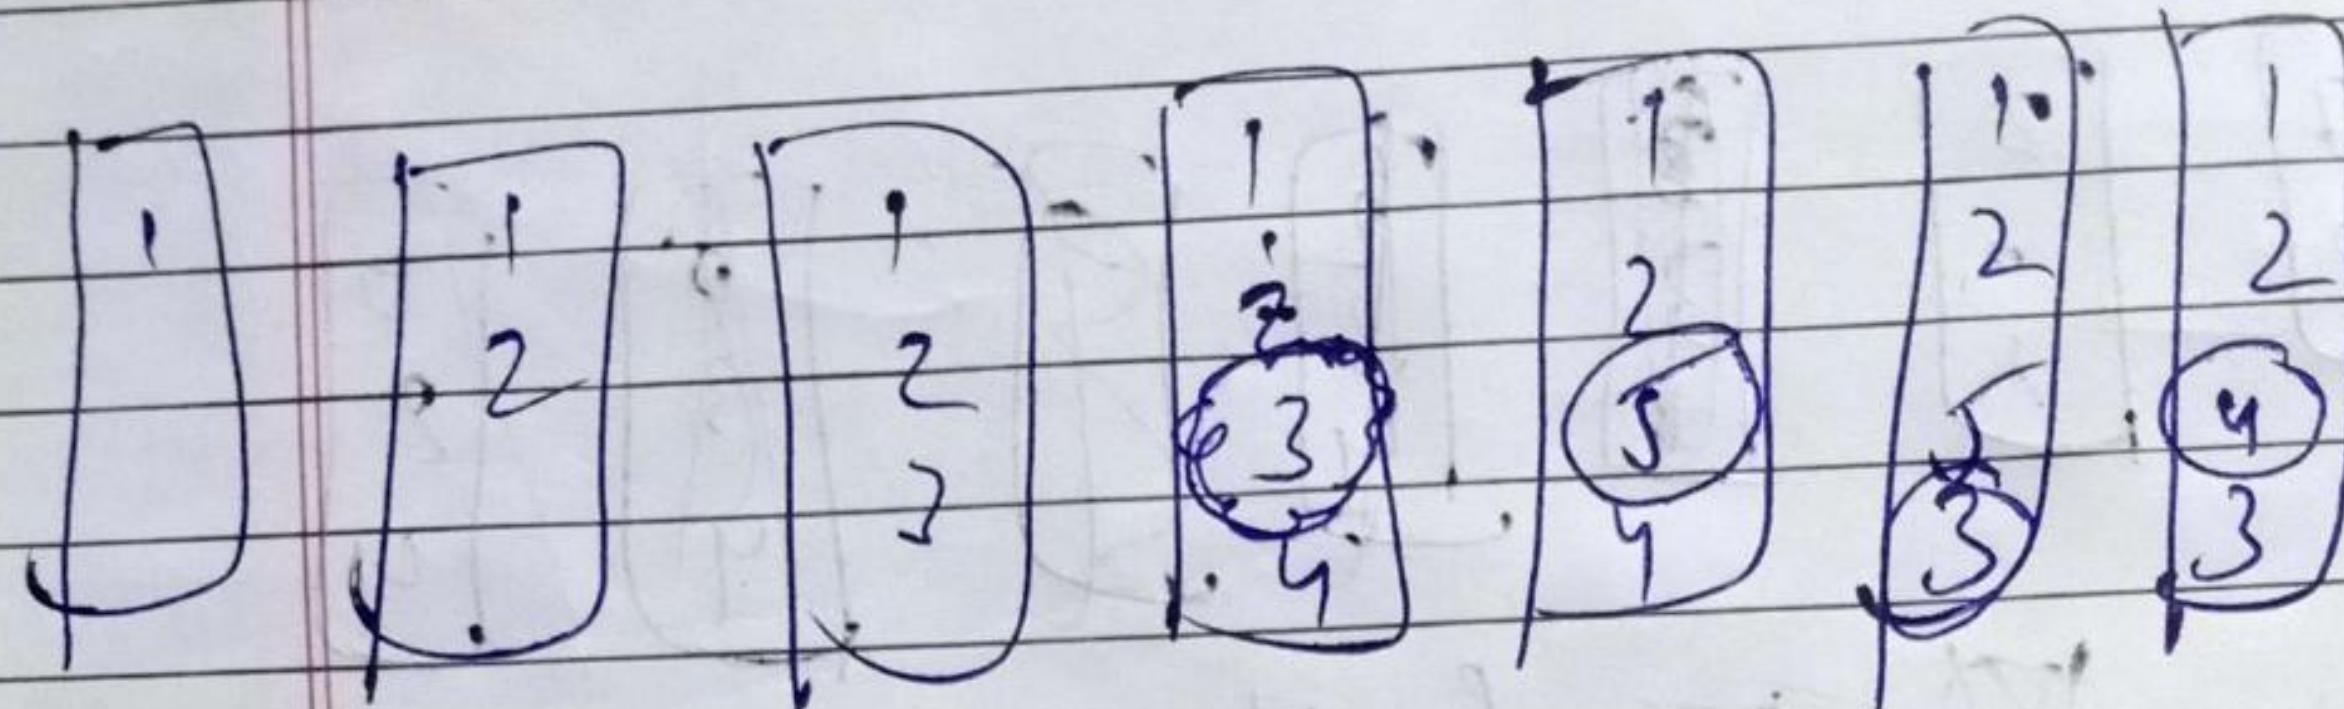
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
in frames



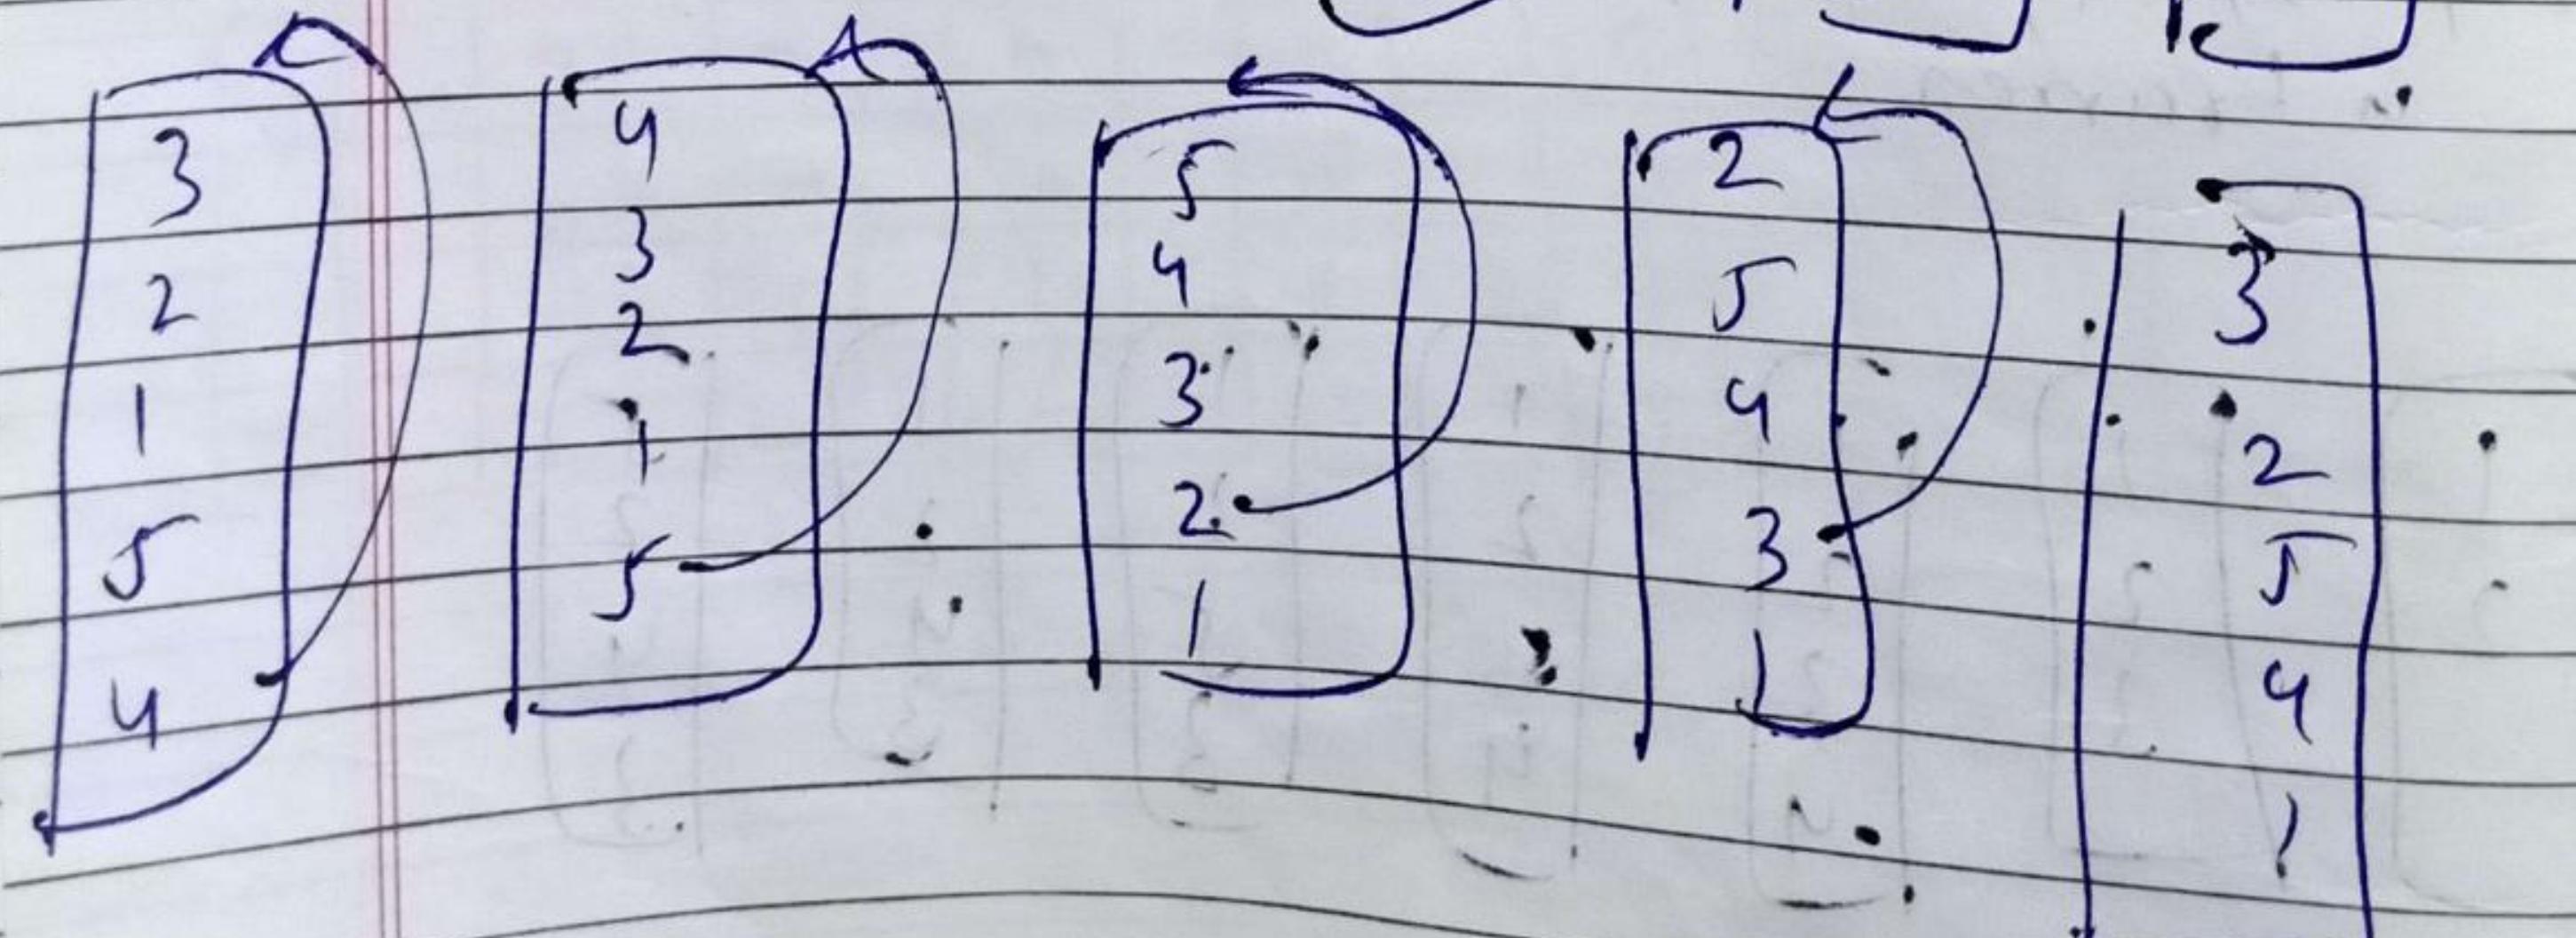
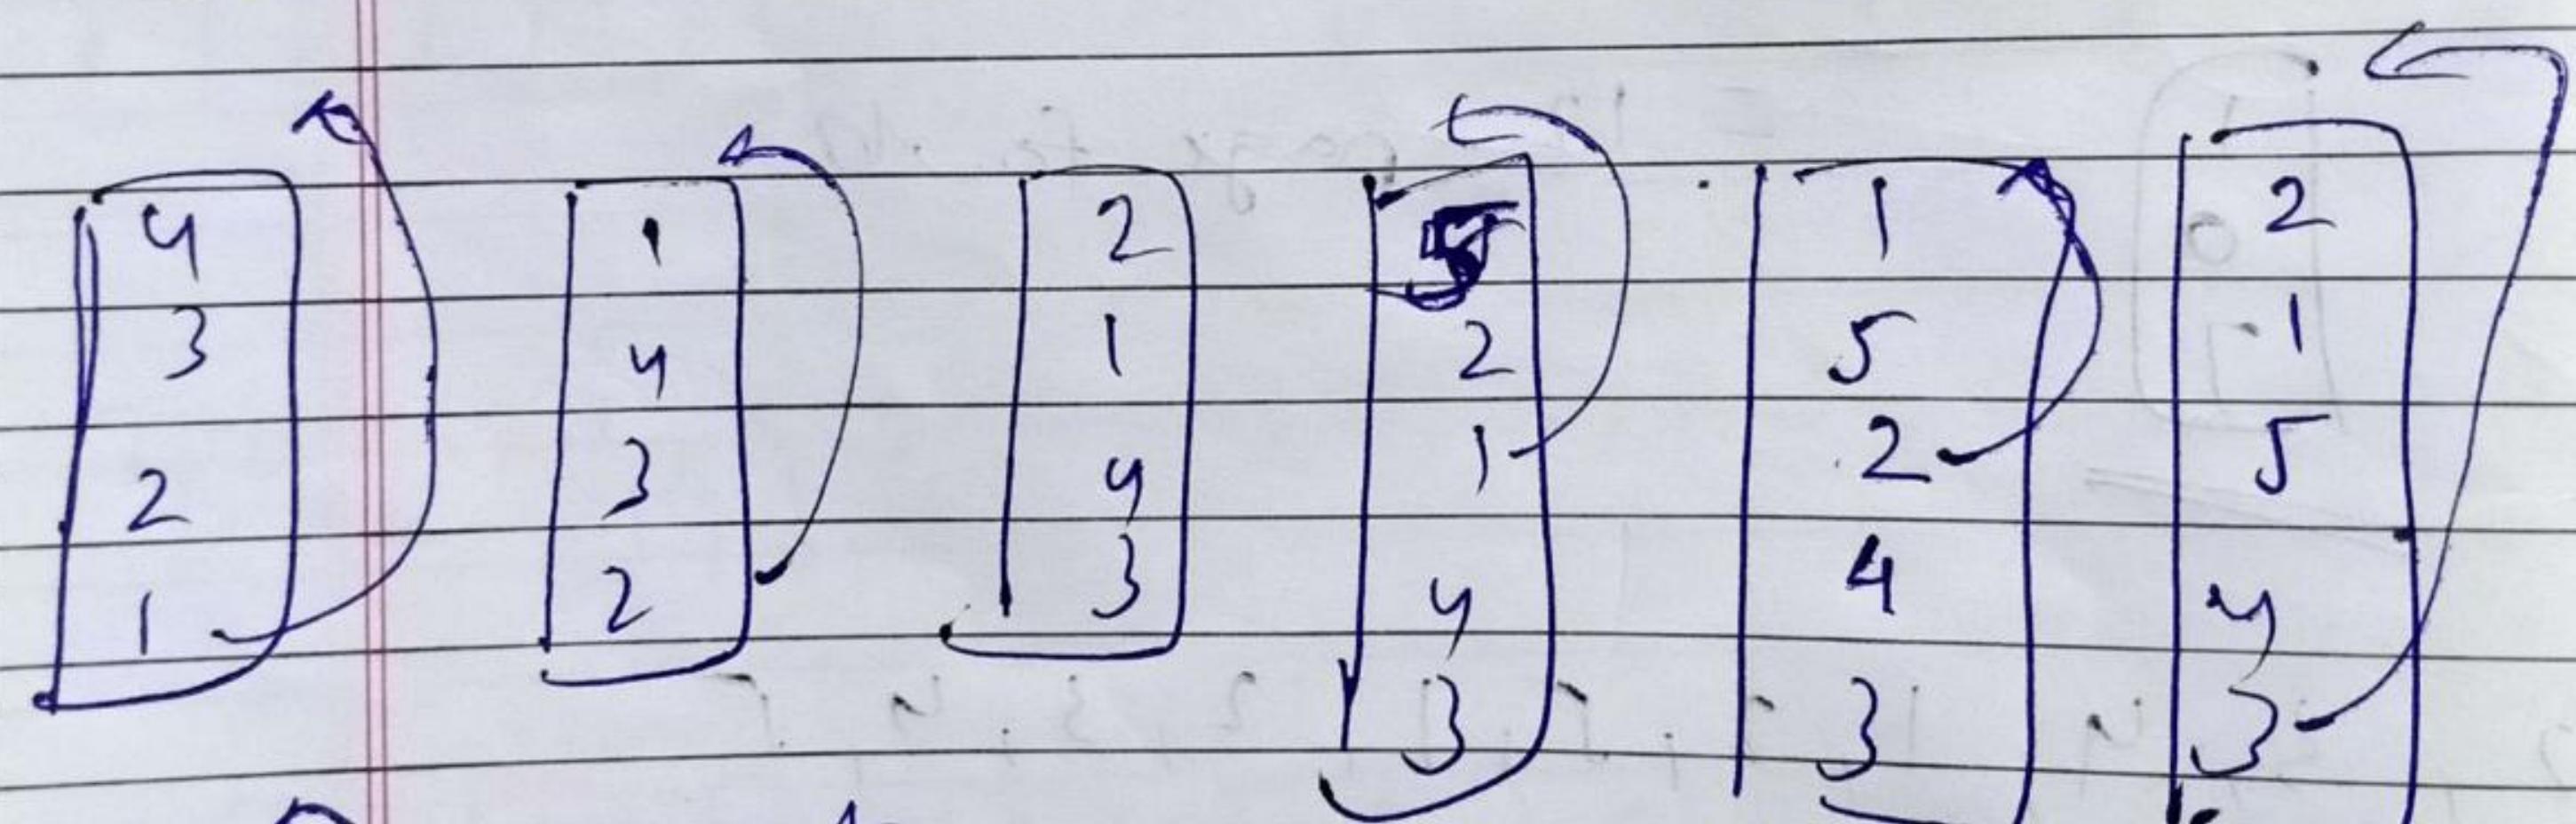
= 8 page faults

LRU implementation.

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5, 2, 3
4 frames



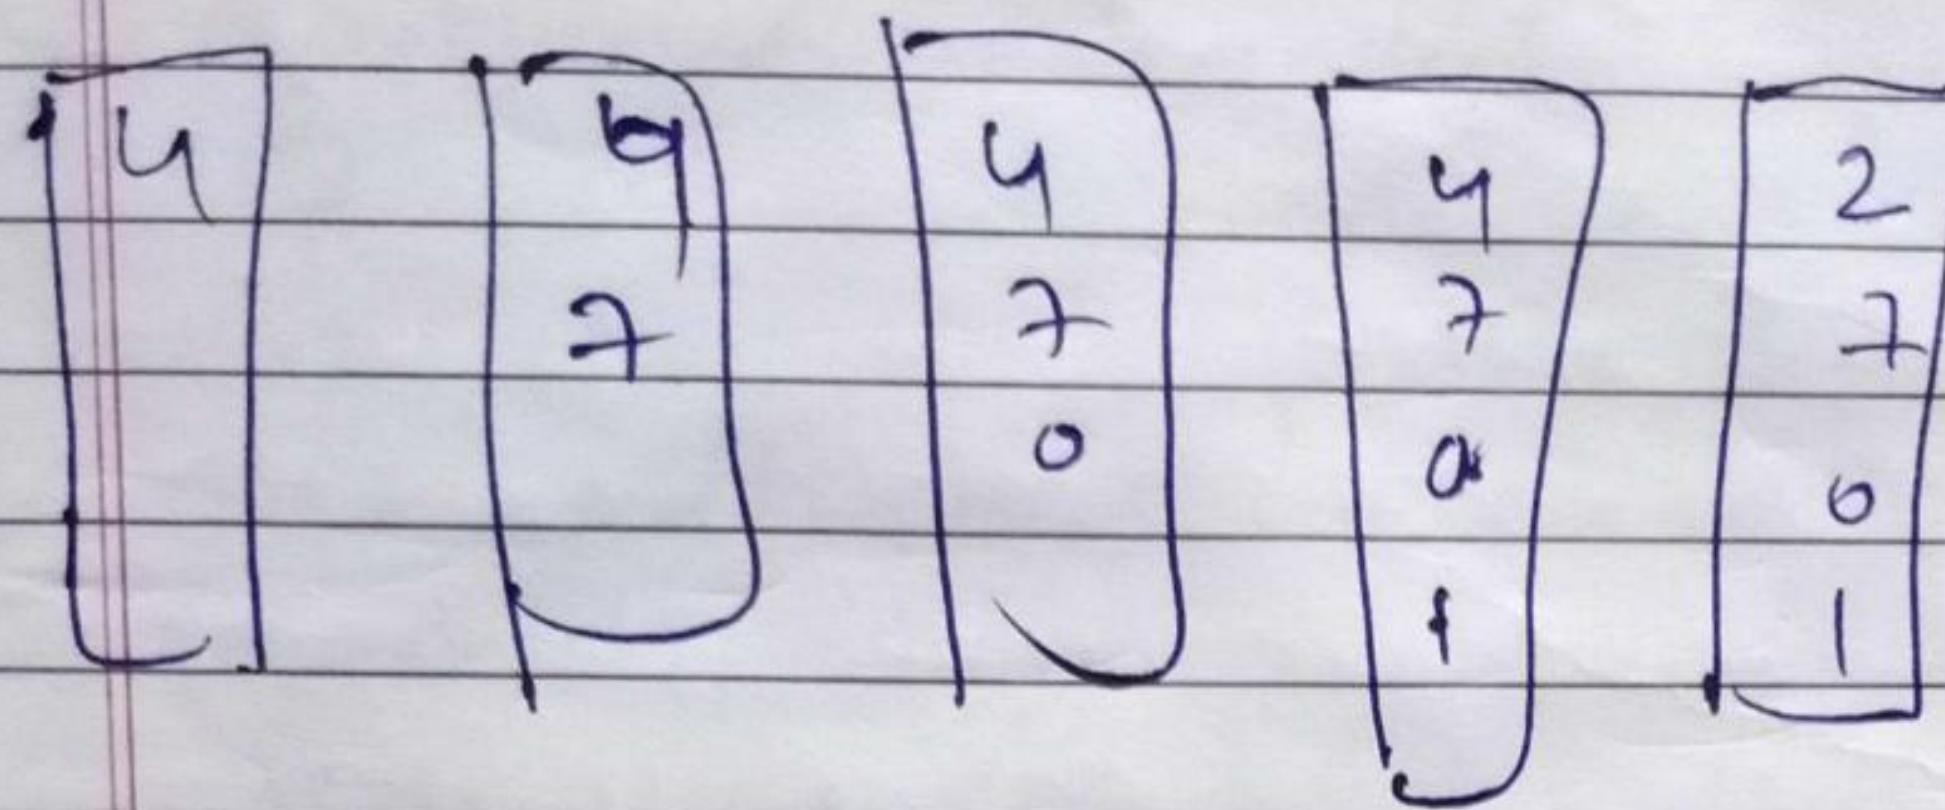
① Stack implementation:



CW implementar

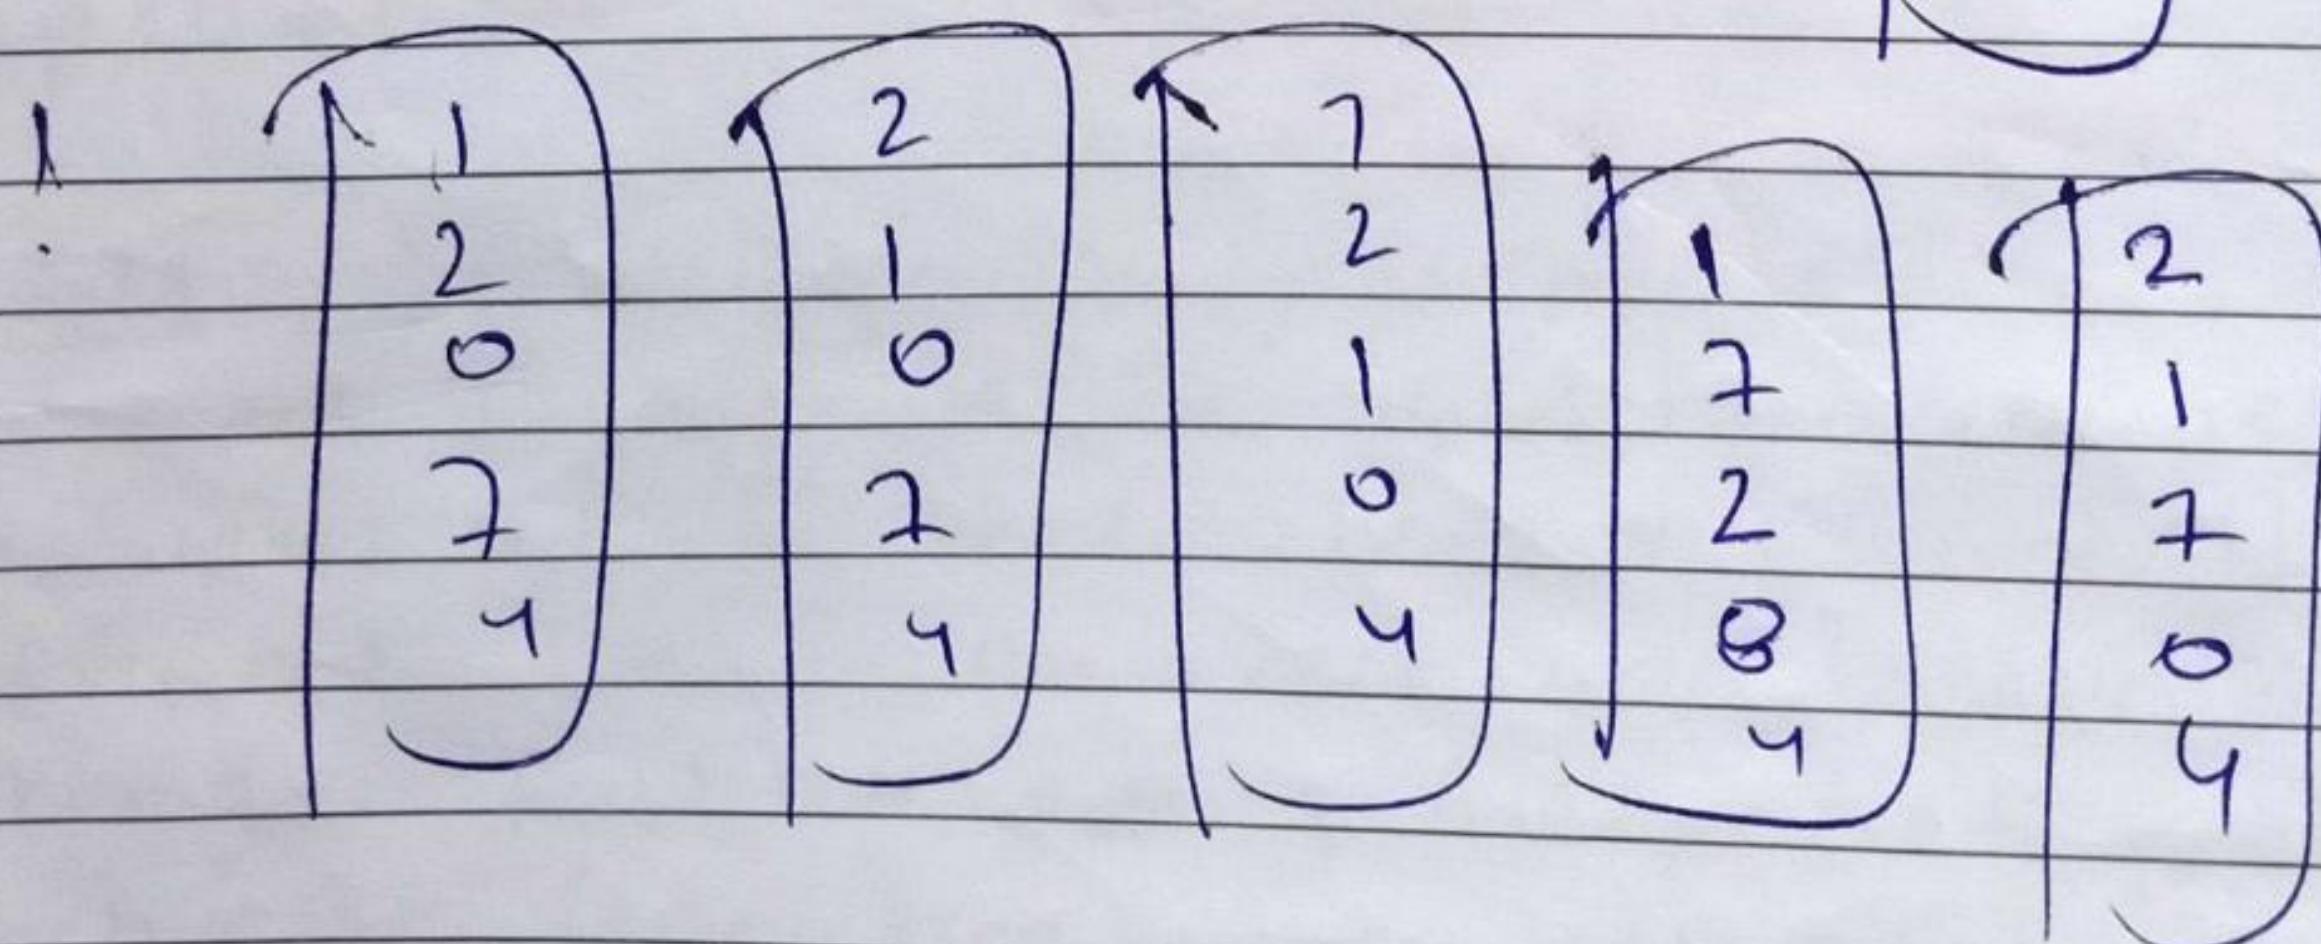
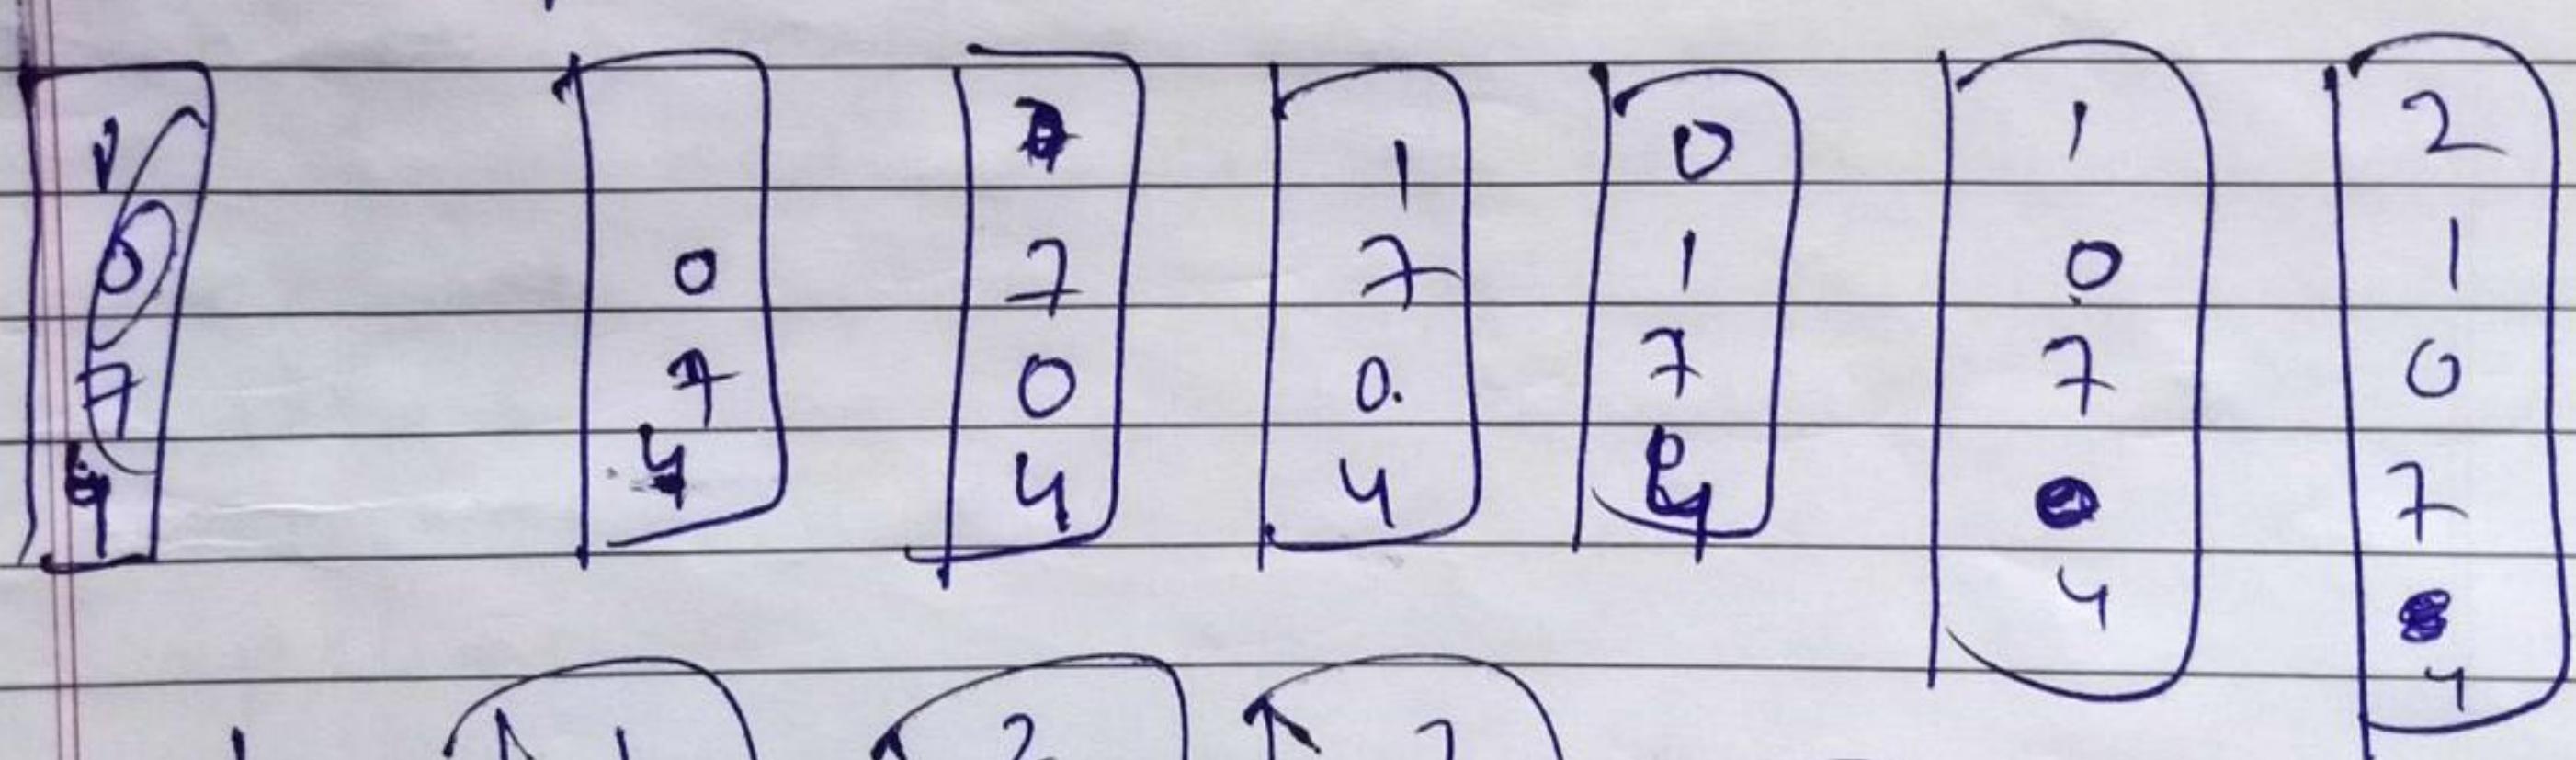
4, 7, 0, 7, 1, 0, 1, 2, 1, 2, 7, 1, 2

Page No.	17	1	2
Date			



= 5 page fa

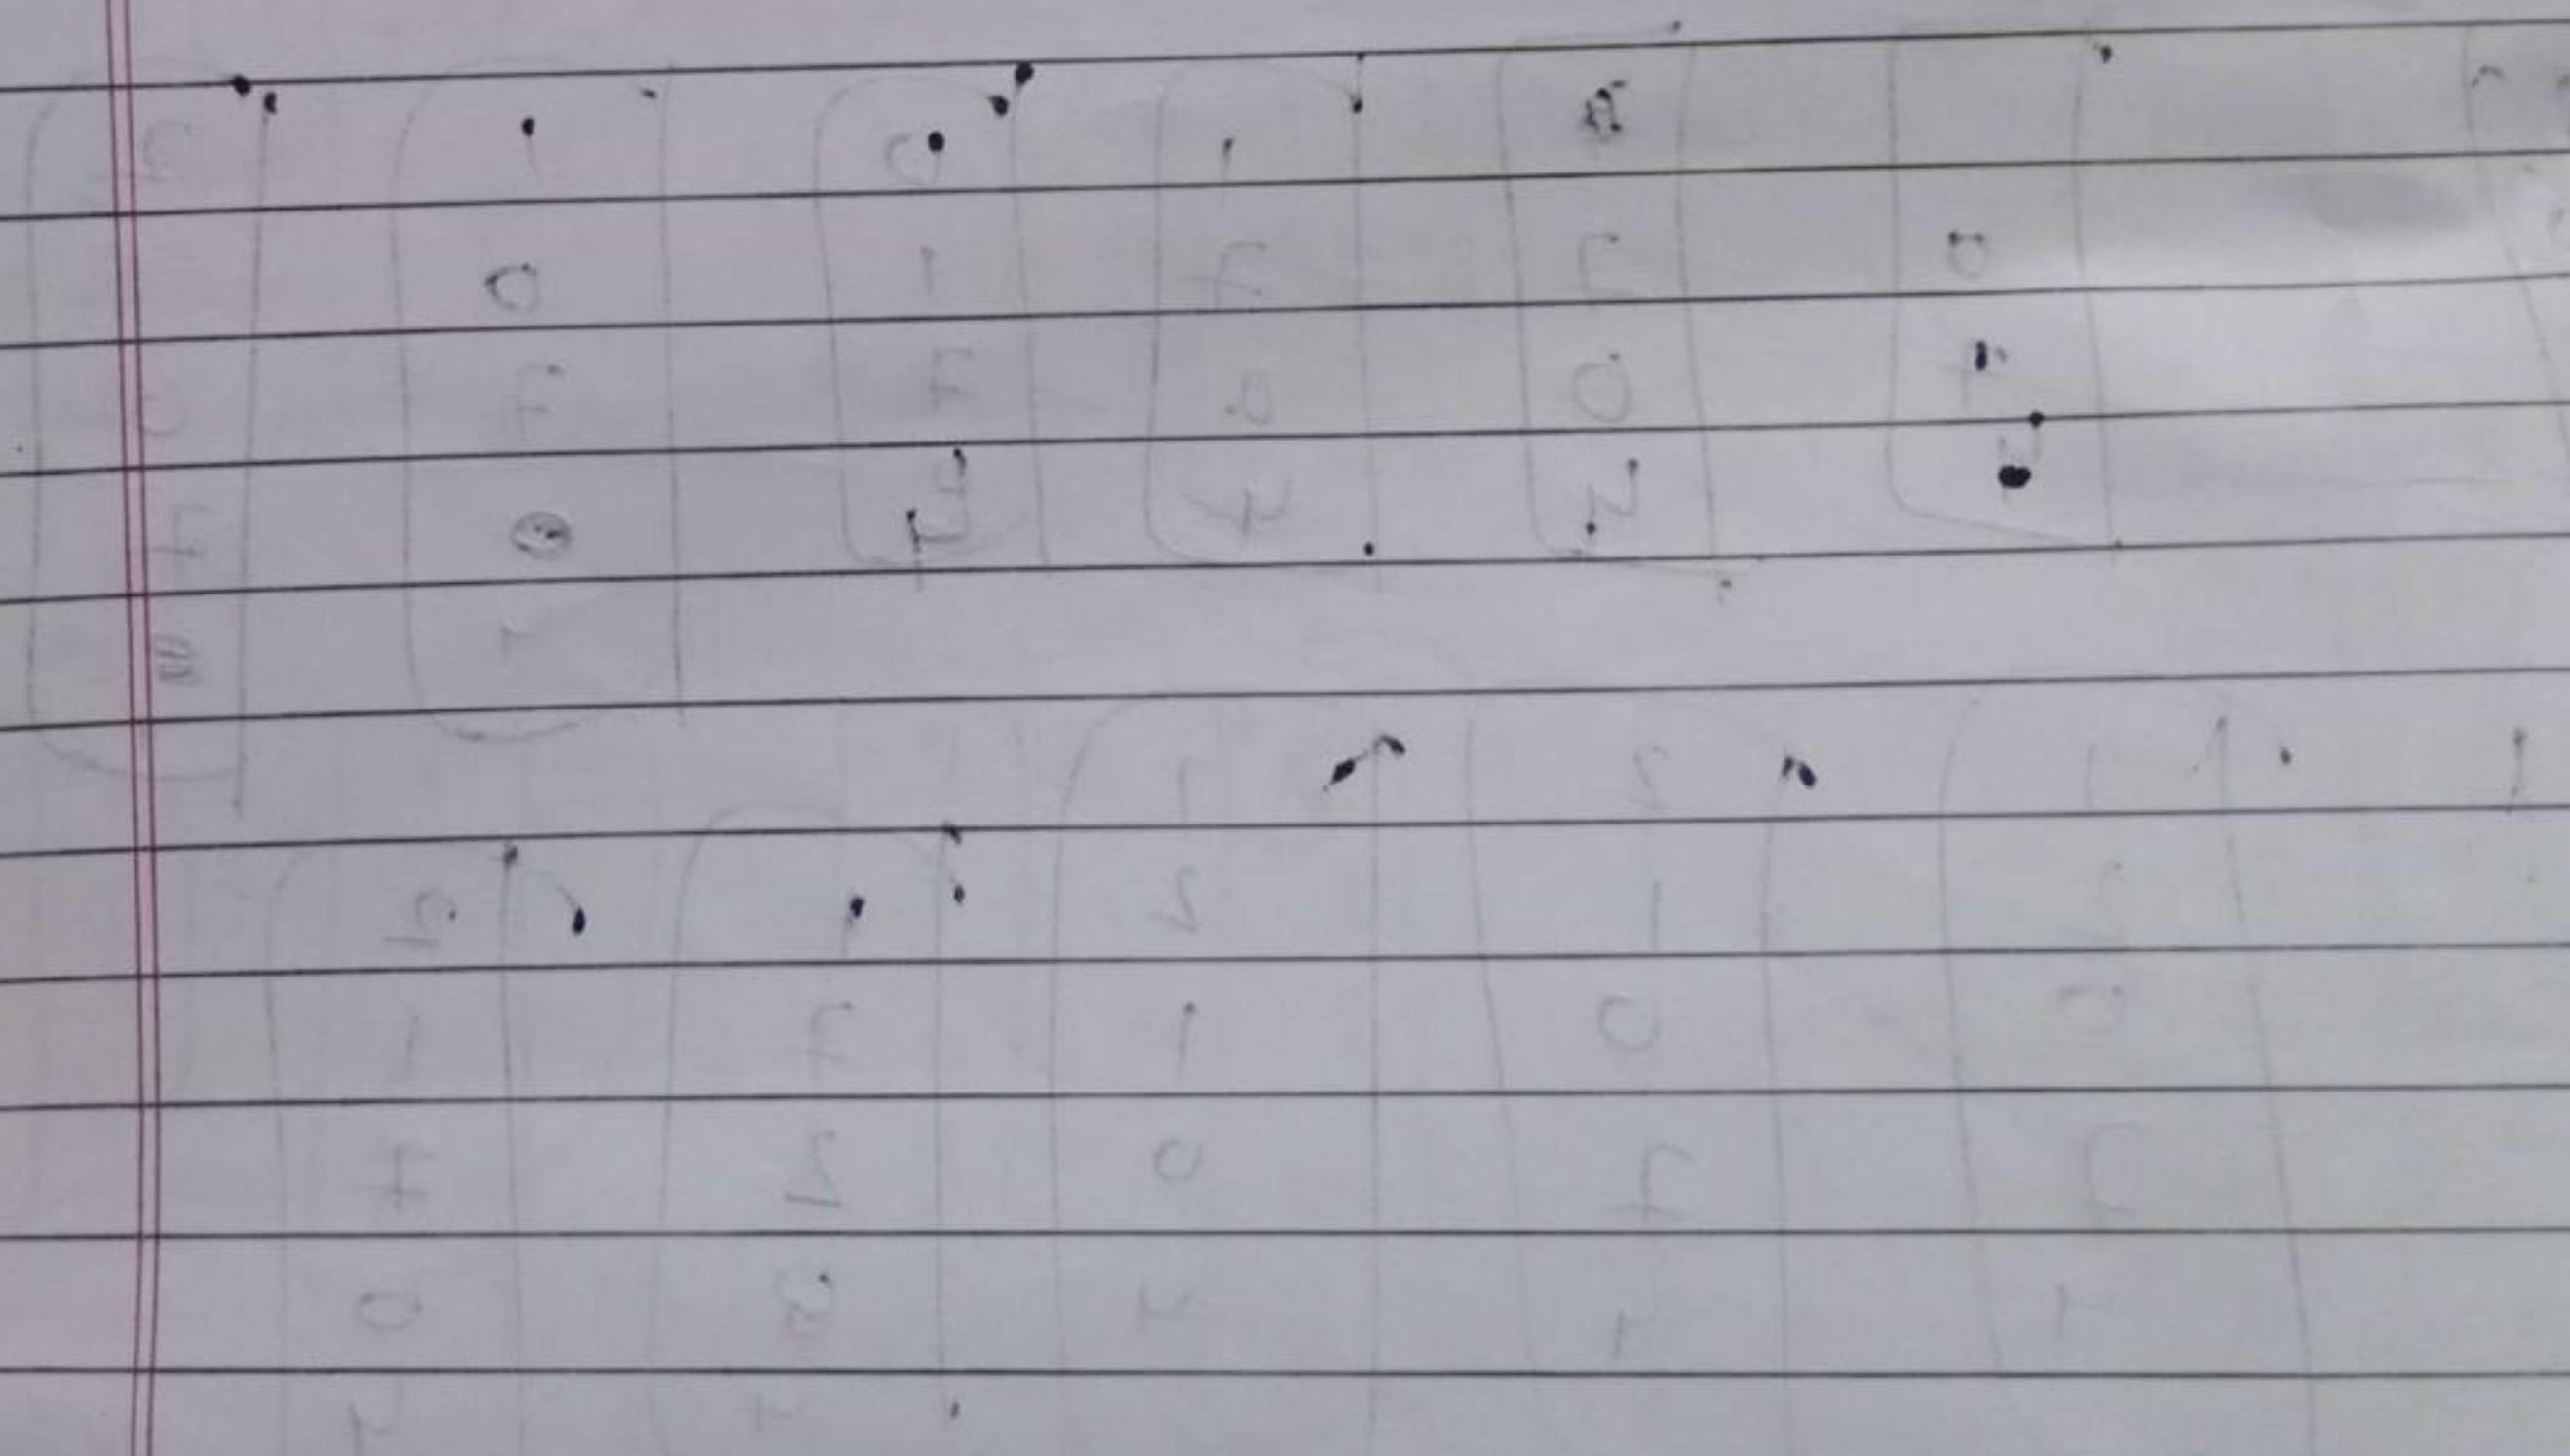
stack implementar



A counter implementation.

Page No. _____
Date _____

A second chance algorithm



Distributed system

Collection of independent computers that appear to

be used as a single coherent system

consist of computers that are autonomous

The differences between computers and the ways
in which they communicate are hidden from the
users. Users and applications can interact with a
distributed system in a consistent and uniform
way regardless of where and when interaction takes
place. Distributed systems should also be
relatively easy to expand or scale.

A distributed system will be continuously available.
Users and applications should not notice that
parts are being replaced or fixed, or that new
parts are added to serve more users or
applications.

* Goals: 1) making resources available
To make it easy for the users to access remote
resources and to share them in a controlled and
efficient way. Resources can be printers, computers,
storage facilities, data, files, web pages and
networks. Connecting users, users and resources
also makes it easier to collaborate and exchange
information. As connectivity and sharing increases,
security becomes important. So systems must
provide protection against eavesdropping or
intrusion in communication.

VVVJP

* Distribution Transparency: To hide that the fact
that its processes and resources are physically
distributed across multiple computers.

A distributed system that is able to present itself to users and applications as if it were only a single computer system is said to be transparent

Page No.	
Date	

	Description
process	Hide differences in data representation and how a resource is accessed.
location	Hide where a resource is located
migration	Hide that a resource may move to another location
relocation	Hide that a resource may be moved to another location while in use
replication	Hide that a resource is replicated
concurrency	Hide that a resource may be shared by several competitive users
failure	Hide the failure and recovery of a resource.

- 2 1) c) page ✓
- 2) a) internal X
- 3) b) hit ratio ✓
- 4) c) valid - invalid ✓
- 5) has data in it X
- 6) d) segment length ✓
- 7) over address

16) c) ✓

17) a) ✓

b) To not execute another form
deadlock X

c) ✓

20) c) max - wildcard ✓

sort by default nature:- first whitespace characters
then wildcard characters.

3) openness:- An open distributed system is a system that offers services distributed according to standard rules that describe the syntax and semantics of those services - In computer networks, standard rules govern the format, contents and meaning of messages sent and received. Such rules are formalized in protocols. In distributed systems, services are specified through interfaces, which are described in an Interface definition language (IDL)

IDLs described an interface in a language-independent way, enabling communication between software components that do not share one language. Interface definitions specify the names of the functions that are available, with types of the parameters, return values, possible exceptions that can be raised.

If properly specified, an interface definition allows an arbitrary process that needs a certain interface to talk to another process that provides that interface. An IDL is used to set up communications between client and server in remote procedure calls (RPC).

Another imp' goal for an open distributed system is that it should be easy to configure the system out of different components.

should be easy to add new components or replace existing ones without affecting those components.

that stay in place. An open distributed system should also be extensible. To achieve flexibility in open distributed systems, it is crucial that the system is organized as a collection of relatively small and easily replaceable or adaptable components.

4) Scalability :- scalability of a system can be measured along the following dimensions. A system can be scalable with respect to its size. We can easily add more users and resources to the system. A geographically scalable system is the one in which the users and resources may lie far apart. A system can be administratively scalable. A system that is scalable in one or more of these dimensions often exhibits some loss of performance as the system scales up.

o Scalability problems : 1) Scaling w.r.t. size. The servers can become a bottleneck as the no. of users and applications grow. Using only a single server is sometimes unavoidable. Eg: managing highly confidential info such as medical records, bank accounts. In such cases, copying the server to several locations to enhance performance maybe out of the question as it would make the services less secure. Many centralized services are implemented by means of only a single server running on a specific machine in the distributed system.

Ex: Even if one has virtually unlimited processing and storage capacity, communication with that server will eventually prohibit further growth.

Advantages of distributed systems over independent PC's

- 1) Accessibility.
- 2) More reliability, reliability.
- 3) More interaction between users.
- 4) Availability
- 5) Easily extendable

Hardware concepts :-

Scalability operation

$$\begin{array}{l} A_0 + B_0 = C_0 \\ A_1 + B_1 = C_1 \\ A_2 + B_2 = C_2 \end{array}$$

$$\begin{array}{l} A_0 + B_0 = C_0 \\ A_1 + B_1 = C_1 \\ A_2 + B_2 = C_2 \end{array}$$

Design Issues

1) Transparency

Kind

meaning

location transparency users cannot tell where resources are located

migration transparency resources can move at will without changing their names.

replication transparency

The users cannot tell how many co-exist multiple users can share resources automatically

concurrency transparency

Activities can happen in parallel without users knowing.

parallelism transparency

Flexibility :

Reliability:

- Availability

Page No.		
Date		

- Reliability
 - Availability
 - Redundancy
 - Security
 - Fault tolerance
- Performance
 - Performance metrics
 - Response time
 - Throughput
 - System utilization
 - Amount of network capacity consumed
 - Are grained parallelism.
 - Coarse grained parallelism.
- Scalability

Challenges: The network is not reliable
the network is not secure
") not homogeneous.

The topology changes
latency due to distance
Bandwidth is not infinite
Transport cost
there is no administrator

File

- A named collection of related info. that is recorded on secondary storage
- Files represent programs and data.
- A file has a certain defined structure according

» Attributes

Name - only info. kept in human readable form

Identifier - this unique tag identifies the file within the file system

Type - needed

location -

size -

protection -

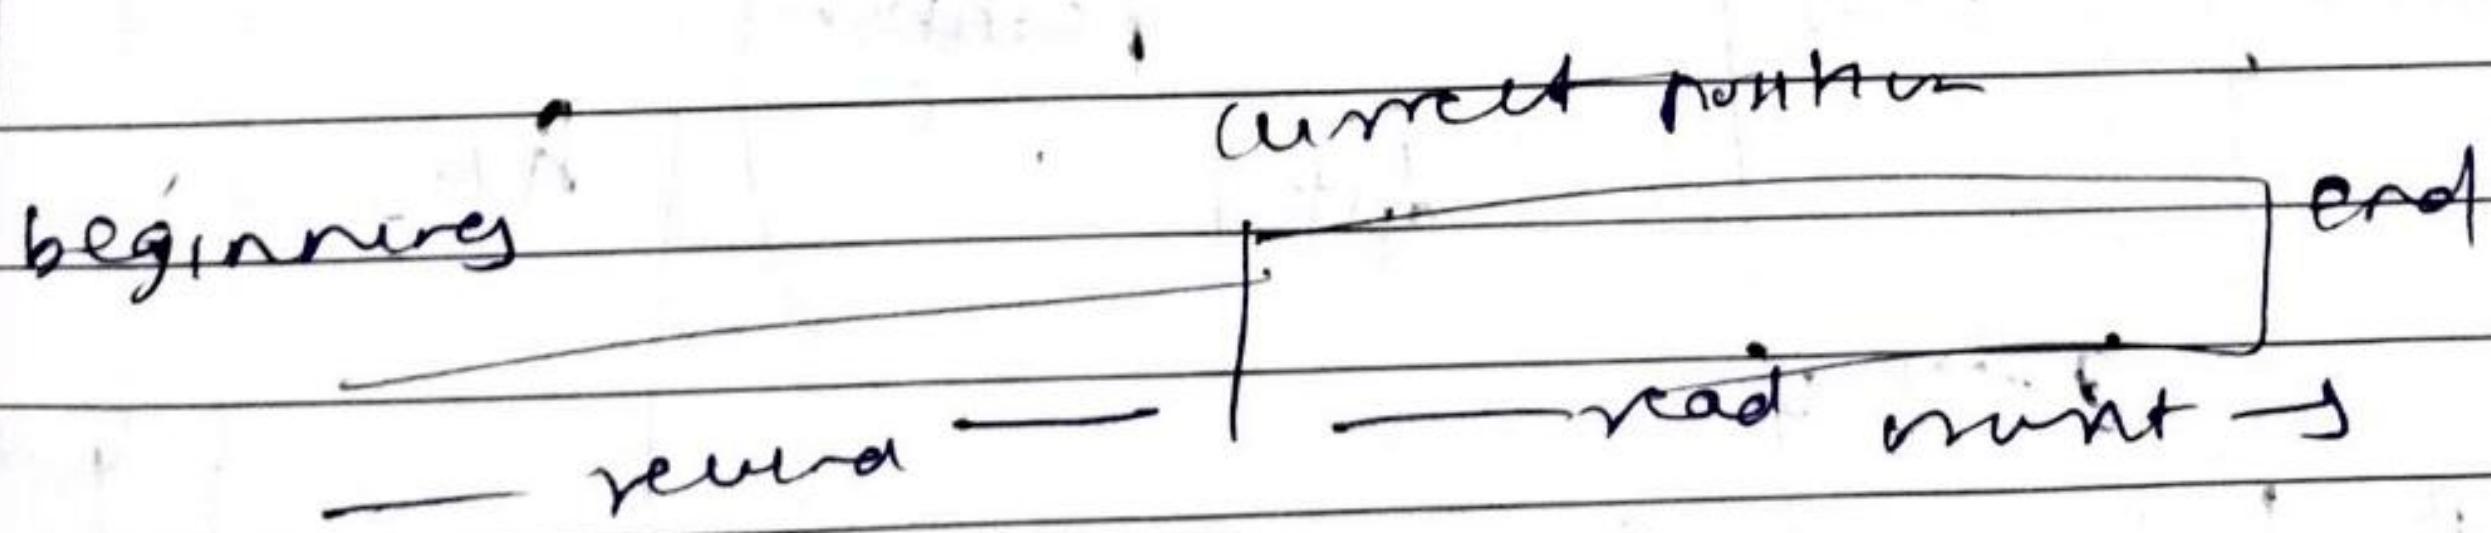
Time, date and user identification

» File operations

create, write, read, reposition within file

delete, Truncate,

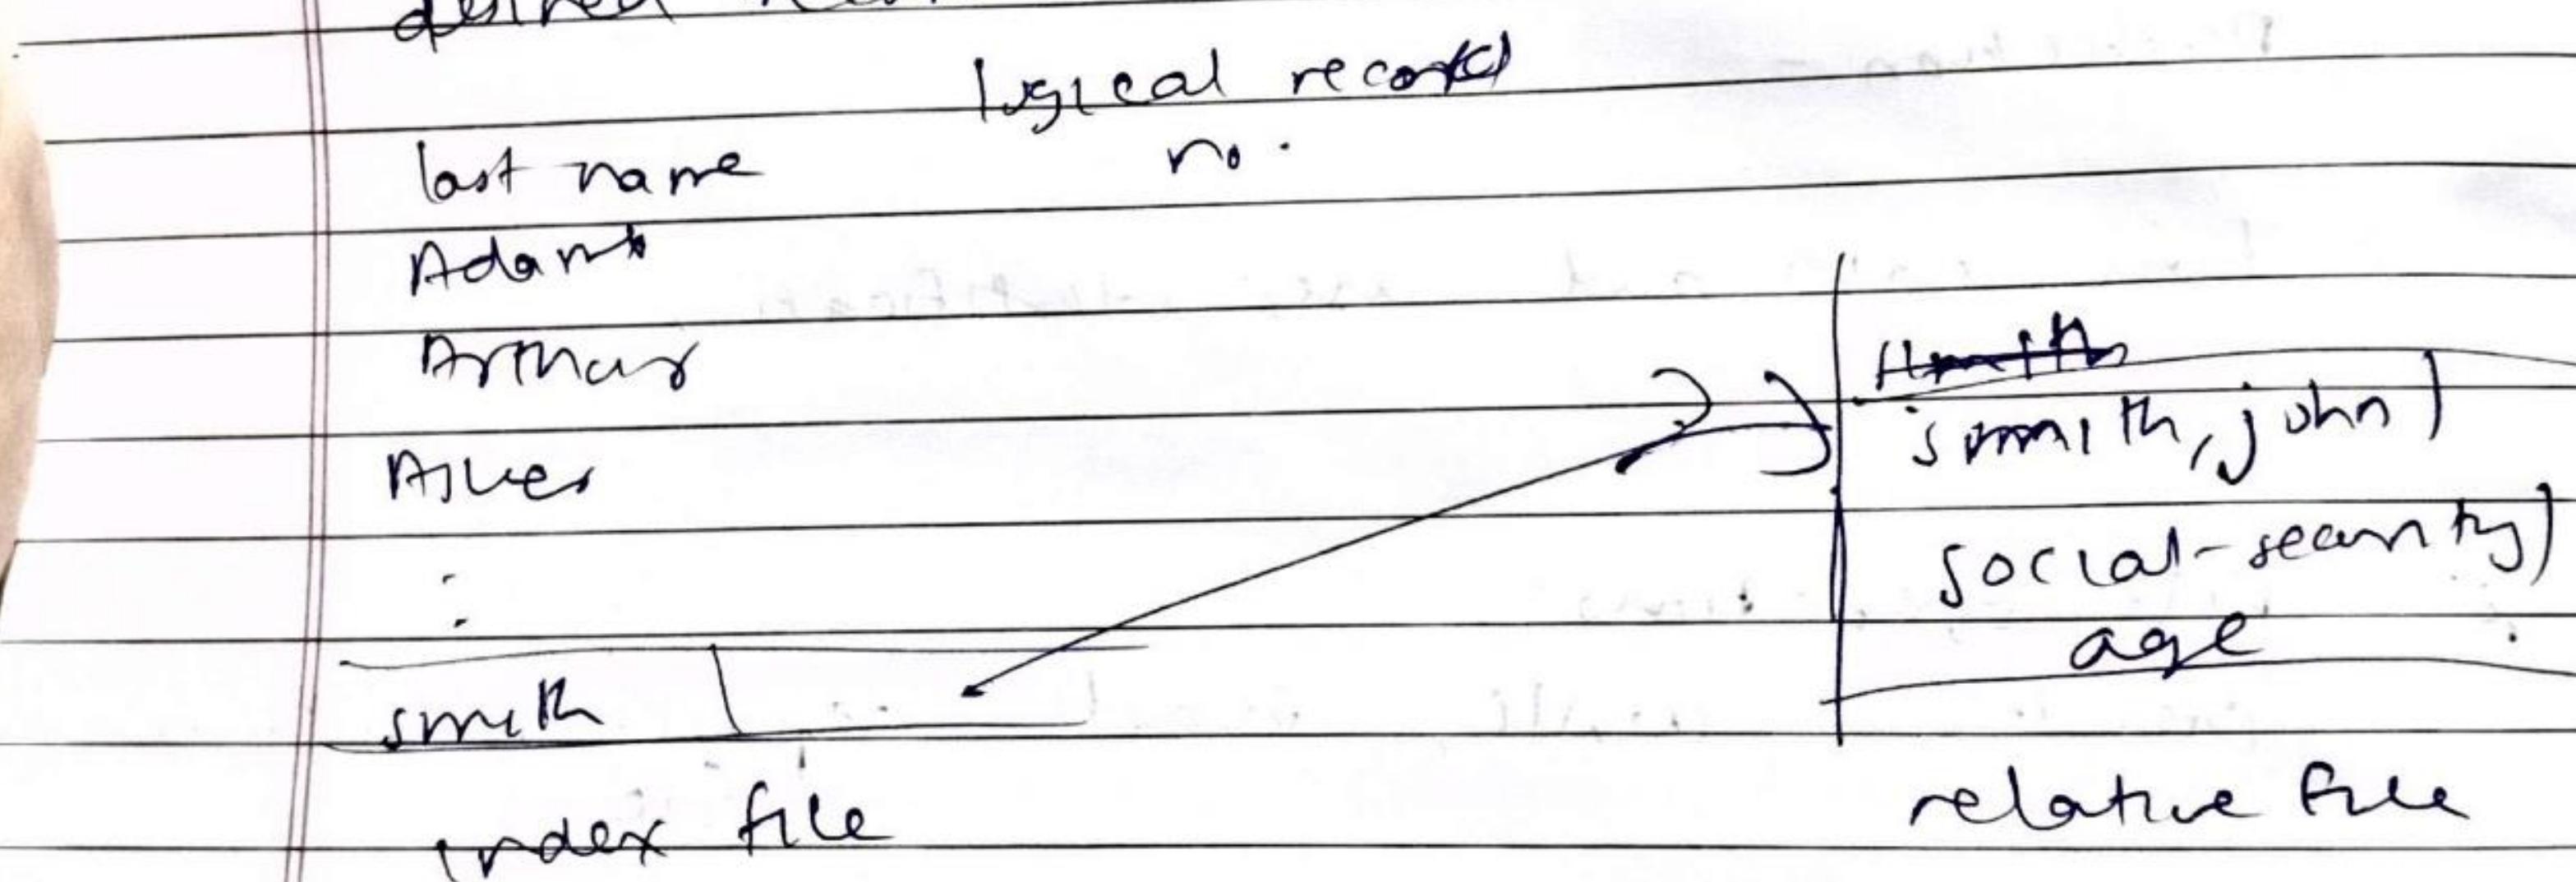
» Access methods - 1) Sequential access - Info. in the file is processed in order, one record after the other.



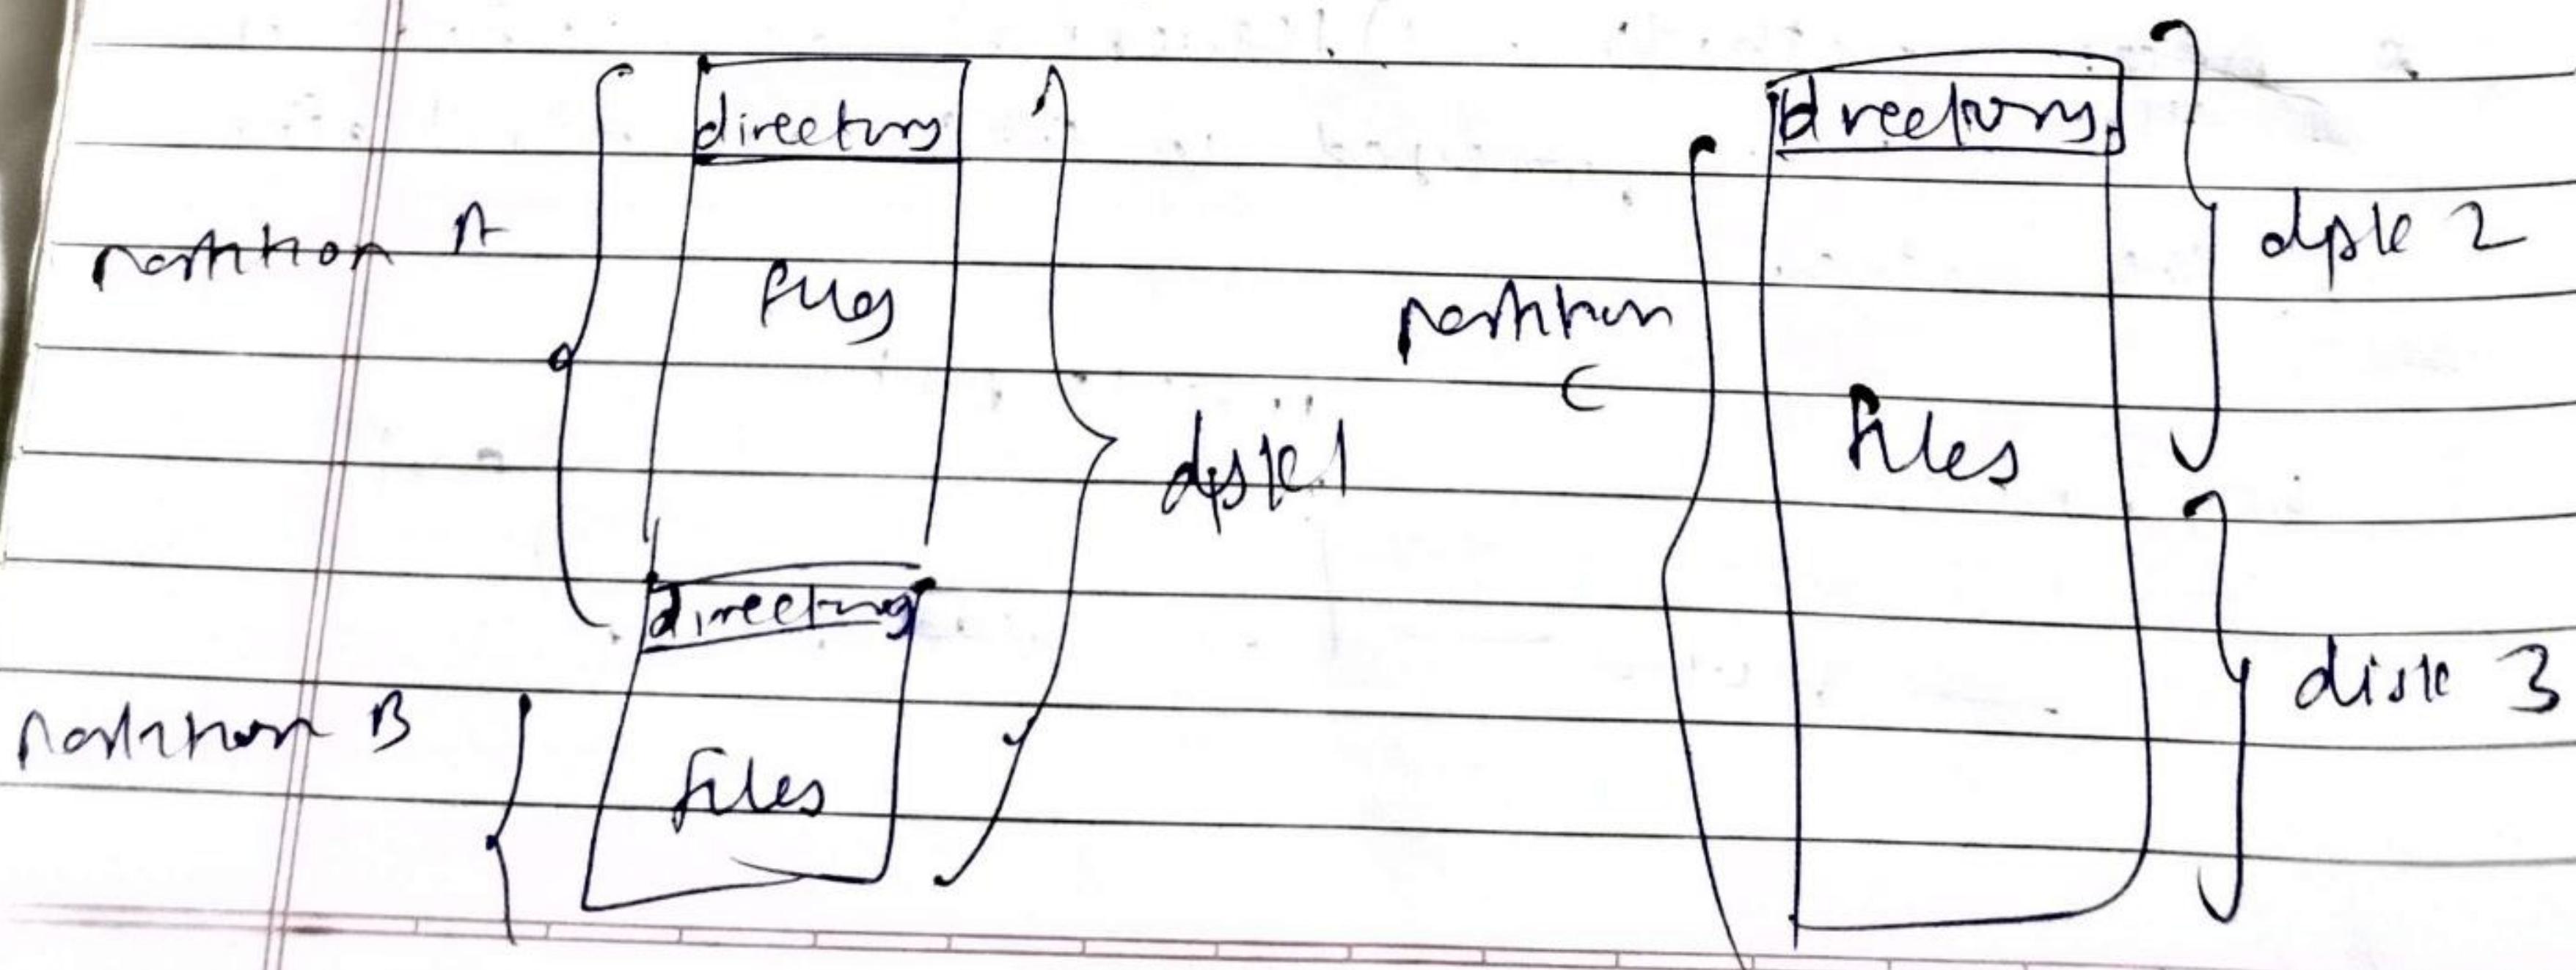
1) Direct access: A file is made up of fixed length logical records that allow programs to read and write records rapidly in no particular order it viewed as a numbered sequence of blocks or records. The direct access method allows random access access to any file block. There is no order in restriction for reading or writing.

Eg: Index and relative files

~~It involves the construction~~ construction of an index for the file. The index contains pointers to access the file directly and find the desired record.

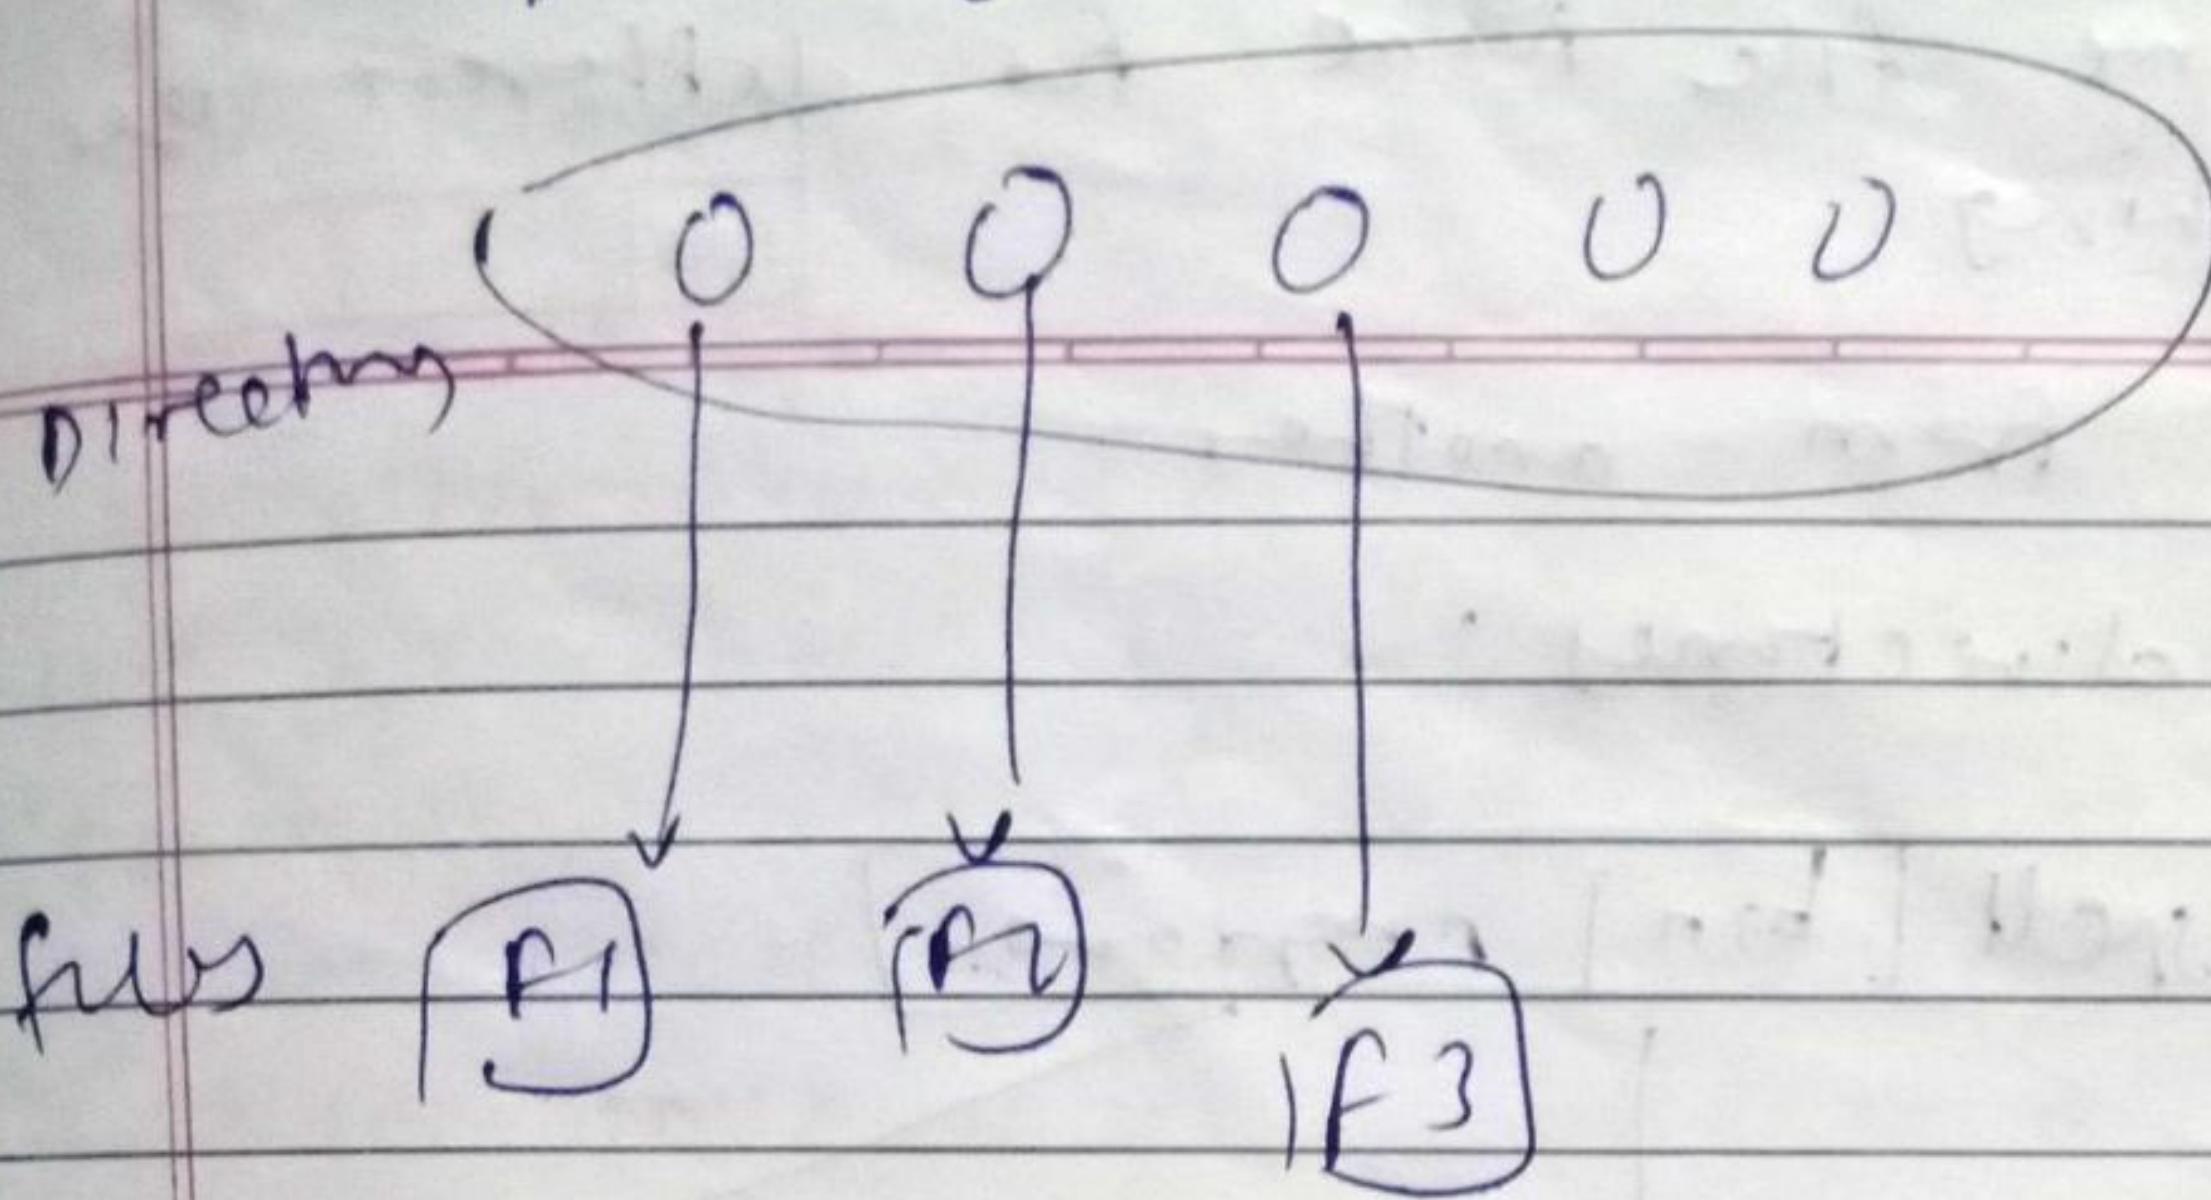


A Typical file-system org.



Directory structure

Page No.		
Date		

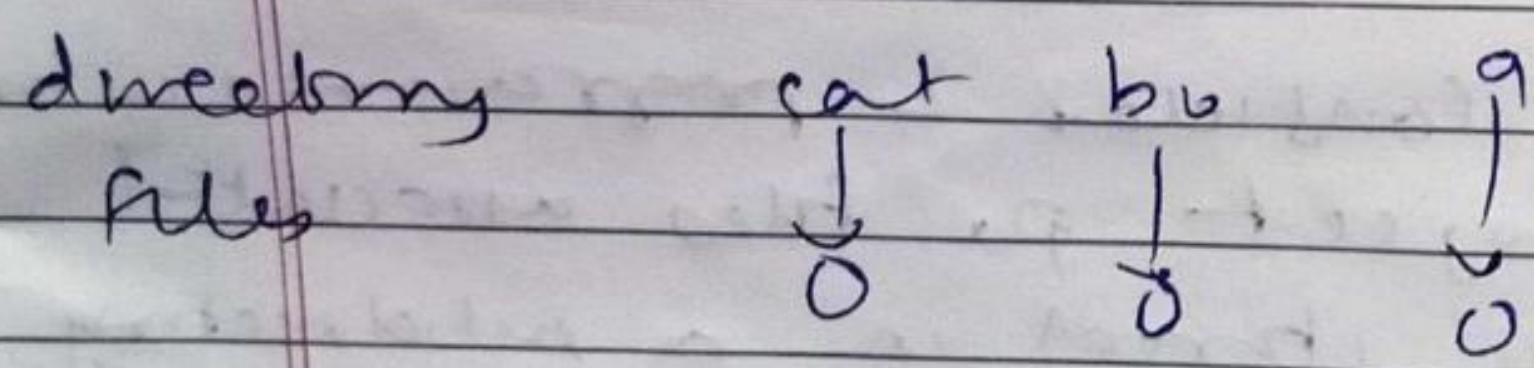


A collection of nodes containing info about all files

Ops. performed on a directory

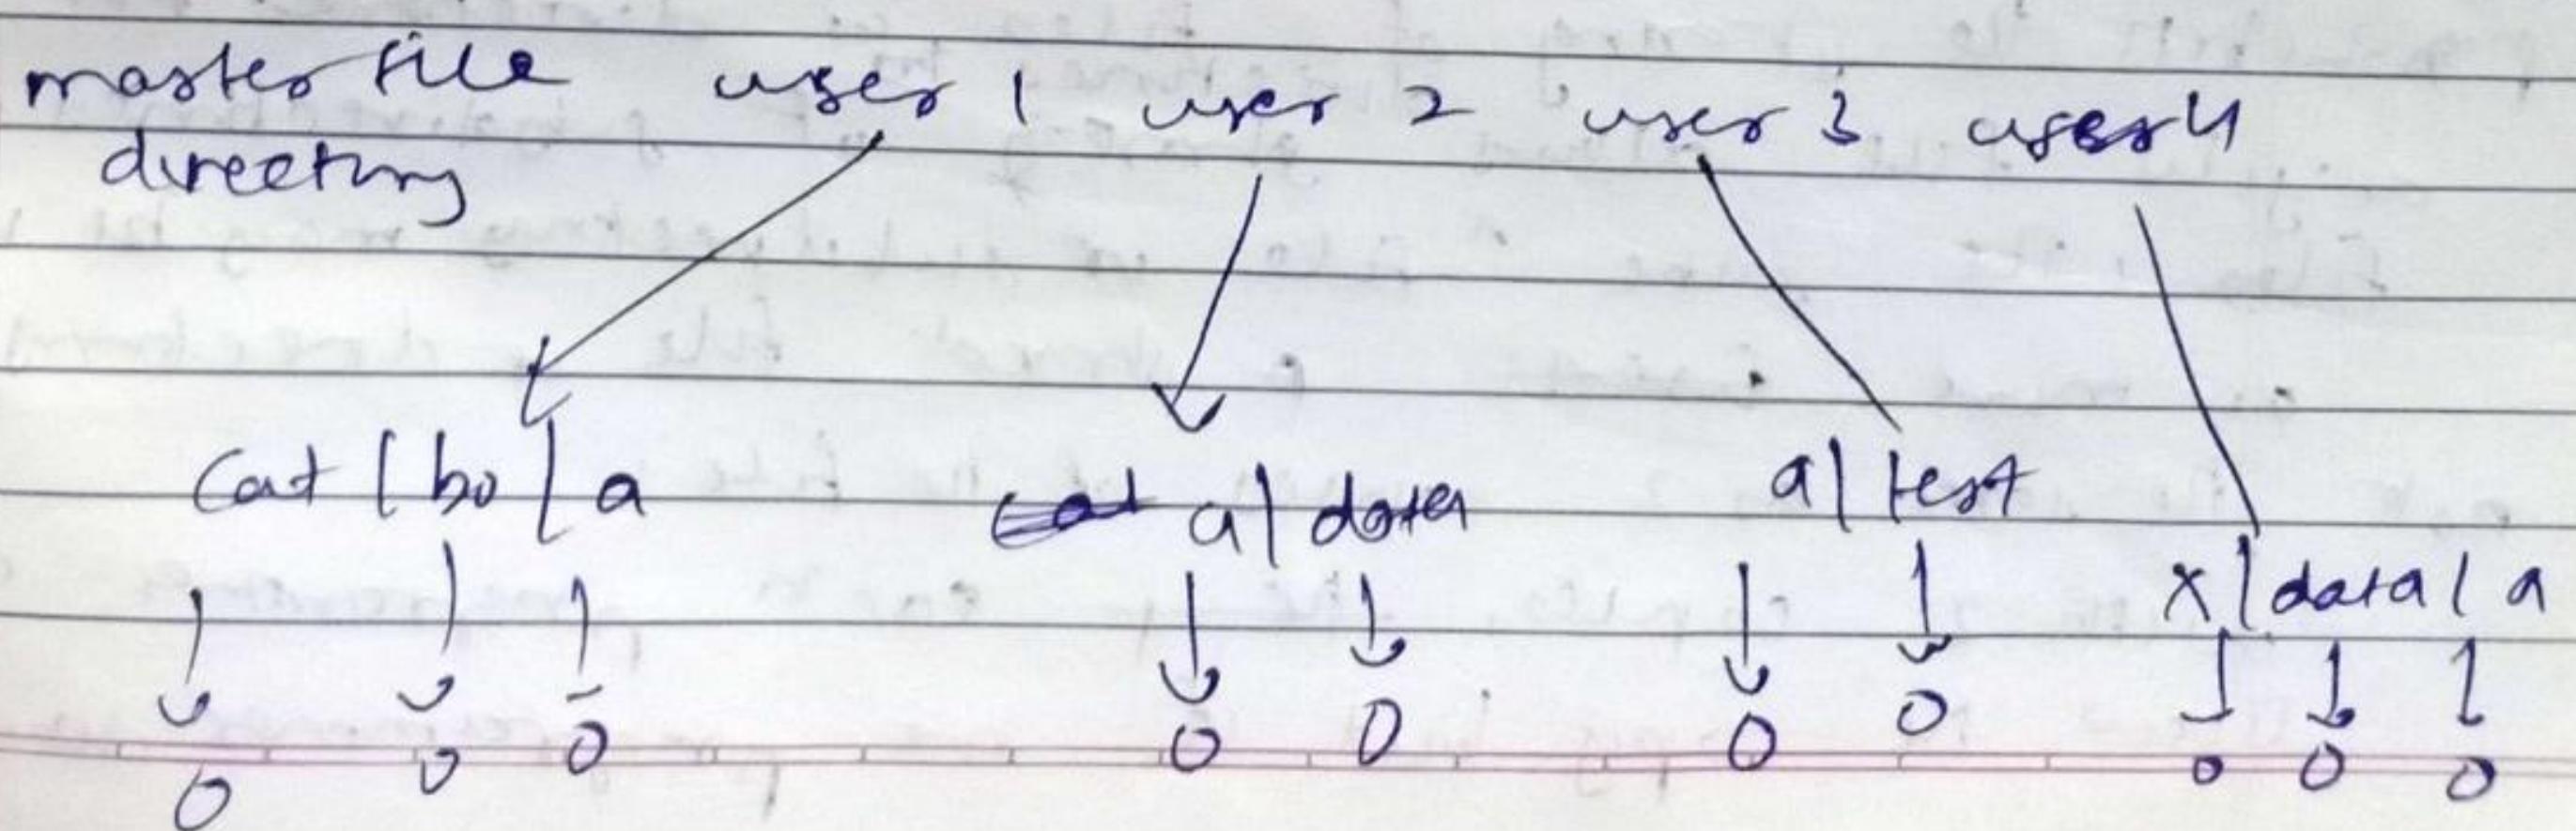
- 1) search for a file
- 2) create a file
- 3) delete a file
- 4) list a directory
- 5) rename a file
- 6) traverse the file system

* single-level directory :- A single directory for all users



~~name naming problem~~
~~grouping problem~~

* Two-level directory :- Separate directory for each user. Each user has its own user file directory (UFD)

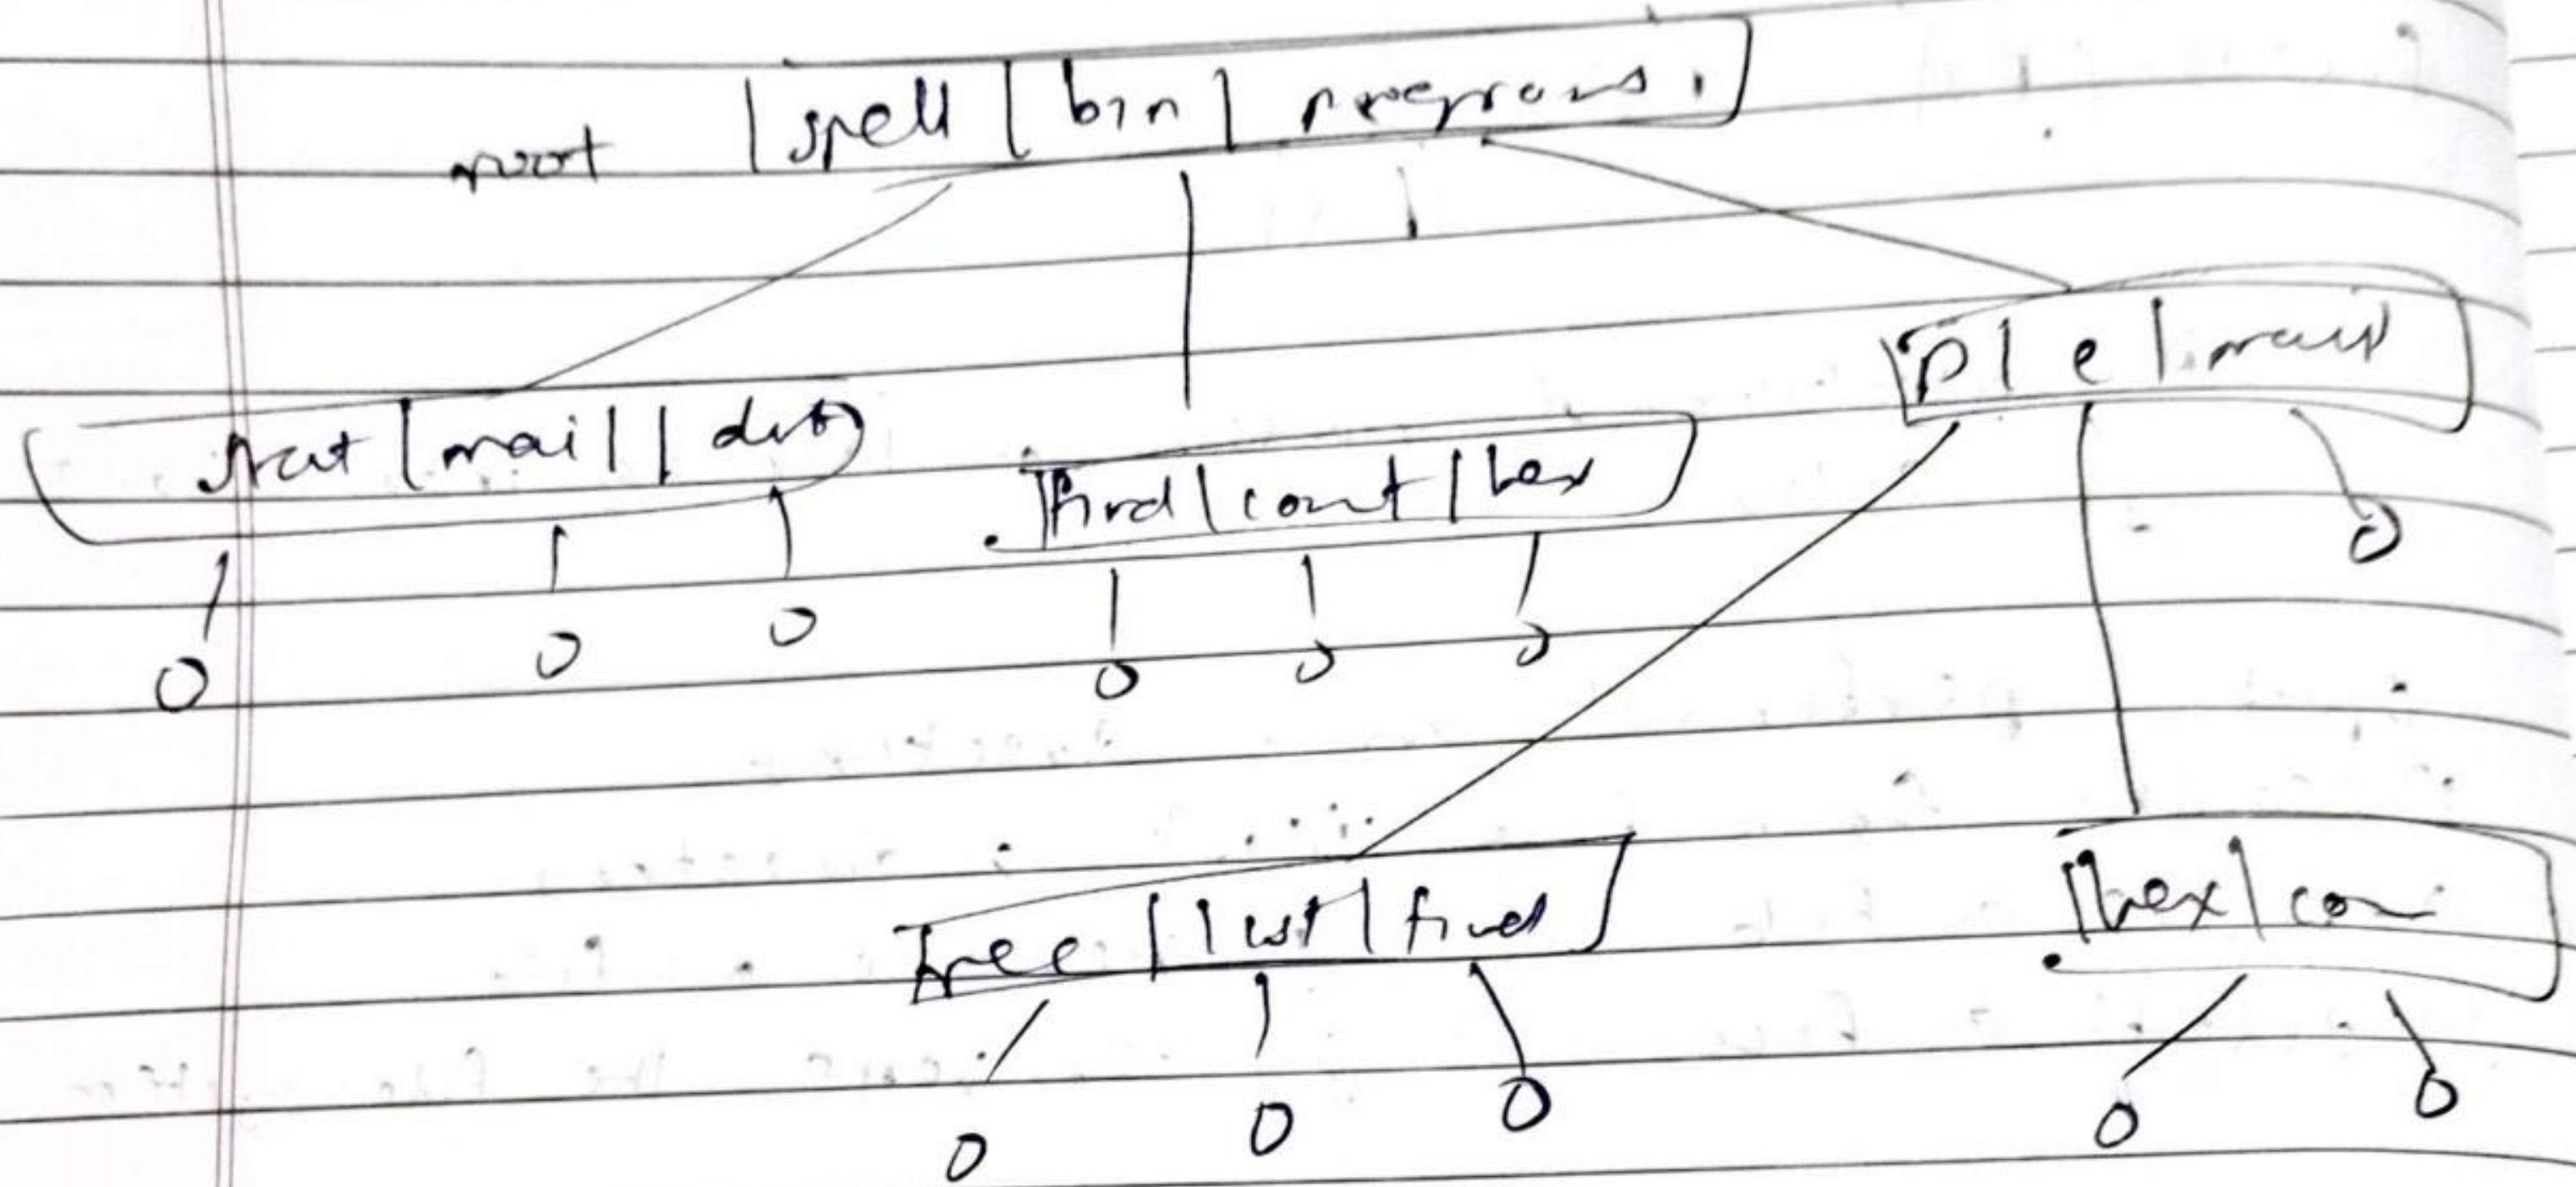


Path name

- Can have the same file name for different user
- Efficient searching
- Isolates one user from another

Page No.		
Date		

• tree-structured directries:



Relative path name, Absolute path name.

* Acyclic - graph directries : Consider 2 programmers who are working on a joint project. If files associated with that project can be stored in a subdirectory, separating them from other projects and files of the 2 programmers. But since both programmers are equally responsible for the project, both want the subdirectory to be in their own directries. The common subdirectory should be shared. A tree structure prohibits the sharing of files or directries. An acyclic file allows ^{directries to} sharing of subdirectories and files. The same file or subdirectory may be in 2 or more ~~files~~. A shared file or directory is not the same as 2 copies of the file. With 2 copies, if each programmer can view the copy but if one programmer changes

to file, the changes will not appear in the other's copy with a shared file, only one actual file exists, so any changes made by one person are immediately visible to the other. Sharing is important for sub-directories; a new file created by one person will automatically appear in all the shared sub-directories. UNIX systems create a new directory entry called a link. A link in UNIX is a pointer to a file. Creating links is a kind of shortcut to access a file. Links ~~allow~~ allow more than 1 path name (absolute or relative) to refer to the same file, ~~anywhere~~ elsewhere. When a ref. when a reference to a file is made, we search the directory. If the directory entry is marked as a link, then the name of the real file is included in the link information. We resolve the link by using that path name to locate the real file. Inode structure and soft and hard links - briefly An acyclic-graph directory structure is more flexible than a simple tree structure, but it is also more complex.

why
it
my,
he
we
can